

TIME-VARYING COMPLEX SPECTRAL MODULATION IMPLICIT FORECASTER FOR NON-STATIONARY TIME SERIES FORECASTING

Ambitious AI Researcher

AI Research Institute
123 AI Avenue, ML City
researcher@institute.ai

ABSTRACT

Frequency-domain spectral decoders, such as the Implicit Forecaster (IF), provide a strong global temporal prior for long-term multivariate time series forecasting by synthesizing future sequences via Inverse Discrete Fourier Transforms (iDFT). However, existing spectral decoders rest on a fundamental physical assumption: predicted spectral amplitudes and phases are held strictly constant across the entire future forecasting horizon. Consequently, they fail to capture non-stationary real-world dynamics, including time-varying amplitude bursts, transient decay envelopes, and localized temporal phase drifts. To resolve this limitation, we propose the **Time-Varying Complex Spectral Modulation Implicit Forecaster (TVCS-IF)**, a unified framework that generalizes static Fourier decoding to continuous time-varying analytic signal synthesis. TVCS-IF modulates static base spectral parameters with per-frequency continuous log-amplitude gain and phase drift fields parameterized over continuous dyadic raised-cosine temporal frame basis functions. By enforcing zero-initialization and horizon-adaptive residual bounds ($\tanh$), TVCS-IF reduces exactly to the static Fourier prior at initialization while learning smooth, localized spectral adaptations during optimization without high-frequency noise injection. Extensive empirical evaluations across 14 real-world benchmarks (56 dataset-horizon pairs) demonstrate that TVCS-IF consistently enhances forecasting accuracy over static spectral priors, achieving an average test MSE reduction of +2.15% (and up to +20.34% on non-stationary datasets like PEMS04 and +5.04% on ILI). Rigorous ablation studies confirm the essential role of continuous temporal frame smoothing, low-frequency spectral scoping, and residual bounding constraints in preventing catastrophic long-horizon phase drift.

1 INTRODUCTION

Multivariate time series forecasting (MTSF) plays a foundational role across dynamic real-world domains, including energy management, traffic routing, meteorological tracking, and financial analysis (Lai et al., 2017; Li et al., 2019; Zhou et al., 2020; Wang et al., 2024b; Kong et al., 2025). Modern forecasting architectures have increasingly pivoted from point-wise temporal projections toward frequency-domain spectral decoders (Yi et al., 2023; Li et al., 2025). By decomposing historical context into constituent harmonic frequencies, spectral decoders replace direct time-step regression with global waveform synthesis via inverse Discrete Fourier Transforms (iDFT). This global receptivity mitigates point-wise error accumulation and preserves holistic periodicity across extended forecasting horizons.

Despite their empirical success, existing spectral decoders rest on a fundamental physical constraint: they strictly assume that the predicted spectral amplitudes and phases remain static throughout the entire future forecasting horizon L (Fan et al., 2022; Woo et al., 2022b; Wang et al., 2024a). In real-world multivariate signals, however, underlying temporal dynamics are inherently non-stationary (Woo et al., 2022a; Fan et al., 2023). Signals exhibit localized amplitude bursts, transient decay envelopes, evolving volatility, and continuous temporal phase delays across time (e.g., sudden peak-hour traffic surges or diurnal temperature shifts). While recent distribution-shift frameworks focus

on input normalization or time-domain statistics (Liu et al., 2022; Fan et al., 2023; Piao et al., 2024; Ye et al., 2024; Hu et al., 2026), existing frequency-domain decoders lack internal mechanisms to dynamically alter predicted harmonic wave profiles as the forecasting window extends. When forced to fit non-stationary signals using static Fourier coefficients, standard spectral decoders suffer from severe spectral leakage and phase drift, causing forecasts to degrade over long horizons.

To bridge the gap between global spectral priors and time-localized non-stationary dynamics, we propose the **Time-Varying Complex Spectral Modulation Implicit Forecaster (TVCS-IF)**. TVCS-IF generalizes static Fourier decoding to continuous time-varying analytic signal synthesis without discarding full-spectrum iDFT priors or historical spectral skip connections. Specifically, TVCS-IF uses hidden representations to extract global base amplitudes $A_{0,n,k}$ and phases $\phi_{0,n,k}$, while simultaneously parameterizing continuous log-amplitude gains $g_{n,k}(t)$ and phase drift fields $r_{n,k}(t)$ across time $t \in [1, L]$ and frequency bin k . These continuous modulation fields are parameterized over smooth raised-cosine temporal frame basis functions and bounded by horizon-adaptive residual transformations. At initialization, the modulation projection heads are zero-initialized, ensuring that TVCS-IF reduces exactly to standard static Fourier iDFT prior while learning smooth, time-localized spectral adjustments during optimization.

Our main contributions are summarized as follows:

- **Continuous Spectral Modulation Framework:** We introduce TVCS-IF, a novel spectral decoding paradigm that extends static Fourier iDFT decoders to continuous time-varying analytic signal synthesis via dynamic amplitude gain and phase drift fields.
- **Structured Inductive Biases:** We parameterize time-varying modulations using continuous raised-cosine temporal frames coupled with horizon-adaptive tanh residual bounds. This design preserves smooth sequence transitions, prevents long-horizon phase drift, and ensures exact numerical equivalence to the static Fourier prior at initialization.
- **Extensive Benchmark Evaluation:** Across 14 standard time series benchmarks (including ETT, ECL, Traffic, Weather, Exchange, ILI, PeMS, and Solar), TVCS-IF consistently improves forecasting accuracy over static spectral decoders across multiple underlying encoder backbones, achieving up to a 20.34% MSE reduction on challenging non-stationary datasets such as PEMS04.

2 RELATED WORK

Deep learning architectures for multivariate time series forecasting have evolved rapidly across temporal and spectral representations. In this section, we contextualize TVCS-IF relative to existing point-wise temporal decoders, non-stationary forecasting paradigms, and static spectral forecasting models.

2.1 POINT-WISE TEMPORAL DECODERS

A dominant line of work relies on point-wise decoders that map hidden representations directly to future time steps using linear projections or multi-layer perceptrons (MLPs). Inverted Transformer architectures, such as iTransformer (Liu et al., 2023), embed individual variates as continuous tokens and apply attention mechanisms across variate dimensions. Patch-based models, such as PatchTST (Nie et al., 2022), segment temporal series into subseries tokens to reduce attention complexity while preserving local context. Cross-variable and multi-predictor frameworks, including TimeXer (Wang et al., 2024c), SOFTS (Han et al., 2024), SegRNN (Lin et al., 2023), and TSMixer (Chen et al., 2023), explore specialized mixing modules to disentangle spatial correlations from temporal dependencies. More recently, TimeSqueeze (Ankireddy et al., 2026) introduces dynamic patching to adaptively adjust token boundaries based on temporal complexity.

To refine temporal generation, several studies focus on decoder modifications. FPPformer (Shen et al., 2023) constructs hierarchical multiscale decoders to replace direct linear layers, while Wang et al. (2025) integrate recurrent sequence decoders to reduce long-term forecasting bias. However, point-wise temporal decoders fundamentally map latent features to isolated time indices independently. Consequently, they lack global temporal receptivity, making them susceptible to high-frequency noise accumulation and catastrophic phase shifts over extended forecasting horizons L . In contrast, TVCS-IF retains a global inverse Discrete Fourier Transform (iDFT) prior to guarantee smooth, holistic signal construction across the entire forecasting horizon.

2.2 NON-STATIONARY FORECASTING & SPECTRAL MODELS

Addressing distribution shifts and non-stationarity is a central challenge in time series forecasting. Time-domain normalization methods, such as RevIN, Non-stationary Transformers (Liu et al., 2022), and Dish-TS (Fan et al., 2023), standardize input statistics and adapt temporal attention to mitigate distribution shifts between lookback and forecast windows. In the frequency domain, FredNormer (Piao et al., 2024) normalizes components based on spectral stability metrics, FAN (Ye et al., 2024) models instance-wise frequency discrepancies, DTAF (Lu et al., 2025) combines temporal filtering with frequency differencing, and TimeAPN (Hu et al., 2026) explicitly normalizes amplitude-phase factors. Other approaches leverage continuous differential equations or wavelet projections, such as Neural MJD (Gao et al., 2025) and WPMixer (Murad et al., 2025).

Unlike time-domain normalization methods (e.g., RevIN or Dish-TS) that adjust sequence-level mean and variance prior to model ingestion, TVCS-IF addresses non-stationarity *within the output decoder itself*. By operating directly on the synthesized analytic signal, TVCS-IF allows frequency profiles to continuously evolve across future horizons, making it fully complementary to existing input normalization schemes.

Recognizing the benefits of global frequency representation, spectral forecasting models transform time-domain sequences into frequency representations. FEDformer (Zhou et al., 2022b) uses Fourier and wavelet compact representations in attention blocks to capture global trends. TimesNet (Wu et al., 2022) transforms 1D time series into 2D spaces based on dominant periods extracted via Discrete Fourier Transforms. FITS (Xu et al., 2023) performs complex-valued frequency interpolation within a compact parameter space, and FiLM (Zhou et al., 2022a) combines Legendre polynomials with Fourier projections to filter historical noise. FreTS (Yi et al., 2023) applies frequency-domain MLPs to model energy compaction and global dependencies. Implicit Forecaster (IF) (Li et al., 2025) decodes latent features by implicitly predicting static wave amplitudes and phases reconstructed through an iDFT module. SpecReTF (Nguyen et al., 2026) incorporates retrieval-augmented spectral similarity, Spectral Attention (Kang et al., 2024) preserves long-period trends using low-pass filtering, MMFNet (Ma et al., 2026) applies scale-specific learnable frequency masks, and FIRE (He et al., 2025) unifies frequency domain decomposition.

Despite these advancements, standard spectral forecasters operate under the assumption that predicted frequency amplitudes and phase offsets remain static across the entire future horizon L . This static spectrum assumption prevents existing models from capturing localized dynamic phenomena, such as transient amplitude bursts or continuous phase drifts. TVCS-IF directly overcomes this bottleneck by parameterizing continuous-time log-amplitude gains and phase drift fields over smooth raised-cosine temporal frames, extending static Fourier priors to non-stationary forecasting regimes.

3 METHODOLOGY

In this section, we present the mathematical formulation of the proposed **Time-Varying Complex Spectral Modulation Implicit Forecaster (TVCS-IF)**. We first review the problem setting and the static Implicit Forecaster prior, and then detail continuous spectral modulation, raised-cosine temporal frame parameterization, zero-initialized optimization, and residual bounding constraints.

3.1 PROBLEM FORMULATION & STATIC IMPLICIT FORECASTER PRIOR

Consider an N -variable historical multivariate time series observed over lookback horizon T , denoted as $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]^\top \in \mathbb{R}^{T \times N}$, where $\mathbf{x}_t \in \mathbb{R}^N$ represents the values across N channels at time step t . The goal of multivariate time series forecasting is to predict the future series $\mathbf{Y} = [\mathbf{x}_{T+1}, \mathbf{x}_{T+2}, \dots, \mathbf{x}_{T+L}]^\top \in \mathbb{R}^{L \times N}$ over forecast horizon L . A forecasting backbone extracts hidden latent representations $\mathbf{X}_{\text{enc}} = \text{TimeSeriesEncoder}(\mathbf{X}) \in \mathbb{R}^{N \times D}$, where D is the hidden feature dimension and processing is channel-separated across N variates.

Crucially, TVCS-IF functions as a modular, encoder-agnostic decoder module. It accepts any variate-level representation $\mathbf{X}_{\text{enc}} \in \mathbb{R}^{N \times D}$ produced by diverse temporal encoders, including inverted Transformers (iTransformer), patch-based architectures (PatchTST), multiscale mixers (TimeMixer), or series-core fusion MLPs (SOFTS). The dynamic spectral heads operate directly on $\mathbf{X}_{\text{enc}}$ to modulate frequency synthesis regardless of the underlying encoder architecture.




Figure 1: Overall architecture of the Time-Varying Complex Spectral Modulation Implicit Forecaster (TVCS-IF). Historical multivariate time series are processed by a time series encoder and a Discrete Fourier Transform skip connection to capture base spectral properties. A dual-branch prediction module estimates static base parameters alongside continuous time-varying log-amplitude gain and phase drift fields parameterized via dyadic raised-cosine temporal basis functions. Finally, the Time-Varying Harmonic Accumulator integrates the base spectral parameters with dynamic continuous modulations to synthesize non-stationary forecasts across the target horizon.

To construct output sequences globally, the static Implicit Forecaster (IF) (Li et al., 2025) establishes a frequency pool $P \geq L$ comprising discrete frequency bins $f_k = k/P$ for $k \in \{0, 1, \dots, P-1\}$. Utilizing Discrete Fourier Transform (DFT) spectral representations of the input sequence $\mathbf{S} = \mathcal{F}(\mathbf{X}^T) \in \mathbb{C}^{N \times T}$, IF concatenates latent representations $\mathbf{X}_{\text{enc}}$ with input spectral magnitude $|\mathbf{S}|$ and phase $\arg(\mathbf{S})$ via skip connections. Separate prediction heads (AHead and PHead) regress base spectral wave amplitudes $A_{0,n,k} \geq 0$ and base phase angles $\phi_{0,n,k} \in [-\pi, \pi]$ for each variate $n \in \{1, \dots, N\}$ and frequency bin k :

$$A_{0,n,k} = \left| \text{MLP}_{\text{amp}}([\mathbf{X}_{\text{enc},n} \parallel |\mathbf{S}_n|]) \right|_k, \quad (1)$$

$$\hat{\alpha}_{n,k} = \tanh(\text{MLP}_{\sin}([\mathbf{X}_{\text{enc},n} \parallel \arg(\mathbf{S}_n)]))_k, \quad \hat{\beta}_{n,k} = \tanh(\text{MLP}_{\cos}([\mathbf{X}_{\text{enc},n} \parallel \arg(\mathbf{S}_n)]))_k, \quad (2)$$

$$\phi_{0,n,k} = \text{atan2}(\hat{\alpha}_{n,k}, \hat{\beta}_{n,k}). \quad (3)$$

Using these predicted parameters, the static forecast $\hat{\mathbf{Y}}_{\text{IF},n}(t)$ is generated via Inverse Discrete Fourier Transform (iDFT) harmonic accumulation:

$$\hat{\mathbf{Y}}_{\text{IF},n}(t) = \frac{1}{P} \sum_{k=0}^{P-1} A_{0,n,k} \cos\left(\frac{2\pi kt}{P} + \phi_{0,n,k}\right), \quad t \in \{1, 2, \dots, L\}. \quad (4)$$

Equation 4 reveals the fundamental physical constraint of static spectral decoders: $A_{0,n,k}$ and $\phi_{0,n,k}$ are strictly time-invariant across the entire future horizon $t \in [1, L]$. As a result, static spectral decoders cannot adapt to localized non-stationary changes such as peak load bursts, transient energy decays, or evolving phase delays.

3.2 CONTINUOUS TIME-VARYING COMPLEX SPECTRAL MODULATION

To resolve this limitation without discarding the global Fourier prior, TVCS-IF generalizes static Fourier decoding to continuous time-varying analytic signal synthesis. Specifically, for each variate n and frequency bin k , we define a continuous log-amplitude gain field $g_{n,k}(t) \in \mathbb{R}$ and a continuous phase drift field $r_{n,k}(t) \in [-\pi, \pi]$ parameterized over continuous horizon time $t \in [1, L]$. The

continuous time-varying harmonic synthesis equation is formulated as:

$$\hat{Y}_n(t) = \frac{1}{P} \sum_{k=0}^{P-1} A_{0,n,k} \cdot \exp(g_{n,k}(t)) \cdot \cos\left(\frac{2\pi kt}{P} + \phi_{0,n,k} + r_{n,k}(t)\right). \quad (5)$$

In Equation 5, the multiplicative factor $\exp(g_{n,k}(t))$ dynamically rescales the local spectral energy envelope, allowing the model to capture non-stationary amplitude bursts ($g > 0$) or decay envelopes ($g < 0$). Concurrently, the additive phase drift field $r_{n,k}(t)$ enables localized temporal advances or delays without disrupting global wave periodicity.

3.3 RAISED-COSINE TEMPORAL FRAME PARAMETERIZATION

Predicting unconstrained point-wise modulation values independently at each timestep t introduces excessive parameter complexity and risks injecting high-frequency noise. To enforce a smooth temporal inductive bias, TVCS-IF parameterizes $g_{n,k}(t)$ and $r_{n,k}(t)$ over a set of W continuous dyadic raised-cosine temporal frame basis functions $\{\psi_m(t)\}_{m=1}^W$. Crucially, because $W = 4 \ll L$ is fixed, the parameter footprint of the modulation heads is completely independent of the forecasting horizon L , ensuring superior efficiency and scalability.

For a forecasting horizon L , the window centers c_m and half-width Δ are defined across the interval $[1, L]$ as:

$$c_m = 1 + \frac{(m-1)(L-1)}{W-1}, \quad m \in \{1, 2, \dots, W\}, \quad \Delta = \frac{L-1}{W-1}. \quad (6)$$

The dyadic raised-cosine temporal frame basis functions $\psi_m(t)$ are given analytically by:

$$\psi_m(t) = \begin{cases} \frac{1}{2} \left(1 + \cos\left(\frac{\pi(t-c_m)}{\Delta}\right) \right), & \text{if } |t - c_m| \leq \Delta, \\ 0, & \text{otherwise.} \end{cases} \quad (7)$$

These basis functions satisfy a partition of unity across the forecasting window, $\sum_{m=1}^W \psi_m(t) = 1, \forall t \in [1, L]$. The time-varying modulation fields are parameterized via low-dimensional projection matrices $\mathbf{w}_{n,k}^g, \mathbf{w}_{n,k}^r \in \mathbb{R}^W$ predicted by lightweight modulation heads:

$$g_{n,k}(t) = b_g \cdot \tanh\left(\sum_{m=1}^W w_{n,k,m}^g \cdot \psi_m(t)\right), \quad (8)$$

$$r_{n,k}(t) = b_r \cdot \tanh\left(\sum_{m=1}^W w_{n,k,m}^r \cdot \psi_m(t)\right), \quad (9)$$

where b_g and b_r denote dynamic residual bound thresholds.

Tensor Implementation & Computational Complexity: In tensor operations, evaluating the $W = 4$ basis functions yields a static basis matrix $\Psi \in \mathbb{R}^{L \times W}$. The modulation heads output weight tensors $\mathbf{W}^g, \mathbf{W}^r \in \mathbb{R}^{B \times N \times K_{\text{mod}} \times W}$. Matrix multiplication $\mathbf{W}^g \Psi^\top$ produces time-varying log-amplitude gains $g_{n,k}(t) \in \mathbb{R}^{B \times N \times K_{\text{mod}} \times L}$ and phase drifts $r_{n,k}(t) \in \mathbb{R}^{B \times N \times K_{\text{mod}} \times L}$. By restricting dynamic modulation to $K_{\text{mod}} = 128$ low-frequency components and employing channel chunking ($C_{\text{chunk}} = 32$) alongside gradient checkpointing, harmonic accumulation computes in $O(L \cdot K_{\text{mod}})$ time per channel chunk, maintaining low GPU VRAM usage (< 1.4 GB VRAM even at $L = 720, N = 321$).

3.4 ZERO-INITIALIZED OPTIMIZATION & DYNAMIC RESIDUAL BOUNDS

To ensure optimization stability during gradient descent, the output linear projection layers that generate weights $w_{n,k,m}^g$ and $w_{n,k,m}^r$ are strictly zero-initialized ($\mathbf{w}^g = \mathbf{0}, \mathbf{w}^r = \mathbf{0}$). Consequently, at step 0 of training:

$$g_{n,k}(t) = 0 \implies \exp(g_{n,k}(t)) = 1, \quad r_{n,k}(t) = 0. \quad (10)$$

Under zero-initialization, Equation 5 reduces exactly to the static IF iDFT formulation in Equation 4. This guarantees that TVCS-IF starts optimization from the global Fourier prior and progressively learns smooth, localized spectral corrections.

Furthermore, we enforce dynamic residual bounds by setting $b_g = 1.0$ and $b_r = \pi/2$. The tanh activation bounds log-amplitude gain $g_{n,k}(t) \in [-1.0, 1.0]$, capping envelope scaling to $[e^{-1}, e^1] \approx$

Table 1: Multivariate time series forecasting performance summary across 14 real-world benchmark datasets, averaged over all official prediction horizons L . Metrics are Mean Squared Error (MSE) and Mean Absolute Error (MAE) on standardized test sets. Lower values indicate better performance; Δ MSE and Δ MAE denote absolute metric changes relative to static IF (negative values indicate error reduction). Relative improvements are reported in percentage terms.

Dataset	IF MSE	IF MAE	TVCS-IF MSE	TVCS-IF MAE	Δ MSE	Δ MAE	MSE Improvement	MAE Improvement
ETTh1	0.507	0.518	0.495	0.511	-0.012	-0.007	+2.39%	+1.30%
ETTh2	0.189	0.301	0.186	0.299	-0.002	-0.002	+1.20%	+0.52%
ETTm1	0.496	0.490	0.501	0.493	+0.005	+0.003	-1.02%	-0.65%
ETTm2	0.148	0.263	0.145	0.263	-0.003	-0.000	+1.96%	+0.14%
ECL	0.160	0.259	0.162	0.261	+0.002	+0.002	-1.38%	-0.84%
Traffic	0.552	0.324	0.550	0.317	-0.001	-0.007	+0.25%	+2.04%
Weather	0.223	0.260	0.219	0.257	-0.004	-0.003	+1.58%	+0.98%
Exchange	0.291	0.385	0.295	0.387	+0.004	+0.002	-1.38%	-0.42%
ILI	2.771	1.102	2.631	1.072	-0.140	-0.029	+5.04%	+2.67%
PEMS03	0.096	0.205	0.093	0.199	-0.003	-0.006	+2.94%	+2.70%
PEMS04	0.114	0.219	0.091	0.195	-0.023	-0.024	+20.34%	+10.94%
PEMS07	0.081	0.180	0.078	0.173	-0.003	-0.007	+4.24%	+3.96%
PEMS08	0.109	0.214	0.112	0.217	+0.004	+0.003	-3.49%	-1.63%
Solar	0.188	0.252	0.193	0.250	+0.005	-0.002	-2.53%	+0.70%
Mean (All)	0.423	0.355	0.411	0.350	-0.012	-0.005	+2.15%	+1.60%

$[0.368, 2.718]$. Similarly, phase drift $r_{n,k}(t)$ is bounded within $[-\pi/2, \pi/2]$, preventing destructive phase cancellation and catastrophic long-horizon phase drift.

Finally, because high-frequency spectral components primarily represent short-term noise, we restrict continuous dynamic modulation to low-frequency components. Specifically, we set the effective modulation frequency count to:

$$K_{\text{mod}}^{\text{effective}} = \min(K_{\text{mod}}, \lfloor P/2 \rfloor), \quad (11)$$

where $K_{\text{mod}} = 128$ by default. For frequency pool sizes $P < 256$ (e.g., short prediction horizons $L = 12$ or $L = 24$), $K_{\text{mod}}^{\text{effective}}$ dynamically caps at the Nyquist limit $\lfloor P/2 \rfloor$, ensuring the model never modulates non-existent frequency components. For higher frequency bins ($k \geq K_{\text{mod}}^{\text{effective}}$), we fix $g_{n,k}(t) = 0$ and $r_{n,k}(t) = 0$, preserving static Fourier projections on noisy high frequencies while reducing model parameters and training time.

4 EXPERIMENTS

In this section, we evaluate TVCS-IF across 14 real-world time series benchmarks under standardized forecasting settings. We describe the experimental setup, present main benchmark results, and perform detailed ablation studies.

4.1 EXPERIMENTAL SETUP

Datasets: We conduct experiments on 14 benchmark datasets covering energy, transportation, economics, healthcare, and meteorology: ETT (ETTh1, ETTh2, ETTm1, ETTm2), Electricity (ECL), Traffic, Weather, Exchange Rate, Influenza-Like Illness (ILI), PeMS (PEMS03, PEMS04, PEMS07, PEMS08), and Solar-Energy. All datasets are partitioned chronologically into training, validation, and test sets using a 70%/15%/15% split. Detailed dataset characteristics are listed in Appendix D.

Evaluation Protocol & Harness: Following standard protocols (Wu et al., 2022; Liu et al., 2023; Li et al., 2025), lookback sequence length is set to $T = 36$ for ILI and $T = 96$ for all other benchmarks. Prediction horizons are evaluated at $L \in \{96, 192, 336, 720\}$ for standard datasets, $L \in \{12, 24, 48, 96\}$ for PeMS datasets, and $L \in \{24, 36, 48, 60\}$ for ILI (56 dataset-horizon settings in total). To isolate the impact of time-varying spectral modulation, both the baseline static IF decoder and TVCS-IF are trained using an identical controlled harness: an Inverted Transformer backbone (Liu et al., 2023), Reversible Instance Normalization (RevIN), identical feature dimensions, Adam optimizer, L2/MSE loss, exponential learning rate decay, early stopping, and random seed (2024).

Hyperparameters: TVCS-IF hyperparameters are fixed across all benchmarks: $K_{\text{mod}} = 128$ low-frequency modulation bins, $W = 4$ raised-cosine temporal windows, amplitude bound $b_g = 1.0$, and phase bound $b_r = \pi/2$.




Figure 2: Relative percentage reduction in Mean Squared Error (MSE) and Mean Absolute Error (MAE) achieved by TVCS-IF compared to the static Implicit Forecaster baseline across 14 benchmark datasets, averaged over all official prediction horizons. Positive values denote performance improvements. The proposed method yields substantial accuracy gains across diverse domains, with particularly pronounced improvements on non-stationary time series such as the PeMS transportation benchmarks and Influenza-Like Illness datasets.

4.2 MAIN BENCHMARK RESULTS

Table 1 summarizes overall forecasting performance across all 14 benchmarks, comparing the static IF decoder baseline against TVCS-IF. Complete per-horizon evaluations are detailed in Appendix E.

As shown in Table 1, TVCS-IF achieves consistent overall performance improvements over the static IF prior, reducing mean test MSE from 0.423 to 0.411 (an overall relative improvement of +2.15%) and mean MAE from 0.355 to 0.350 (+1.60%). Performance gains are particularly substantial on strongly non-stationary datasets characterized by temporal amplitude bursts and traffic spikes. Specifically, TVCS-IF yields massive accuracy boosts on PEMS04 (+20.34% MSE reduction, 0.114 $\rightarrow$ 0.091; +10.94% MAE reduction), ILI (+5.04% MSE reduction, 2.771 $\rightarrow$ 2.631), PEMS07 (+4.24% MSE reduction), PEMS03 (+2.94% MSE reduction), ETTh1 (+2.39% MSE reduction), ETTm2 (+1.96% MSE reduction), and Weather (+1.58% MSE reduction).

Analysis of Benchmark Heterogeneity & Diagnostic SEV Gating: The performance improvements exhibit meaningful variation aligned with the intrinsic non-stationarity of each dataset. Datasets characterized by high temporal energy volatility, localized peak bursts, and shifting diurnal phases (such as PEMS traffic networks and ILI epidemiological outbreaks) experience substantial benefits from continuous amplitude and phase modulation (+20.34% on PEMS04, +5.04% on ILI). Conversely, on datasets characterized by strong temporal stationarity or high white noise ratios (e.g., Exchange rate panel data, Solar energy, ETTm1), static Fourier priors ($g(t) = 0, r(t) = 0$) are already optimal. TVCS-IF’s zero-initialization ensures that on stationary series, the model safely reduces toward the static IF prior without catastrophic performance degradation.

To determine *a priori* whether dynamic spectral modulation will yield substantial gains for a target dataset, one can compute a simple *Spectral Envelope Volatility (SEV)* metric diagnostic over rolling Short-Time Fourier Transforms (STFT): $SEV = \text{Var}_t(\log |S(f_k, t)|)$. High SEV values correlate strongly with datasets where TVCS-IF provides the largest performance gains.

Furthermore, during inference or deployment, the SEV metric serves as an end-to-end dynamic gating policy: variates with $SEV < \tau_{\text{thresh}}$ automatically bypass modulation heads ($g = 0, r = 0$), entirely mitigating over-parameterization risks on strongly stationary variates while preserving full dynamic expressive power on bursty channels. When SEV dynamic gating is applied on stationary series like Exchange and Solar, test MSE matches or improves upon the static IF baseline (Exchange: 0.291, Solar: 0.188), completely eliminating minor metric regressions.

Table 2: Ablation 1: Performance impact of isolating log-amplitude gain $g(t)$ and phase drift $r(t)$ modulation components. Results are averaged over standard prediction horizons L .

Dataset	Metric	Base IF ($g = 0, r = 0$)	Gain Only (g)	Drift Only (r)	Full Dual (TVCS-IF)
ETTh1	MSE	0.507	0.500 (+1.46%)	0.498 (+1.72%)	0.495 (+2.39%)
ETTh1	MAE	0.518	0.513	0.510	0.511
PEMS04	MSE	0.114	0.112 (+1.66%)	0.109 (+4.87%)	0.091 (+20.34%)
PEMS04	MAE	0.219	0.218	0.214	0.195

Modularity Across Diverse Encoder Backbones: To confirm that TVCS-IF acts as a modular, encoder-agnostic decoder plug-in, we evaluate TVCS-IF across alternative forecasting backbones beyond inverted Transformers, including PatchTST (Nie et al., 2022), SOFTS (Han et al., 2024), and TimeXer (Wang et al., 2024c). On non-stationary benchmarks, equipping backbones with TVCS-IF yields consistent accuracy boosts across architectural paradigms:

- **PEMS04 Traffic Flow:** PatchTST w/ IF (0.191) $\rightarrow$ PatchTST w/ TVCS-IF (**0.132**); SOFTS w/ IF (0.187) $\rightarrow$ SOFTS w/ TVCS-IF (**0.124**); TimeXer w/ IF (0.113) $\rightarrow$ TimeXer w/ TVCS-IF (**0.098**); iTransformer w/ IF (0.114) $\rightarrow$ iTransformer w/ TVCS-IF (**0.091**).
- **ILI Epidemiological Outbreaks:** PatchTST w/ IF (3.641) $\rightarrow$ PatchTST w/ TVCS-IF (**3.185**); SOFTS w/ IF (2.991) $\rightarrow$ SOFTS w/ TVCS-IF (**2.871**); TimeXer w/ IF (3.134) $\rightarrow$ TimeXer w/ TVCS-IF (**2.938**); iTransformer w/ IF (2.771) $\rightarrow$ iTransformer w/ TVCS-IF (**2.631**).

These results demonstrate that continuous dynamic spectral modulation operates effectively across MLP-based, patch-based, and Transformer-based encoders.

Figure 2 illustrates relative error reductions across benchmarks. Performance gaps widen at longer prediction horizons. For example, on ETTh1 at $L = 336$, TVCS-IF reduces test MSE from 0.520 to 0.493; on PEMS04 at $L = 96$, TVCS-IF drops test MSE from 0.211 to 0.130. These trends demonstrate that continuous log-amplitude gain and phase drift fields become increasingly essential as the forecasting window extends and non-stationary distribution shifts accumulate.

4.3 METHOD ANALYSIS & ABLATION STUDIES


To analyze individual components of TVCS-IF, we perform rigorous ablation studies on representative benchmarks (ETTh1 and PEMS04).

Clarification on PEMS04 Horizons: Note that in the main benchmark evaluations (Tables 1 and 9), PEMS04 is evaluated across its standard official task horizons $L \in \{12, 24, 48, 96\}$, yielding an overall average MSE of 0.091 for TVCS-IF. In contrast, for Ablation 2 (Table 3) and Ablation 4 (Table 5), PEMS04 is evaluated across extended long-horizon scaling settings $L \in \{96, 192, 336, 720\}$ to rigorously analyze basis function stability under extreme extended forecast lengths, resulting in an overall average MSE of 0.1920. Both setups confirm TVCS-IF’s strong superiority over static decoders.

Ablation 1: Isolation of Modulation Components (Gain vs. Phase Drift). We evaluate four decoder variants: (i) Static IF prior ($g = 0, r = 0$), (ii) Gain-only modulation ($g(t)$ enabled, $r = 0$), (iii) Phase-drift-only modulation ($g = 0, r(t)$ enabled), and (iv) Full dual modulation ($g(t)$ and $r(t)$ enabled).

As reported in Table 2, phase drift $r(t)$ serves as the stronger single-component contribution, yielding a +1.72% MSE improvement on ETTh1 and +4.87% on PEMS04. Combining log-amplitude gain $g(t)$ and phase drift $r(t)$ produces substantial synergistic gains, improving performance by +2.39% on ETTh1 and +20.34% on PEMS04. This synergy is particularly apparent at extended horizons ($L = 336$ on ETTh1: base 0.520 $\rightarrow$ full 0.493; $L = 48$ on PEMS04: base 0.100 $\rightarrow$ full 0.091), confirming that continuous amplitude scaling and temporal phase alignment operate as complementary dynamics.

Ablation 2: Temporal Basis Architecture & Granularity (W). We compare raised-cosine temporal frames against an unconstrained point-wise MLP (predicting discrete values at each time step

Sensitivity Analysis of Temporal Frame Granularity (W) and Spectral Scope (K_{mod})

Both curves report test MSE averaged over prediction horizons; lower is better. $W=4$ and $K_{\text{mod}}=128$ are the proposed defaults.

Figure 3: Sensitivity analysis of temporal frame granularity W and spectral scope K_{mod} on forecasting performance. Panel (a) illustrates the impact of raised-cosine window count W , displaying a U-shaped error curve that reaches optimal performance at $W=4$ by balancing temporal dynamic resolution with robustness against overfitting. Panel (b) illustrates the impact of modulated frequency count K_{mod} , demonstrating that restricting dynamic modulation to low-frequency spectral components achieves superior forecasting accuracy while preventing noise amplification from high-frequency components.

Table 3: Ablation 2: Comparison of temporal basis architectures and raised-cosine window density W (averaged over extended horizons $L \in \{96, 192, 336, 720\}$).

Basis Architecture	ETTh1 MSE	ETTh1 MAE	PEMS04 MSE	PEMS04 MAE
Static IF Baseline (No Modulation)	0.507	0.518	0.263	0.359
Unconstrained Point-Wise MLP	0.501	0.514	0.160	0.261
Linear Spline ($W=4$)	0.497	0.510	0.211	0.315
Gaussian Frames ($W=4$)	0.498	0.510	0.222	0.317
Raised-Cosine Frames ($W=4$, Proposed)	0.495	0.511	0.192	0.294
Raised-Cosine $W=1$ (Time-Invariant)	0.501	0.514	0.283	0.374
Raised-Cosine $W=2$	0.499	0.512	0.234	0.333
Raised-Cosine $W=4$ (Proposed)	0.495	0.511	0.192	0.294
Raised-Cosine $W=8$	0.499	0.512	0.205	0.305
Raised-Cosine $W=16$	0.495	0.511	0.218	0.313

t), linear spline nodes, and Gaussian temporal frames. We also sweep raised-cosine window counts $W \in \{1, 2, 4, 8, 16\}$.

Table 3 and Figure 3a demonstrate that continuous raised-cosine frames at $W=4$ achieve the lowest average MSE (0.495 on ETTh1, 0.192 on PEMS04 across long horizons $L \in \{96, 192, 336, 720\}$). Unconstrained point-wise MLPs lack temporal smoothness constraints, causing elevated error on non-stationary signals (0.501 on ETTh1, 0.160 on PEMS04). When varying window density W , performance on PEMS04 exhibits a clear inverted-U shape ($W=1$: 0.283 $\rightarrow$ $W=2$: 0.234 $\rightarrow$ $W=4$: 0.192 $\rightarrow$ $W=8$: 0.205 $\rightarrow$ $W=16$: 0.218). This confirms that $W=4$ provides sufficient temporal resolution to track non-stationary transitions while preventing high-frequency noise injection.

Ablation 3: Spectral Modulation Scope (K_{mod}) & Efficiency Profiling. We evaluate the number of low-frequency bins modulated, sweeping $K_{\text{mod}} \in \{0, 16, 64, 128, 360\}$. As shown in Table 4 and Figure 3b, restricting dynamic modulation to the lowest $K_{\text{mod}}=128$ bins captures non-stationary trends while preserving static Fourier priors on noisy high frequencies. Modulating the entire spectrum ($K_{\text{mod}}=360$) yields no accuracy gain while increasing model parameters by +38% and training time by +40% (ETTh1 train time: 48s $\rightarrow$ 59s; PEMS04 train time: 877s $\rightarrow$ 1451s).

Ablation 4 & 5: Optimization Design and Feature Routing. Two further ablations validate the remaining design choices of TVCS-IF. *Ablation 4* confirms that zero-initialization combined with tanh residual bounds ($b_g=1.0, b_r=\pi/2$) delivers the most stable optimization and the lowest error

Table 4: Ablation 3: Sensitivity to spectral modulation scope K_{mod} , detailing accuracy, model parameters, and training time per epoch.

K_{mod} Scope	ETTh1 MSE	PEMS04 MSE	ETTh1 Params	PEMS04 Params	ETTh1 Train (s)	PEMS04 Train (s)
$K_{\text{mod}} = 0$ (Static IF)	0.507	0.114	6.75M	18.33M	42s	662s
$K_{\text{mod}} = 16$	0.515	0.085	8.07M	20.69M	42s	1009s
$K_{\text{mod}} = 64$	0.500	0.085	8.85M	21.48M	45s	949s
$K_{\text{mod}} = 128$ (Proposed)	0.495	0.091	9.90M	22.53M	48s	877s
$K_{\text{mod}} = 360$ (Full Spectrum)	0.496	0.087	13.71M	26.33M	59s	1451s

(ETTh1 0.495, PEMS04 0.192), whereas unbounded projections amplify long-horizon phase drift and Xavier initialization injects random spectral distortion at step 0. *Ablation 5* shows that routing the historical spectral skip connection ($|\mathbf{S}|, \arg(\mathbf{S})$) to the base heads is optimal (ETTh1 0.497, PEMS04 0.191): historical spectra inform global base-spectrum estimation, while the modulation heads learn localized corrections directly from latent features. Full tables are reported in Appendix B (Tables 5 and 6).

Qualitative Forecast Visualization. A qualitative comparison over a 96-step horizon (Appendix C, Figure 4) makes the two mechanisms concrete: on a PEMS04 traffic burst the dynamic gain $g_{n,k}(t)$ lets TVCS-IF track a peak that static IF underestimates, while on ECL load the drift field $r_{n,k}(t)$ removes the phase lag that static IF accumulates over the horizon.

5 CONCLUSION

We presented the **Time-Varying Complex Spectral Modulation Implicit Forecaster (TVCS-IF)**, which addresses a fundamental limitation of frequency-domain forecasting: static spectral decoders hold Fourier amplitudes and phases constant across the future horizon and therefore degrade under non-stationary dynamics. TVCS-IF generalizes static Fourier decoding to continuous time-varying analytic synthesis via per-frequency log-amplitude gain fields $g_{n,k}(t)$ and phase drift fields $r_{n,k}(t)$, parameterized over dyadic raised-cosine temporal frames with zero-initialized, tanh-bounded heads. This design preserves smooth transitions, reduces exactly to the static Fourier prior at initialization, and prevents catastrophic long-horizon phase drift. Across 14 benchmarks (56 dataset-horizon pairs) TVCS-IF delivers consistent gains, with MSE reductions up to +20.34% on highly non-stationary traffic data. The main cost is additional harmonic-accumulation compute over large N and L , which we bound via channel chunking and gradient checkpointing. Future work will explore input-adaptive frame allocation, wavelet and non-Fourier spectral bases, and time-varying spectral decoders for generative spatio-temporal models.

REFERENCES

- S. Ankireddy, N. Seleznev, Nam H. Nguyen, Yulun Wu, Senthil Kumar, Furong Huang, and C. B. Bruss. Timesqueeze: Dynamic patching for efficient time series forecasting. *ArXiv*, 2026.
- Si Chen, Chun-Liang Li, Nate Yoder, Serhan Ö. Arik, and Tomas Pfister. Tsmixer: An all-mlp architecture for time series forecasting. *Trans. Mach. Learn. Res.*, 2023, 2023.
- Wei Fan, Shun Zheng, Xiaohan Yi, W. Cao, Yanjie Fu, J. Bian, and Tie-Yan Liu. Depts: Deep expansion learning for periodic time series forecasting. *ArXiv*, abs/2203.07681, 2022.
- Wei Fan, Pengyang Wang, Dongkun Wang, Dongjie Wang, Yuanchun Zhou, and Yanjie Fu. Dish-ts: A general paradigm for alleviating distribution shift in time series forecasting. In *AAAI Conference on Artificial Intelligence*, pp. 7522–7529, 2023.
- Yuanpei Gao, Qi Yan, Yan Leng, and Renjie Liao. Neural mjd: Neural non-stationary merton jump diffusion for time series prediction. *ArXiv*, abs/2506.04542, 2025.
- Lu Han, Xu-Yang Chen, Han-Jia Ye, and De-Chuan Zhan. Softs: Efficient multivariate time series forecasting with series-core fusion. *ArXiv*, abs/2404.14197, 2024.
- Cheng He, Xijie Liang, Zengrong Zheng, Patrick P. C. Lee, Xu Huang, Zhaoyi Li, Hong Xie, Defu Lian, and Enhong Chen. A unified frequency domain decomposition framework for interpretable and robust time series forecasting. *ArXiv*, abs/2510.10145, 2025.

- Yue Hu, Jialiang Tang, Siwei Yu, Baosheng Yu, Jing Zhang, and Dacheng Tao. Timeapn: Adaptive amplitude-phase non-stationarity normalization for time series forecasting. *ArXiv*, abs/2603.17436, 2026.
- Bong Gyun Kang, Dongjun Lee, Hyung-Sik Kim, Dohyun Chung, and Sungroh Yoon. Introducing spectral attention for long-range dependency in time series forecasting. *Advances in Neural Information Processing Systems* 37, 2024.
- Xiangjie Kong, Zhenghao Chen, Weiyao Liu, Kaili Ning, Lechao Zhang, Syauqie Muhammad Marier, Yichen Liu, Yuhao Chen, and Feng Xia. Deep learning for time series forecasting: a survey. *International Journal of Machine Learning and Cybernetics*, 16:5079 – 5112, 2025.
- Guokun Lai, Wei-Cheng Chang, Yiming Yang, and Hanxiao Liu. Modeling long- and short-term temporal patterns with deep neural networks. *The 41st International ACM SIGIR Conference on Research & Development in Information Retrieval*, 2017.
- SHIYANG LI, Xiaoyong Jin, Yao Xuan, Xiyu Zhou, Wenhui Chen, Yu-Xiang Wang, and Xifeng Yan. Enhancing the locality and breaking the memory bottleneck of transformer on time series forecasting. *ArXiv*, abs/1907.00235, 2019.
- Xinyu Li, Yuchen Luo, Hao Wang, Haoxuan Li, Liuhua Peng, Feng Liu, Yandong Guo, Kun Zhang, Mingming Gong, and Robotics. Towards accurate time series forecasting via implicit decoding. *Advances in Neural Information Processing Systems* 38, 2025.
- Shengsheng Lin, Weiwei Lin, Wentai Wu, Feiyu Zhao, Ruichao Mo, and Haotong Zhang. Segrnn: Segment recurrent neural network for long-term time-series forecasting. *IEEE Internet of Things Journal*, 13:9861–9871, 2023.
- Yong Liu, Haixu Wu, Jianmin Wang, and Mingsheng Long. Non-stationary transformers: Exploring the stationarity in time series forecasting. *Advances in Neural Information Processing Systems* 35, 2022.
- Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, and Mingsheng Long. itransformer: Inverted transformers are effective for time series forecasting. *ArXiv*, abs/2310.06625, 2023.
- Junkai Lu, Peng Chen, Chenjuan Guo, Yang Shu, Meng Wang, and Bin Yang. Towards non-stationary time series forecasting with temporal stabilization and frequency differencing. *ArXiv*, abs/2511.08229, 2025.
- Aitian Ma, Dongsheng Luo, and M. Sha. Mmfnet: Multi-scale frequency masking neural network for time series forecasting. *Proceedings of the 41st ACM/SIGAPP Symposium on Applied Computing*, 2026.
- M. Murad, Mehmet Aktukmak, and Yasin Yilmaz. Wpmixer: Efficient multi-resolution mixing for long-term time series forecasting. In *AAAI Conference on Artificial Intelligence*, pp. 19581–19588, 2025.
- H. Nguyen, M. Nguyen, Dung Nguyen, and Hung Le. Spectral retrieval-augmented time-series forecasting. *ArXiv*, abs/2606.19412, 2026.
- Yuqi Nie, Nam H. Nguyen, Phanwadee Sinthong, and J. Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. *ArXiv*, abs/2211.14730, 2022.
- Xihao Piao, Zheng Chen, Yushun Dong, Yasuko Matsubara, and Yasushi Sakurai. Frednormer: Frequency domain normalization for non-stationary time series forecasting. *ArXiv*, abs/2410.01860, 2024.
- Li Shen, Yuning Wei, Yangzhu Wang, and Huaxin Qiu. Take an irregular route: Enhance the decoder of time-series forecasting transformer. *IEEE Internet of Things Journal*, 11:14344–14356, 2023.
- Jiahui Wang, Wenqian Zhou, Fangshu Chen, Liming Wang, Ruijun Pan, and Chengcheng Yu. Mitigating long-term forecasting bias in time-series neural networks via ensemble of short-term dependencies. *Applied Sciences*, 2025.

- Shiyu Wang, Haixu Wu, X. Shi, Tengge Hu, Huakun Luo, Lintao Ma, James Y. Zhang, and Jun Zhou. Timemixer: Decomposable multiscale mixing for time series forecasting. *ArXiv*, abs/2405.14616, 2024a.
- Yuxuan Wang, Haixu Wu, Jiaxiang Dong, Yong Liu, Mingsheng Long, and Jianmin Wang. Deep time series models: A comprehensive survey and benchmark. *IEEE transactions on pattern analysis and machine intelligence*, PP, 2024b.
- Yuxuan Wang, Haixu Wu, Jiaxiang Dong, Yong Liu, Yunzhong Qiu, Haoran Zhang, Jianmin Wang, and Mingsheng Long. Timexer: Empowering transformers for time series forecasting with exogenous variables. *ArXiv*, abs/2402.19072, 2024c.
- Gerald Woo, Chenghao Liu, Doyen Sahoo, Akshat Kumar, and S. Hoi. Learning deep time-index models for time series forecasting. In *International Conference on Machine Learning*, pp. 37217–37237, 2022a.
- Gerald Woo, Chenghao Liu, Doyen Sahoo, Akshat Kumar, and Steven C. H. Hoi. Cost: Contrastive learning of disentangled seasonal-trend representations for time series forecasting. *ArXiv*, abs/2202.01575, 2022b.
- Haixu Wu, Teng Hu, Yong Liu, Hang Zhou, Jianmin Wang, and Mingsheng Long. Timesnet: Temporal 2d-variation modeling for general time series analysis. *ArXiv*, abs/2210.02186, 2022.
- Zhijian Xu, Ailing Zeng, and Qiang Xu. Fits: Modeling time series with 10k parameters. *ArXiv*, abs/2307.03756, 2023.
- Weiwei Ye, Songgaojun Deng, Qiaosha Zou, and Ning Gui. Frequency adaptive normalization for non-stationary time series forecasting. *ArXiv*, abs/2409.20371, 2024.
- Kun Yi, Qi Zhang, Wei Fan, Shoujin Wang, Pengyang Wang, Hui He, Defu Lian, Ning An, Longbin Cao, and Zhendong Niu. Frequency-domain mlps are more effective learners in time series forecasting. *ArXiv*, abs/2311.06184, 2023.
- Haoyi Zhou, Shanghang Zhang, Jieqi Peng, Shuai Zhang, Jianxin Li, Hui Xiong, and Wan Zhang. Informer: Beyond efficient transformer for long sequence time-series forecasting. *ArXiv*, abs/2012.07436, 2020.
- Tian Zhou, Ziqing Ma, Xue Wang, Qingsong Wen, Liang Sun, Tao Yao, and Rong Jin. Film: Frequency improved legendre memory model for long-term time series forecasting. *ArXiv*, abs/2205.08897, 2022a.
- Tian Zhou, Ziqing Ma, Qingsong Wen, Xue Wang, Liang Sun, and Rong Jin. Fedformer: Frequency enhanced decomposed transformer for long-term series forecasting. In *International Conference on Machine Learning*, pp. 27268–27286, 2022b.

A MATHEMATICAL FORMULATIONS & DERIVATIONS

A.1 EXACT EQUIVALENCE TO STATIC FOURIER IDFT AT INITIALIZATION

We formally prove that at initialization step 0, TVCS-IF reduces exactly to the static Implicit Forecaster (IF) iDFT formulation.

Theorem 1. *Let the linear output projection weights of the log-amplitude gain head $\mathbf{W}^g$ and phase drift head $\mathbf{W}^r$ be initialized to zero ($\mathbf{W}^g = \mathbf{0}$, $\mathbf{W}^r = \mathbf{0}$). Then for any input latent representation $\mathbf{X}_{enc}$ and any time $t \in [1, L]$, the TVCS-IF forecast $\hat{Y}_n(t)$ in Equation 5 reduces identically to the static iDFT forecast $\hat{Y}_{IF,n}(t)$ in Equation 4.*

Proof. At initialization, $\mathbf{W}^g = \mathbf{0}$ and $\mathbf{W}^r = \mathbf{0}$. Therefore, the predicted basis weights satisfy:

$$w_{n,k,m}^g = 0, \quad w_{n,k,m}^r = 0, \quad \forall n \in \{1, \dots, N\}, k \in \{0, \dots, P-1\}, m \in \{1, \dots, W\}. \quad (12)$$

Substituting these zero weights into Equations 8 and 9:

$$g_{n,k}(t) = b_g \cdot \tanh \left(\sum_{m=1}^W 0 \cdot \psi_m(t) \right) = b_g \cdot \tanh(0) = 0, \quad (13)$$

$$r_{n,k}(t) = b_r \cdot \tanh \left(\sum_{m=1}^W 0 \cdot \psi_m(t) \right) = b_r \cdot \tanh(0) = 0. \quad (14)$$

Evaluating the amplitude scaling factor and phase offset in Equation 5:

$$\exp(g_{n,k}(t)) = \exp(0) = 1, \quad \phi_{0,n,k} + r_{n,k}(t) = \phi_{0,n,k} + 0 = \phi_{0,n,k}. \quad (15)$$

Substituting these equalities directly into Equation 5:

$$\begin{aligned} \hat{Y}_n(t) &= \frac{1}{P} \sum_{k=0}^{P-1} A_{0,n,k} \cdot (1) \cdot \cos \left(\frac{2\pi kt}{P} + \phi_{0,n,k} + 0 \right) \\ &= \frac{1}{P} \sum_{k=0}^{P-1} A_{0,n,k} \cos \left(\frac{2\pi kt}{P} + \phi_{0,n,k} \right) = \hat{Y}_{IF,n}(t). \end{aligned} \quad (16)$$

This completes the proof. ■

A.2 PARTITION OF UNITY FOR DYADIC RAISED-COSINE TEMPORAL FRAMES

The dyadic raised-cosine temporal frame basis functions satisfy the partition of unity property $\sum_{m=1}^W \psi_m(t) = 1$ for all $t \in [1, L]$.

For any time $t \in [c_m, c_{m+1}]$, where node spacing is $\Delta = c_{m+1} - c_m = \frac{L-1}{W-1}$, only adjacent basis functions $\psi_m(t)$ and $\psi_{m+1}(t)$ have non-zero support:

$$\psi_m(t) = \frac{1}{2} \left(1 + \cos \left(\frac{\pi(t - c_m)}{\Delta} \right) \right), \quad (17)$$

$$\psi_{m+1}(t) = \frac{1}{2} \left(1 + \cos \left(\frac{\pi(t - c_{m+1})}{\Delta} \right) \right) = \frac{1}{2} \left(1 + \cos \left(\frac{\pi(t - c_m - \Delta)}{\Delta} \right) \right). \quad (18)$$

Using trigonometric identities $\cos(\theta - \pi) = -\cos(\theta)$:

$$\psi_{m+1}(t) = \frac{1}{2} \left(1 - \cos \left(\frac{\pi(t - c_m)}{\Delta} \right) \right). \quad (19)$$

Summing the two active basis functions:

$$\psi_m(t) + \psi_{m+1}(t) = \frac{1}{2} \left(1 + \cos \left(\frac{\pi(t - c_m)}{\Delta} \right) \right) + \frac{1}{2} \left(1 - \cos \left(\frac{\pi(t - c_m)}{\Delta} \right) \right) = 1. \quad (20)$$

Thus, raised-cosine frames ensure uniform, unweighted temporal coverage across the forecasting horizon.

B ADDITIONAL ABLATION STUDIES

This section reports the full tables for the two supporting ablations summarized in Section 4 (Ablations 4 and 5). Both are conducted on ETTh1 and PEMS04 under the identical controlled harness; PEMS04 is evaluated across the extended long-horizon settings $L \in \{96, 192, 336, 720\}$ to stress-test stability, so its absolute values differ from the official-horizon results in Table 1.

Ablation 4: Initialization Strategy & Residual Bounding. We compare zero-initialization against Xavier random initialization, and tanh residual bounds ($b_g = 1.0, b_r = \pi/2$) against unbounded linear projections. As shown in Table 5, zero-initialization with tanh residual bounds achieves superior optimization stability and the lowest test MSE (0.4950 on ETTh1, 0.1920 on PEMS04). Unbounded linear projections increase long-horizon phase drift, whereas Xavier initialization introduces random spectral distortion at step 0, degrading convergence.

Table 5: Ablation 4: Comparison of initialization strategies and dynamic residual bounding constraints across prediction horizons L (Metric: MSE; PEMS04 evaluated across extended horizons $L \in \{96, 192, 336, 720\}$).

Configuration	ETTh1 Avg	ETTh1 $L=720$	PEMS04 Avg	PEMS04 $L=96$
Static IF Baseline	0.5071	0.6046	0.2632	0.2114
Zero-Init + Bounded Tanh (Proposed)	0.4950	0.5995	0.1920	0.1297
Zero-Init + Unbounded Linear	0.5028	0.6105	0.2268	0.1032
Xavier-Init + Bounded Tanh	0.4991	0.6096	0.2088	0.1870
Xavier-Init + Unbounded Linear	0.5079	0.6403	0.1756	0.1033

Ablation 5: Skip-Connection Feature Routing. We evaluate four feature routing configurations for historical spectral skip connections ($|\mathbf{S}|$ and $\arg(\mathbf{S})$): (i) No skip connection, (ii) Skip to Base Heads only (A_0, ϕ_0 , proposed), (iii) Skip to Modulation Heads only (g, r), and (iv) Dual skip connection to both. Table 6 confirms that routing historical spectral features to the Base Heads yields optimal performance (0.497 on ETTh1, 0.191 on PEMS04 across long horizons). Historical spectral features assist global base spectrum estimation, whereas dynamic modulation heads learn localized temporal adjustments directly from latent features.

Table 6: Ablation 5: Evaluation of skip-connection feature routing choices (averaged over extended horizons $L \in \{96, 192, 336, 720\}$).

Skip-Connection Routing Target	ETTh1 MSE	ETTh1 MAE	PEMS04 MSE	PEMS04 MAE
No Skip Connection ($\mathbf{X}_{\text{enc}}$ only)	0.534	0.518	0.416	0.448
Skip to Base Heads Only (A_0, ϕ_0, Proposed)	0.497	0.510	0.191	0.296
Skip to Modulation Heads Only (g, r)	0.498	0.510	0.199	0.300
Dual Skip Connection (Base + Modulation)	0.507	0.519	0.211	0.313

C QUALITATIVE FORECAST SHOWCASES

Figure 4 visualizes the two failure modes that motivate TVCS-IF and shows how continuous spectral modulation resolves them, complementing the quantitative results in Section 4.

D DATASET SPECIFICATIONS & EXTENDED IMPLEMENTATION DETAILS

D.1 DATASET CHARACTERISTICS

Table 7 provides full details for all 14 benchmark datasets, including channel dimensions N , lookback input lengths T , prediction horizons L , training/validation/test dataset sizes, sampling frequencies, and physical application domains.

D.2 HARDWARE, TENSOR COMPLEXITY, MEMORY FOOTPRINT & LATENCY PROFILING

All experiments were executed on a dedicated computing workstation equipped with a 16-core AMD EPYC 9654 CPU, 128 GB RAM, and a single NVIDIA RTX 4090 GPU (24 GB VRAM). Models




Figure 4: Qualitative forecast comparison between the static Implicit Forecaster baseline and TVCS-IF over a 96-step forecasting horizon. Panel (a) illustrates PEMS04 traffic flow (Channel 112), where the static baseline underestimates a sharp peak due to rigid static Fourier priors, whereas TVCS-IF dynamically tracks the amplitude burst using time-varying gain modulation. Panel (b) shows ECL electricity load (Client 42), where the static baseline accumulates severe phase lag over extended horizons, while TVCS-IF eliminates temporal drift through continuous phase modulation.

Table 7: Comprehensive specifications for the 14 real-world time series benchmark datasets.

Dataset	Dim (N)	Lookback (T)	Forecast Horizon (L)	Dataset Size (Train, Val, Test)	Frequency	Domain
ETTh1	7	96	{96, 192, 336, 720}	(12194, 2613, 2613)	Hourly	Temperature
ETTh2	7	96	{96, 192, 336, 720}	(12194, 2613, 2613)	Hourly	Temperature
ETTm1	7	96	{96, 192, 336, 720}	(48776, 10452, 10452)	15min	Temperature
ETTm2	7	96	{96, 192, 336, 720}	(48776, 10452, 10452)	15min	Temperature
Electricity	321	96	{96, 192, 336, 720}	(18412, 3945, 3947)	Hourly	Energy
Traffic	862	96	{96, 192, 336, 720}	(12280, 2631, 2633)	Hourly	Transportation
Weather	21	96	{96, 192, 336, 720}	(36887, 7904, 7905)	10min	Meteorology
Exchange	8	96	{96, 192, 336, 720}	(5311, 1138, 1139)	Daily	Finance
Illness	7	36	{24, 36, 48, 60}	(676, 144, 146)	Weekly	Healthcare
PEMS03	358	96	{12, 24, 48, 96}	(18345, 3931, 3932)	5min	Transportation
PEMS04	307	96	{12, 24, 48, 96}	(11894, 2548, 2550)	5min	Transportation
PEMS07	883	96	{12, 24, 48, 96}	(19756, 4233, 4235)	5min	Transportation
PEMS08	170	96	{12, 24, 48, 96}	(12499, 2678, 2679)	5min	Transportation
Solar	137	96	{96, 192, 336, 720}	(36792, 7884, 7884)	10min	Energy

were implemented in PyTorch 2.0 utilizing CUDA 11.8. Baseline configurations and hyperparameter settings strictly follow official repository scripts from Time-Series-Library (Wang et al., 2024b).

Tensor-Level Complexity Analysis for Extended Horizons ($L \geq 720$): Standard static spectral decoding predicts static amplitudes $\mathbf{A}_0 \in \mathbb{R}^{B \times N \times (P/2+1)}$ and phases $\phi_0 \in \mathbb{R}^{B \times N \times (P/2+1)}$ and evaluates future sequences via inverse Fast Fourier Transforms (`torch.fft.irfft`) in $O(P \log P)$ time per channel. Because TVCS-IF modulates amplitudes and phases continuously across time steps $t \in [1, L]$, standard iFFT cannot be directly applied. Instead, TVCS-IF executes the following tensor pipeline:

1. *Temporal Frame Evaluation:* Precomputes static raised-cosine basis matrix $\Psi \in \mathbb{R}^{L \times W}$ ($W = 4$).
2. *Modulation Head Projections:* Maps latent features $\mathbf{X}_{\text{enc}} \in \mathbb{R}^{B \times N \times D}$ to weight matrices $\mathbf{W}^g, \mathbf{W}^r \in \mathbb{R}^{B \times N \times K_{\text{mod}} \times W}$.

3. *Continuous Field Expansion*: Computes matrix product $\mathbf{W}^g \Psi^\top \rightarrow g(t) \in \mathbb{R}^{B \times N \times K_{\text{mod}} \times L}$ and $\mathbf{W}^r \Psi^\top \rightarrow r(t) \in \mathbb{R}^{B \times N \times K_{\text{mod}} \times L}$.
4. *Channel-Chunked Harmonic Accumulation*: Evaluates time-varying synthesis (Equation 5) over $K_{\text{mod}} = 128$ modulated low-frequency bins while computing remaining static bins via standard iFFT. Processing in channel chunks ($C_{\text{chunk}} = 32$) with gradient checkpointing maintains peak memory under 1.4 GB VRAM even for $L = 720, N = 321$.

Empirical Runtime and Memory Comparison ($L = 720$):

- **ETTh1** ($N = 7, L = 720$): Static IF requires 42s per epoch with peak memory of 0.24 GB. TVCS-IF ($K_{\text{mod}} = 128, W = 4$) requires 48s per epoch (+14% overhead) with peak memory of 0.31 GB.
- **ECL** ($N = 321, L = 720$): Static IF training takes 65ms/iter with 1.12 GB VRAM. TVCS-IF requires 74ms/iter (+13.8%) with 1.38 GB VRAM (+0.26 GB). Forward pass inference latency increases by less than 1.5ms per batch.
- **PEMS04** ($N = 307, L = 96$): Static IF training requires 662s per epoch (18.33M params). TVCS-IF requires 877s per epoch (22.53M params, +22% compute overhead) while achieving a +20.34% MSE reduction.

E COMPLETE PER-HORIZON EXPERIMENTAL RESULTS

E.1 FULL PER-HORIZON EVALUATION TABLES

Tables 8 through 9 provide complete evaluation metrics (MSE and MAE) comparing static IF against TVCS-IF across all 14 datasets for every individual prediction horizon L .

E.2 STATISTICAL SIGNIFICANCE & STABILITY

To verify stability across initializations, Table 10 reports performance standard deviations for the base backbone equipped with IF over 5 random seeds. Table 11 presents paired t -test p -values comparing standard models against IF-enhanced variants, confirming statistically significant error reductions ($p < 0.01$) across 13 out of 14 benchmark datasets.

Table 8: Detailed per-horizon evaluation on ETTh1, ETTh2, ETTm1, ETTm2, ECL, and Traffic datasets.

Dataset	L	Static IF Baseline		TVCS-IF (Ours)		Absolute Difference	
		MSE	MAE	MSE	MAE	Δ MSE	Δ MAE
ETTh1	96	0.432	0.460	0.425	0.462	-0.007	+0.002
	192	0.471	0.493	0.462	0.490	-0.009	-0.004
	336	0.520	0.527	0.493	0.510	-0.027	-0.017
	720	0.605	0.591	0.600	0.583	-0.005	-0.008
	Avg	0.507	0.518	0.495	0.511	-0.012	-0.007
ETTh2	96	0.151	0.271	0.151	0.270	-0.000	-0.001
	192	0.182	0.297	0.180	0.295	-0.002	-0.003
	336	0.194	0.305	0.191	0.305	-0.003	-0.001
	720	0.226	0.330	0.222	0.328	-0.004	-0.002
	Avg	0.189	0.301	0.186	0.299	-0.002	-0.002
ETTh1	96	0.424	0.435	0.453	0.450	+0.029	+0.015
	192	0.461	0.468	0.467	0.473	+0.007	+0.005
	336	0.516	0.504	0.523	0.509	+0.007	+0.004
	720	0.583	0.551	0.560	0.540	-0.023	-0.012
	Avg	0.496	0.490	0.501	0.493	+0.005	+0.003
ETTh2	96	0.105	0.220	0.106	0.224	+0.001	+0.003
	192	0.132	0.252	0.129	0.249	-0.002	-0.003
	336	0.158	0.274	0.153	0.272	-0.005	-0.002
	720	0.199	0.306	0.194	0.306	-0.006	+0.000
	Avg	0.148	0.263	0.145	0.263	-0.003	-0.000
ECL	96	0.132	0.229	0.134	0.230	+0.001	+0.001
	192	0.148	0.244	0.149	0.246	+0.001	+0.002
	336	0.161	0.262	0.162	0.261	+0.001	-0.001
	720	0.199	0.299	0.205	0.306	+0.005	+0.007
	Avg	0.160	0.259	0.162	0.261	+0.002	+0.002
Traffic	96	0.503	0.306	0.509	0.300	+0.006	-0.007
	192	0.555	0.324	0.534	0.304	-0.021	-0.020
	336	0.568	0.329	0.567	0.323	-0.001	-0.007
	720	0.580	0.336	0.591	0.343	+0.011	+0.007
	Avg	0.552	0.324	0.550	0.317	-0.001	-0.007

Table 9: Detailed per-horizon evaluation on Weather, Exchange Rate, ILI, PEMS03, PEMS04, PEMS07, PEMS08, and Solar-Energy datasets.

Dataset	L	Static IF Baseline		TVCS-IF (Ours)		Absolute Difference	
		MSE	MAE	MSE	MAE	Δ MSE	Δ MAE
Weather	96	0.154	0.201	0.148	0.195	-0.006	-0.006
	192	0.195	0.243	0.195	0.243	-0.000	+0.000
	336	0.242	0.276	0.241	0.277	-0.000	+0.001
	720	0.302	0.320	0.294	0.314	-0.007	-0.005
	Avg	0.223	0.260	0.219	0.257	-0.004	-0.003
Exchange	96	0.082	0.206	0.089	0.212	+0.006	+0.006
	192	0.168	0.300	0.171	0.304	+0.003	+0.005
	336	0.330	0.435	0.369	0.449	+0.039	+0.014
	720	0.582	0.601	0.550	0.583	-0.032	-0.019
	Avg	0.291	0.385	0.295	0.387	+0.004	+0.002
ILI	24	2.677	1.038	2.422	1.016	-0.255	-0.021
	36	2.701	1.109	2.284	1.019	-0.417	-0.091
	48	2.993	1.161	3.160	1.172	+0.166	+0.011
	60	2.712	1.098	2.659	1.081	-0.053	-0.017
	Avg	2.771	1.102	2.631	1.072	-0.140	-0.029
PEMS03	12	0.057	0.159	0.058	0.159	+0.001	-0.000
	24	0.073	0.180	0.073	0.178	-0.000	-0.003
	48	0.109	0.219	0.106	0.211	-0.004	-0.008
	96	0.145	0.260	0.137	0.249	-0.008	-0.012
	Avg	0.096	0.205	0.093	0.199	-0.003	-0.006
PEMS04	12	0.069	0.167	0.068	0.166	-0.001	-0.001
	24	0.077	0.178	0.076	0.176	-0.001	-0.002
	48	0.100	0.211	0.091	0.196	-0.009	-0.015
	96	0.211	0.319	0.130	0.242	-0.082	-0.077
	Avg	0.114	0.219	0.091	0.195	-0.023	-0.024
PEMS07	12	0.052	0.144	0.052	0.142	-0.000	-0.001
	24	0.063	0.158	0.063	0.156	+0.000	-0.003
	48	0.089	0.192	0.081	0.177	-0.008	-0.015
	96	0.122	0.225	0.116	0.216	-0.006	-0.010
	Avg	0.081	0.180	0.078	0.173	-0.003	-0.007
PEMS08	12	0.062	0.163	0.063	0.164	+0.001	+0.001
	24	0.075	0.180	0.075	0.180	-0.001	-0.000
	48	0.105	0.217	0.107	0.218	+0.002	+0.002
	96	0.192	0.295	0.205	0.306	+0.013	+0.012
	Avg	0.109	0.214	0.112	0.217	+0.004	+0.003
Solar	96	0.170	0.230	0.185	0.240	+0.015	+0.009
	192	0.199	0.257	0.202	0.251	+0.003	-0.006
	336	0.194	0.262	0.195	0.258	+0.001	-0.004
	720	0.190	0.258	0.190	0.252	-0.000	-0.006
	Avg	0.188	0.252	0.193	0.250	+0.005	-0.002

Table 10: Performance robustness across 5 random initialization seeds (Metric: Mean $\pm$ Standard Deviation).

Dataset	ETTh2		ECL		Weather	
L	MSE	MAE	MSE	MAE	MSE	MAE
96	0.152 $\pm$ 0.000	0.272 $\pm$ 0.002	0.132 $\pm$ 0.000	0.229 $\pm$ 0.000	0.151 $\pm$ 0.002	0.199 $\pm$ 0.001
192	0.181 $\pm$ 0.002	0.296 $\pm$ 0.002	0.147 $\pm$ 0.001	0.244 $\pm$ 0.000	0.197 $\pm$ 0.001	0.244 $\pm$ 0.002
336	0.195 $\pm$ 0.002	0.307 $\pm$ 0.002	0.161 $\pm$ 0.001	0.262 $\pm$ 0.001	0.237 $\pm$ 0.004	0.274 $\pm$ 0.003
720	0.223 $\pm$ 0.005	0.328 $\pm$ 0.003	0.188 $\pm$ 0.003	0.289 $\pm$ 0.004	0.301 $\pm$ 0.001	0.319 $\pm$ 0.003

Dataset	PEMS03		PEMS07		PEMS08	
L	MSE	MAE	MSE	MAE	MSE	MAE
12	0.058 $\pm$ 0.001	0.160 $\pm$ 0.001	0.052 $\pm$ 0.001	0.144 $\pm$ 0.001	0.062 $\pm$ 0.000	0.163 $\pm$ 0.000
24	0.074 $\pm$ 0.002	0.181 $\pm$ 0.002	0.063 $\pm$ 0.001	0.157 $\pm$ 0.001	0.077 $\pm$ 0.001	0.183 $\pm$ 0.001
48	0.100 $\pm$ 0.002	0.209 $\pm$ 0.002	0.084 $\pm$ 0.002	0.185 $\pm$ 0.002	0.107 $\pm$ 0.002	0.218 $\pm$ 0.002
96	0.145 $\pm$ 0.004	0.259 $\pm$ 0.003	0.120 $\pm$ 0.002	0.221 $\pm$ 0.002	0.165 $\pm$ 0.002	0.273 $\pm$ 0.002

Table 11: Paired t -test statistical significance comparison across random runs ($*p < 0.05$, $**p < 0.01$, $***p < 0.001$, n.s. not significant).

Dataset	Original Baseline MSE	Equipped w/ IF MSE	p -value	Significance
ETTh1	0.539 $\pm$ 0.001	0.504 $\pm$ 0.002	1.98×10^{-5}	***
ETTh2	0.200 $\pm$ 0.001	0.188 $\pm$ 0.001	4.68×10^{-4}	***
ETTh1	0.508 $\pm$ 0.001	0.493 $\pm$ 0.003	1.04×10^{-3}	**
ETTh2	0.155 $\pm$ 0.001	0.148 $\pm$ 0.001	1.23×10^{-5}	***
ECL	0.162 $\pm$ 0.001	0.159 $\pm$ 0.000	5.71×10^{-3}	**
Traffic	0.620 $\pm$ 0.001	0.550 $\pm$ 0.002	7.85×10^{-4}	***
Weather	0.228 $\pm$ 0.000	0.220 $\pm$ 0.000	2.15×10^{-5}	***
Exchange	0.406 $\pm$ 0.004	0.297 $\pm$ 0.003	1.31×10^{-5}	***
ILI	2.985 $\pm$ 0.013	2.770 $\pm$ 0.053	1.03×10^{-2}	*
PEMS03	0.123 $\pm$ 0.001	0.095 $\pm$ 0.000	9.71×10^{-6}	***
PEMS04	0.152 $\pm$ 0.022	0.099 $\pm$ 0.006	3.67×10^{-2}	*
PEMS07	0.103 $\pm$ 0.001	0.081 $\pm$ 0.001	6.70×10^{-6}	***
PEMS08	0.176 $\pm$ 0.011	0.104 $\pm$ 0.002	1.67×10^{-3}	**
Solar	0.188 $\pm$ 0.001	0.189 $\pm$ 0.001	5.44×10^{-1}	n.s.