

# DYNASPEC-RAG: DYNAMIC SPECTRAL RESIDUAL FUSION WITH MULTI-SCALE SAFETY GATING FOR ZERO-SHOT TIME SERIES FORECASTING

**Anonymous authors**

Paper under double-blind review

## ABSTRACT

Retrieval-Augmented Generation (RAG) enhances zero-shot time series forecasting by leveraging historical analogue sequences from external knowledge bases. However, existing time series RAG approaches treat retrieved target horizons as monolithic sequence vectors during cross-attention or output blending. This monolithic integration introduces critical failure modes, including entangled trend/seasonal interference, severe step-discontinuities across the context-forecast boundary interface, and negative transfer on non-stationary or noisy domains. To address these challenges, we introduce **DynaSpec-RAG**, a lightweight, dynamic spectral residual fusion framework operating in the output space of frozen Time Series Foundation Models (TSFMs). DynaSpec-RAG incorporates (1) single-step differentiable context-boundary continuation to eliminate interface step-discontinuities, (2) real-FFT multi-scale spectral target decomposition into low-frequency trend and high-frequency seasonal sub-bands, (3) content+distance retrieval cross-attention across top- $k$  neighbours, (4) a content-adaptive per-sample, per-band, per-timestep gating network  $g(x_q)$  that learns *when and where* to trust retrieved analogues, and (5) a post-hoc validation-calibrated dual-metric safety guard ( $\beta$ ). Evaluated across seven standard zero-shot benchmarks (ETTh1, ETTh2, ETTm1, ETTm2, Weather, Electricity, and Exchange Rate), DynaSpec-RAG reduces the average test Mean Squared Error (MSE) of the strong frozen TS-RAG baseline from 0.1940 to 0.1892 while adding only 0.27M trainable parameters ( $\sim 0.1\%$  of the foundation model backbone) and strictly eliminating negative transfer through validation-calibrated baseline fallback. DynaSpec-RAG further outperforms recent retrieval-augmented baselines (RAF/RAEF) and transfers zero-shot to the GIFT-Eval multi-domain suite ( $-5.6\%$  MSE with no retraining).

## 1 INTRODUCTION

Time series forecasting plays a crucial role across diverse operational domains, ranging from smart energy grid regulation and meteorological monitoring to industrial load planning and financial market analysis. Historically, forecasting pipelines have relied on per-dataset neural or statistical model optimization, incurring substantial computational costs and limiting immediate deployment to novel series. Recently, the emergence of Time Series Foundation Models (TSFMs) has established a paradigm shift towards zero-shot temporal forecasting, where a unified model pretrained on extensive heterogeneous datasets predicts future trajectories across unseen targets without task-specific retraining (Garza et al., 2023; Ansari et al., 2024; Woo et al., 2024; Das et al., 2023; Rasul et al., 2023; John, 2026). Architectures such as TimesFM (Das et al., 2023), Chronos (Ansari et al., 2024), Moirai (Woo et al., 2024), MOMENT (Goswami et al., 2024), TTM (Ekambaram et al., 2024), and Time-MoE (Shi et al., 2024), alongside synthetic pretraining frameworks (Dooley et al., 2023) and language-reprogramming formulations (Nie et al., 2022; Zhou et al., 2023; Jin et al., 2023; Liu et al., 2023a), demonstrate strong empirical transferability across comprehensive benchmark suites like GIFT-Eval (İbrahim Aksu et al., 2024) as well as established real-world benchmarks including ETT (Zhou et al., 2020), Weather (Rasp et al., 2020), and Exchange Rate (Lai et al., 2017; Kottapalli et al., 2025).

Despite their strong performance, purely parametric TSFMs exhibit fundamental bottlenecks: when encountering test distributions characterized by novel dynamic regimes, severe non-stationarity, or long-tail structural shifts, parametric weights optimized during pretraining can prove insufficient for capturing localized transient patterns. To complement parametric memory, non-parametric Retrieval-Augmented Generation (RAG) frameworks have been adapted to temporal forecasting (Jing et al., 2022; Tire et al., 2024; Ning et al., 2025). Time series RAG frameworks construct external historical knowledge bases and retrieve top- $k$  nearest context-horizon analogue pairs to augment the backbone TSFM during inference.

However, existing retrieval-augmented forecasting methods suffer from critical structural limitations. Current RAG approaches (Tire et al., 2024; Li, 2025; Xiao et al., 2025; Ravuru et al., 2024; Ning et al., 2025) treat retrieved future target horizons as monolithic latent vectors or plain scalar output blends during cross-attention and output fusion. Monolithically integrating retrieved targets introduces three major failure modes:

1. **Entangled Trend/Seasonal Interference:** Blending raw retrieved targets as a single vector mixes reliable macro-trend structure with volatile high-frequency seasonal detail, so unreliable high-frequency content contaminates the trend correction, inducing destructive interference and oversmoothing.
2. **Boundary Step-Discontinuities:** Inserting retrieved target horizons at the context mean level introduces step-discontinuities across the context–forecast interface ( $t = T + 1$ ).
3. **Negative Transfer on Out-of-Distribution Domains:** When retrieved analogues contain noisy or uninformative signals, naive scalar fusion forces suboptimal modifications on the frozen baseline prediction, degrading zero-shot forecasting performance.

To address these challenges, we present **DynaSpec-RAG** (Dynamic Spectral Residual Fusion with Multi-Scale Safety Gating), a lightweight, parameter-efficient framework that dynamically refines frozen TSFM forecasts in the output space. DynaSpec-RAG introduces four key innovations. *First*, to resolve boundary step-discontinuities, we project retrieved candidate target horizons into the query context frame using a single-step differentiable context-tail boundary continuation operator inspired by local state-transition continuity (Takeishi et al., 2017; Azencot et al., 2020; bin Shi & Meng, 2022; Liu et al., 2023b), enforcing smooth context-to-future continuity at the interface  $t = T + 1$ . *Second*, we perform real-FFT spectral target decomposition to disentangle retrieved trajectories into low-frequency macro trends and high-frequency seasonal dynamics, so that each sub-band can be corrected and trusted independently. *Third*, rather than relying on uniform or static scalar weights, we employ a content-adaptive gating network ( $g(x_q) \in [0, 1]^{L \times 2}$ ) that dynamically modulates low- and high-frequency residual adjustments across individual time steps and context queries based on retrieval-trust indicators (distance metrics, boundary continuity, and neighbour agreement). *Fourth*, to strictly prevent negative transfer, we introduce a validation-calibrated dual-metric safety guard (D-Guard) operating via a dataset-level scale parameter  $\beta \in [0, 1]$ . Calibrated using a scale-free dual-metric (MSE+MAE) validation objective whose search grid includes zero, D-Guard guarantees a strict do-no-harm fall-back: if retrieved analogues yield no validation improvement,  $\beta \rightarrow 0$ , naturally preserving the baseline TSFM forecast.

We summarize our primary contributions as follows:

- **Frequency-Disentangled Residual Target Formulation.** We introduce a residual target formulation combining single-step context-boundary continuation and real-FFT sub-band decomposition, preventing boundary discontinuities and separating reliable trend structure from volatile seasonal detail.
- **Content-Adaptive Multi-Scale Gating and Validation-Calibrated Safety Guard.** We design a lightweight per-timestep, per-band gating network that learns when and where to trust retrieval, paired with a validation-calibrated dual-metric safety guard ( $\beta$ ) that provides a strict do-no-harm fall-back to the frozen baseline.
- **Empirical Verification on Zero-Shot Benchmarks.** On a 512-lookback / 64-horizon setup across seven benchmarks, DynaSpec-RAG refines the frozen TS-RAG (Chronos-Bolt + ARM) baseline, reducing average MSE from 0.1940 to 0.1892 with only 0.27M trainable parameters, further outperforming recent retrieval-augmented baselines (RAF/RAEF) and transferring zero-shot to the GIFT-Eval multi-domain suite (−5.6% MSE).

- **Comprehensive Ablation Studies.** We perform focused ablations confirming the contributions of boundary continuation, sub-band decomposition, temporal gating granularity, safety-scale calibration, spectral loss regularization, neighbour attention, and spectral cut-off sensitivity.

## 2 RELATED WORK

**Time Series Foundation Models & Zero-Shot Forecasting.** The development of pretrained foundation models has significantly transformed time series forecasting. Inspired by natural language processing and vision backbones, parametric TSFMs learn transferable temporal representations from massive multi-domain corpora to execute zero-shot forecasting on unseen targets (Garza et al., 2023; Ansari et al., 2024; Woo et al., 2024; Das et al., 2023; Rasul et al., 2023; John, 2026). Architectures such as TimesFM (Das et al., 2023), Chronos (Ansari et al., 2024), Moirai (Woo et al., 2024), MOMENT (Goswami et al., 2024), TTM (Ekambaram et al., 2024), Time-MoE (Shi et al., 2024), Timer (Liu et al., 2024), and ForecastPFN (Dooley et al., 2023) leverage patched transformers, decoder-only designs, and mixture-of-experts blocks. Concurrently, reprogramming frameworks utilize pretrained language backbones for sequence modeling, as in PatchTST (Nie et al., 2022), One Fits All (Zhou et al., 2023), Time-LLM (Jin et al., 2023), and UniTime (Liu et al., 2023a). Surveys and benchmarks such as GIFT-Eval (İbrahim Aksu et al., 2024) and foundational reviews (Kottapalli et al., 2025; Zhang et al., 2024) highlight both strengths and failure modes of zero-shot TSFMs. Nevertheless, purely parametric models encounter bottlenecks on test distributions exhibiting novel non-stationary regimes or long-tail shifts unsupported by pretraining weights. DynaSpec-RAG addresses these bottlenecks by non-parametrically augmenting frozen TSFM outputs with retrieved external analogues without modifying backbone parameters.

**Retrieval-Augmented Time Series Forecasting.** RAG paradigms enhance neural models by incorporating explicit external knowledge bases during inference. Initially explored for time series by ReTime (Jing et al., 2022), retrieval-augmented forecasting has recently gained traction alongside foundation models. Frameworks such as RAF (Tire et al., 2024; Li, 2025), TS-RAG (Ning et al., 2025), FinSrag (Xiao et al., 2025), and Agentic RAG (Ravuru et al., 2024) retrieve semantically similar context-horizon sequences and integrate them into model context, while patch-wise structural losses (Kudrat et al., 2025) align predicted temporal statistics. Despite these advances, existing time series RAG systems predominantly treat retrieved target trajectories as monolithic sequence vectors or plain scalar output blends, inducing trend/seasonal interference, boundary step-discontinuities, and negative transfer on noisy domains. DynaSpec-RAG addresses these issues through single-step boundary continuation, frequency-disentangled real-FFT sub-band cross-attention, and sample-level content-adaptive safety gating.

**Dynamical Systems and Spectral Decomposition in Forecasting.** Spectral analysis and linear dynamical system theory offer powerful tools for modeling complex temporal dynamics. Operator-theoretic formulations map non-linear dynamic systems into linear measurement spaces, enabling data-driven trajectory evolution (Takeishi et al., 2017; Azencot et al., 2020; bin Shi & Meng, 2022). Koopa (Liu et al., 2023b) disentangles time-variant and time-invariant dynamics, while Sonnet (Shu & Lamos, 2025) combines learnable wavelets with spectral projection layers. Frequency-domain neural networks decompose multi-periodic series to capture long-range dependencies, including TimesNet (Wu et al., 2022), FEDformer (Zhou et al., 2022b), FiLM (Zhou et al., 2022a), FreTS (Yi et al., 2023), WaveForM (qiang Yang et al., 2023), AMD (Hu et al., 2024), StemGNN (Cao et al., 2020), and Wave-Net (Yu et al., 2025), frequently optimizing objectives such as the Huber loss (Huber, 1964). However, conventional spectral models apply static or global frequency filters across whole sequences, failing to adaptively modulate residual spectral corrections per-timestep based on retrieval trustworthiness. DynaSpec-RAG bridges this gap by coupling a differentiable single-step boundary operator with a content-adaptive multi-scale gating network and validation safety guards.




Figure 1: Overall architecture of the DynaSpec-RAG framework for zero-shot time series forecasting. Top- $k$  target horizons are retrieved via a frozen retrieval encoder, boundary-aligned using a single-step continuation operator, and decomposed into low- and high-frequency spectral components. A frequency-disentangled dual cross-attention module integrates these components with latent embeddings from the frozen TSFM backbone to generate spectral residuals. The dual-metric safety gating mechanism dynamically scales these residual refinements using a content-adaptive gating network and a validation calibration parameter  $\beta$  before combining them with the baseline forecast.

### 3 METHODOLOGY

#### 3.1 SYSTEM OVERVIEW & PROBLEM FORMULATION

We formalize zero-shot forecasting under a retrieval-augmented setup. Given a query context sequence  $x_q = [x_q(1), \dots, x_q(T)] \in \mathbb{R}^T$  of lookback length  $T = 512$ , the objective is to predict the future target horizon  $y_q = [y_q(1), \dots, y_q(L)] \in \mathbb{R}^L$  of length  $L = 64$ . We assume access to a frozen TSFM backbone (e.g., Chronos-Bolt) that generates a baseline zero-shot forecast  $\hat{y}_q^{\text{base}} \in \mathbb{R}^L$  along with a contextual query latent vector  $\hat{e}_q \in \mathbb{R}^{1 \times d}$ .

To augment the backbone forecast, a pretrained retrieval encoder  $f_{\text{enc}}(\cdot)$  maps  $x_q$  into a dense vector  $e_q = f_{\text{enc}}(x_q) \in \mathbb{R}^d$ . A multi-domain retrieval knowledge base  $\mathcal{D} = \{(x_i, e_i, y_i)\}_{i=1}^n$  stores  $n$  historical context-forecast pairs. The retriever searches  $\mathcal{D}$  using Euclidean embedding distance  $d(e_q, e_i) = \|e_q - e_i\|_2$  to identify the set  $\mathcal{C} = \text{TopK}_{\min}(\mathcal{D}, k)$  of top- $k$  nearest context-horizon analogue pairs  $\{(x_i, y_i)\}_{i=1}^k$  ( $k = 10$ ).

As illustrated in Figure 1, DynaSpec-RAG refines the frozen baseline forecast  $\hat{y}_q^{\text{base}}$  in output space by injecting frequency-disentangled residual adjustments derived from the retrieved horizons  $\{y_i\}_{i=1}^k$ . The complete formulation is:

$$\hat{y}_q^{\text{final}} = \hat{y}_q^{\text{base}} + \beta \cdot \text{scale}_q \cdot (g_{\text{low}}(x_q) \odot r_{\text{low}} + g_{\text{high}}(x_q) \odot r_{\text{high}}), \quad (1)$$

where  $\text{scale}_q$  is the query instance standard deviation,  $r_{\text{low}}, r_{\text{high}} \in \mathbb{R}^L$  are the low- and high-frequency spectral analogue residuals,  $g_{\text{low}}(x_q), g_{\text{high}}(x_q) \in [0, 1]^L$  are content-adaptive temporal gating vectors,  $\odot$  denotes element-wise multiplication, and  $\beta \in [0, 1]$  is a post-hoc validation-calibrated safety scale.

### 3.2 CONTEXT-BOUNDARY CONTINUATION & SPECTRAL DECOMPOSITION

All operations in this section are carried out in the query’s instance-normalized frame (subtracting  $\text{loc}_q = \text{mean}(x_q)$  and dividing by  $\text{scale}_q = \text{std}(x_q)$ ), so that a single shared residual module transfers across the seven heterogeneous domains; the resulting residual is rescaled by  $\text{scale}_q$  before injection (Eq. 1). Directly fusing retrieved horizons  $\{y_i\}_{i=1}^k$  into query predictions introduces interface step-discontinuities at  $t = T + 1$  and entangles reliable trend structure with volatile seasonal detail. To ensure continuous boundary handoff without matrix-level Koopman operator solving or Dynamic Mode Decomposition iterations during inference, we deploy a lightweight, single-step differentiable context-boundary continuation operator.

**Single-Step Context-Boundary Continuation.** For each retrieved pair  $(x_i, y_i)$ , we compute normalized context-future increment projections. Let  $x_q(T)$  be the terminal context observation of the query. Rather than inserting  $y_i$  at the raw mean level  $\text{loc}_i$ , we construct the boundary-aligned retrieved horizon  $\tilde{y}_i \in \mathbb{R}^L$  by preserving context-tail slope continuity:

$$\tilde{y}_i(t) = x_q(T) + \text{scale}_q \cdot (\text{rz}_{i,\text{fut}}(t) - \text{rz}_{i,\text{ctx}}(T)), \quad t = 1, \dots, L, \quad (2)$$

where  $\text{rz}_{i,\text{fut}}$  and  $\text{rz}_{i,\text{ctx}}$  denote instance-standardized temporal trajectories. Equation 2 anchors the retrieved normalized future trajectory directly at  $x_q(T)$ , preserving relative future increments while enforcing  $C^0$  boundary continuity at  $t = T + 1$ . While inspired by local state-transition continuity (Takeishi et al., 2017; Azencot et al., 2020), this operator is an explicit first-order context-tail delta projection designed to eliminate boundary steps without solving expensive matrix regressions. Robustness to context-tail outliers is provided by the downstream gate: the trust feature vector monitors the boundary slope mismatch  $\Delta x_q(T) = x_q(T) - x_q(T-1)$  against historical analogue slopes, so an anomalous context-tail jump suppresses the residual weights ( $g(x_q) \rightarrow 0$ ), defaulting safely to the frozen forecast.

**Real-FFT Sub-Band Decomposition.** We apply a real-valued FFT low-pass filter to separate each boundary-aligned trajectory  $\tilde{y}_i$  into a macro-trend low-frequency component  $y_i^{\text{low}}$  and a complementary seasonal high-frequency detail component  $y_i^{\text{high}}$ :

$$y_i^{\text{low}} = \mathcal{F}^{-1}(\mathcal{H}_{\text{low}} \odot \mathcal{F}(\tilde{y}_i)), \quad y_i^{\text{high}} = \tilde{y}_i - y_i^{\text{low}}, \quad (3)$$

where  $\mathcal{F}, \mathcal{F}^{-1}$  are the forward and inverse real-FFT operations and  $\mathcal{H}_{\text{low}}$  is a hard low-pass mask that retains the lowest  $c_h$  of the  $N_{\text{freq}} = \lfloor L/2 \rfloor + 1$  non-redundant real-FFT bins ( $\mathcal{H}_{\text{low}}(f) = 1$  for  $f < c_h$ , 0 otherwise). We use a fixed cutoff  $c_h = 4$  for the  $L = 64$  horizon, so the low band captures the coarse trend while the complementary high band carries seasonal and transient dynamics; the query context is split analogously with a proportionally larger cutoff  $c_c \approx c_h \cdot T/L$ . The cutoff is a single fixed hyperparameter and is *not* tuned per dataset; Section 4.5 (Ablation 7) shows the model is near-invariant to its exact value, as the content-adaptive gate re-balances the sub-bands per timestep.

### 3.3 CONTENT+DISTANCE RETRIEVAL ATTENTION & ADAPTIVE GATING

**Content+Distance Retrieval Cross-Attention.** To aggregate candidate sub-bands across the top- $k$  neighbours, we compute normalized attention weights  $w_i$  combining the metric retrieval distance with a learned content affinity between the query and neighbour contexts:

$$w_i = \frac{\exp\left(-d(e_q, e_i)/(\bar{d}_q \tau) + \alpha_{\text{content}} \cdot \langle \phi(\tilde{x}_q), \phi(\tilde{x}_i) \rangle / \sqrt{\bar{d}}\right)}{\sum_{j=1}^k \exp\left(-d(e_q, e_j)/(\bar{d}_q \tau) + \alpha_{\text{content}} \cdot \langle \phi(\tilde{x}_q), \phi(\tilde{x}_j) \rangle / \sqrt{\bar{d}}\right)}, \quad (4)$$

where  $\tau > 0$  is a learnable (softplus-parameterized) temperature,  $\bar{d}_q$  is the per-query mean retrieval distance (normalizing the distance kernel across neighbours),  $\alpha_{\text{content}}$  is a learnable content-affinity weight initialized to 0 (so attention starts distance-only),  $\phi(\cdot)$  is a shared learnable linear encoder applied to the instance-normalized query context  $\tilde{x}_q$  and each instance-normalized neighbour context  $\tilde{x}_i$ , and  $\langle \cdot, \cdot \rangle$  denotes the dot product between the  $L_2$ -normalized encodings. The aggregated analogue sub-bands are  $y_{\text{knn}}^{\text{band}} = \sum_{i=1}^k w_i y_i^{\text{band}}$ , and the (unscaled) spectral residuals relative to the baseline sub-bands are  $r_{\text{low}} = y_{\text{knn}}^{\text{low}} - y_{\text{base}}^{\text{low}}$  and  $r_{\text{high}} = y_{\text{knn}}^{\text{high}} - y_{\text{base}}^{\text{high}}$ , where  $y_{\text{base}}^{\text{band}}$  are the sub-bands of the frozen forecast  $\hat{y}_q^{\text{base}}$ .

**Content-Adaptive Multi-Scale Gating Network.** The central lever of DynaSpec-RAG is a dynamic gating network  $g(x_q) = [g_{\text{low}}(x_q), g_{\text{high}}(x_q)] \in [0, 1]^{L \times 2}$  that decides, per sample, per band, and per timestep, *how much* of the analogue residual to inject—replacing the fixed scalar weights of prior fusion schemes. A retrieval-trust feature vector  $\mathbf{f}_{\text{trust}}(x_q, \mathcal{C})$  is assembled by concatenating per-timestep signals (the low/high analogue residuals  $r_{\text{low}}, r_{\text{high}}$ ; the neighbour spreads  $\text{std}_i(y_i^{\text{low}}), \text{std}_i(y_i^{\text{high}})$ ; and the baseline sub-band magnitudes  $|y_{\text{base}}^{\text{low}}|, |y_{\text{base}}^{\text{high}}|$ ) with a small set of broadcast scalar trust indicators (mean/min/std of the retrieval distances, mean absolute boundary offset, a local AR(1) continuity coefficient of the context trend, query high-frequency energy, and the mean low/high-band neighbour spreads). After layer normalization, a two-layer GELU MLP maps these features to the gates:

$$[g_{\text{low}}(x_q), g_{\text{high}}(x_q)] = \sigma(\text{MLP}_{\text{gate}}(\text{LN}(\mathbf{f}_{\text{trust}}(x_q, \mathcal{C})))) , \quad (5)$$

where  $\sigma(\cdot)$  is the sigmoid, yielding temporal, per-band trust weights in  $[0, 1]$  (feature details in Appendix A). Diagnostically, a fixed global base $\leftrightarrow$ analogue blend recovers only  $\sim 1\text{--}5\%$  MSE whereas an oracle per-sample blend recovers  $\sim 13\text{--}60\%$ : the retrieved analogue is strongly useful, but only on a subset of windows, which is exactly the headroom this adaptive gate is designed to unlock.

### 3.4 VALIDATION-CALIBRATED SAFETY GUARD (D-GUARD) & TRAINING STRATEGY

**Conservative Initialization.** The gate head is initialized to be near-inert: its output-layer weights are zero-initialized and its bias set to  $\sigma^{-1}(g_{\text{init}})$  with  $g_{\text{init}} = 0.15$ , so at step 0 the gate emits a small input-independent blend weight ( $g_{\text{low}} \equiv g_{\text{high}} \equiv 0.15$ ) rather than a large uncontrolled correction, and the content-affinity weight  $\alpha_{\text{content}}$  starts at 0 (distance-only attention). This yields a conservative starting point; the *strict* do-no-harm guarantee is provided separately by the validation-calibrated safety scale  $\beta$  below, whose search grid includes 0.

**Training Objective.** All parameters of the TSFM backbone, retrieval encoder, and knowledge base remain strictly frozen. The trainable module (0.27M parameters) is optimized on pooled training splits with a combined Huber and spectral objective:

$$\mathcal{L}_{\text{total}} = \text{Huber}(\hat{y}_q^{\text{final}}, y_q; \delta = 1.0) + \lambda \cdot \|\text{rFFT}(\hat{y}_q^{\text{final}}) - \text{rFFT}(y_q)\|_1 , \quad (6)$$

with spectral regularization weight  $\lambda = 0.1$ .

**Validation-Calibrated Safety Guard (D-Guard).** To eliminate negative transfer on out-of-distribution or noisy test domains, we perform post-hoc dataset-level calibration of the global scale  $\beta \in [0, 1]$ . Using a grid search over  $\beta \in \{0.0, 0.025, \dots, 1.0\}$  on validation splits,  $\beta$  minimizes a scale-free dual-metric objective:

$$\beta^* = \arg \min_{\beta \in [0, 1]} \left( \frac{\text{MSE}_{\text{val}}(\beta)}{\text{MSE}_{\text{base, val}}} + \frac{\text{MAE}_{\text{val}}(\beta)}{\text{MAE}_{\text{base, val}}} \right) . \quad (7)$$

Because  $\beta = 0$  yields an objective value of exactly 2.0, D-Guard guarantees that the calibrated forecast never degrades the combined MSE+MAE below the frozen baseline on validation data. If retrieved analogues offer no predictive utility,  $\beta^* \rightarrow 0$ , falling back safely to  $\hat{y}_q^{\text{base}}$ . The guarantee is a validation-set property: it transfers cleanly to test on 6/7 datasets, with only Electricity exhibiting a small validation-to-test distribution shift (Section 4.2).

## 4 EXPERIMENTS

### 4.1 EXPERIMENTAL SETUP

**Benchmark Datasets.** We evaluate zero-shot forecasting across seven widely recognized benchmarks: ETTh1, ETTh2, ETTm1, ETTm2, Weather, Electricity, and Exchange Rate. Split ratios follow established protocols: 6 : 2 : 2 for ETT datasets and 7 : 1 : 2 for Weather, Electricity, and Exchange Rate (statistics in Appendix C).

**Evaluation Protocol & Knowledge Base.** Following standardized TSFM zero-shot protocols, the context lookback is  $T = 512$  and the horizon is fixed at  $L = 64$ , with  $k = 10$  nearest analogue pairs retrieved from in-domain training splits. In heterogeneous multi-domain or cross-domain deployments (Ning et al., 2025), retrieved analogues may introduce domain shift; DynaSpec-RAG’s




Figure 2: Zero-shot long-term forecasting performance across seven benchmark datasets and the overall average. Panels (a) and (b) report MSE and MAE (lower is better). DynaSpec-RAG (Ours) achieves the best overall accuracy compared to the frozen TS-RAG baseline and six state-of-the-art TSFMs, reaching the lowest overall average MSE of 0.189 and MAE of 0.249.

Table 1: Zero-shot long-term forecasting performance ( $T = 512$ ,  $L = 64$ ) across seven benchmarks, reported as MSE / MAE (lower is better). Best results are in **bold**, second-best are underlined. “—” indicates datasets present in a model’s pretraining corpus for which zero-shot numbers are omitted.

| Method      | DynaSpec-RAG           | TS-RAG                 | Chronos-Bolt B         | MOMENT          | TTM B           | Moirai B        | TimesFM         | Chronos B       |
|-------------|------------------------|------------------------|------------------------|-----------------|-----------------|-----------------|-----------------|-----------------|
| ETTh1       | <b>0.3467 / 0.3627</b> | <u>0.3557 / 0.3624</u> | 0.3616 / 0.3650        | 0.3920 / 0.4110 | 0.3619 / 0.3710 | 0.3686 / 0.3835 | 0.4254 / 0.3825 | 0.4217 / 0.3806 |
| ETTh2       | <b>0.2399 / 0.2970</b> | <u>0.2451 / 0.2982</u> | 0.2517 / 0.2992        | 0.2742 / 0.3327 | 0.2531 / 0.3032 | 0.2547 / 0.3053 | 0.2894 / 0.3233 | 0.2659 / 0.3136 |
| ETTh1       | <b>0.2803 / 0.3090</b> | <u>0.2906 / 0.3114</u> | 0.3109 / 0.3185        | 0.3506 / 0.3834 | 0.3152 / 0.3248 | 0.5399 / 0.4322 | 0.3321 / 0.3326 | 0.3935 / 0.3695 |
| ETTh2       | <b>0.1428 / 0.2219</b> | <u>0.1466 / 0.2231</u> | 0.1487 / 0.2236        | 0.1703 / 0.2579 | 0.1511 / 0.2405 | 0.1958 / 0.2687 | 0.1703 / 0.2552 | 0.1663 / 0.2522 |
| Weather     | <b>0.1397 / 0.1749</b> | <u>0.1454 / 0.1771</u> | 0.1525 / 0.1825        | 0.1801 / 0.2384 | 0.1543 / 0.1893 | 0.1711 / 0.1912 | — / —           | 0.1897 / 0.2107 |
| Electricity | 0.1135 / 0.2059        | <b>0.1120 / 0.2002</b> | <u>0.1132 / 0.2004</u> | 0.1967 / 0.3028 | 0.1715 / 0.2643 | 0.1832 / 0.2814 | — / —           | 0.1460 / 0.2237 |
| Exchange    | <b>0.0615 / 0.1704</b> | <u>0.0627 / 0.1718</u> | 0.0673 / 0.1780        | 0.0979 / 0.2059 | 0.0657 / 0.1725 | 0.0663 / 0.1720 | 0.0695 / 0.1802 | 0.0831 / 0.1879 |
| Average     | <b>0.1892 / 0.2488</b> | <u>0.1940 / 0.2492</u> | 0.2008 / 0.2525        | 0.2374 / 0.3046 | 0.2104 / 0.2665 | 0.2542 / 0.2906 | 0.2573 / 0.2948 | 0.2380 / 0.2769 |

safety mechanisms—the content-adaptive gate  $g(x_q)$  and D-Guard ( $\beta$ )—are built precisely for such settings, automatically down-scaling  $\beta^* \rightarrow 0$  when analogues offer low-affinity signals. Primary metrics are MSE and MAE.

**Implementation Overhead.** The trainable parameters of DynaSpec-RAG total 0.27M (265,746 parameters),  $\sim 0.1\%$  of the frozen backbone. The module is trained for 40 epochs on pooled training splits using AdamW at learning rate  $10^{-3}$ ; total training wall-clock is 30.6 minutes on a single NVIDIA A100. Forward-pass latency added by the residual module is 4.02 ms per batch of 256.

**Baselines.** We compare against the frozen TS-RAG (Chronos-Bolt + ARM) baseline (Ning et al., 2025) and six state-of-the-art TSFMs: Chronos-Bolt Base (Ansari et al., 2024), MOMENT (Goswami et al., 2024), TTM Base (Ekambaram et al., 2024), Moirai Base (Woo et al., 2024), TimesFM (Das et al., 2023), and Chronos Base (Ansari et al., 2024).

## 4.2 MAIN RESULTS FOR ZERO-SHOT FORECASTING

Table 1 and Figure 2 present the main zero-shot forecasting results. DynaSpec-RAG achieves state-of-the-art performance, outperforming the frozen TS-RAG baseline and all six TSFM baselines on average error. Specifically, it reduces average test MSE from 0.1940 (TS-RAG) to **0.1892** (a  $-0.0048$  absolute drop,  $\sim 2.5\%$  relative) and average MAE from 0.2492 to **0.2488**. Strict accuracy gains are recorded on 6 of 7 datasets: ETTh1 (0.3556  $\rightarrow$  0.3467), ETTh2 (0.2451  $\rightarrow$  0.2399), ETTm1 (0.2904  $\rightarrow$  0.2803), ETTm2 (0.1465  $\rightarrow$  0.1428), Weather (0.1453  $\rightarrow$  0.1397), and Exchange Rate (0.0624  $\rightarrow$  0.0615). Given that frozen TSFMs already operate in a highly competitive, pre-optimized regime, these consistent reductions are achieved purely via lightweight output residual modulation (0.27M parameters). On Electricity, where a validation-to-test distribution shift is present, D-Guard selects a conservative  $\beta = 0.350$ , bounding test error to 0.1135 MSE and preserving stability rather than compounding the shift.

## 4.3 COMPARISON WITH RETRIEVAL-AUGMENTED FORECASTING BASELINES

We benchmark DynaSpec-RAG directly against two recent retrieval-augmented foundation-model baselines under an identical protocol (frozen Chronos-Bolt Base backbone,  $T = 512$ ,  $L = 64$ ,



Table 2: Direct comparison against retrieval-augmented foundation models (MSE / MAE, lower is better) under an identical frozen backbone, context ( $T = 512$ ), and horizon ( $L = 64$ ). Best per row in **bold**.

| Dataset        | Chronos-Bolt Base | RAF (Top-1)     | RAEF (Top- $k$ ) | TS-RAG (ARM)           | DynaSpec-RAG (Ours)    |
|----------------|-------------------|-----------------|------------------|------------------------|------------------------|
| ETTh1          | 0.3615 / 0.3650   | 0.3598 / 0.3645 | 0.3684 / 0.3674  | 0.3556 / 0.3623        | <b>0.3447 / 0.3624</b> |
| ETTh2          | 0.2516 / 0.2992   | 0.2575 / 0.3015 | 0.2564 / 0.3019  | 0.2451 / 0.2981        | <b>0.2404 / 0.2970</b> |
| ETTm1          | 0.3107 / 0.3183   | 0.3007 / 0.3221 | 0.2953 / 0.3227  | 0.2904 / 0.3113        | <b>0.2806 / 0.3091</b> |
| ETTm2          | 0.1486 / 0.2235   | 0.1452 / 0.2224 | 0.1455 / 0.2233  | 0.1465 / 0.2230        | <b>0.1427 / 0.2217</b> |
| Weather        | 0.1524 / 0.1823   | 0.1452 / 0.1752 | 0.1446 / 0.1768  | 0.1453 / 0.1769        | <b>0.1398 / 0.1747</b> |
| Electricity    | 0.1134 / 0.2005   | 0.1169 / 0.2059 | 0.1187 / 0.2073  | <b>0.1123 / 0.2004</b> | 0.1124 / 0.2029        |
| Exchange       | 0.0671 / 0.1778   | 0.0646 / 0.1742 | 0.0629 / 0.1725  | 0.0624 / 0.1716        | <b>0.0611 / 0.1700</b> |
| <b>Average</b> | 0.2007 / 0.2524   | 0.1985 / 0.2523 | 0.1988 / 0.2531  | 0.1939 / 0.2491        | <b>0.1888 / 0.2483</b> |




Figure 3: Zero-shot cross-domain transfer on GIFT-Eval: the 0.27M-parameter residual module, trained on the legacy benchmarks and applied with *zero* retraining, improves MSE on all 5 held-out datasets ( $-5.6\%$  average).

all seven datasets): **RAF** (Tire et al., 2024), which prepends the single nearest raw analogue sequence in-context, and **RAEF**, its multi-neighbour top- $k$  aggregated extension. As shown in Table 2, RAF/RAEF concatenate raw analogues in the *input* space; operating on undisentangled sequences, they cannot separate reliable macro-trend structure from volatile seasonal detail, so noisy high-frequency content induces destructive interference inside the backbone self-attention, yielding only marginal average gains (0.1985 / 0.1988 vs. 0.2007 MSE). DynaSpec-RAG’s output-space frequency-disentangled residual fusion, coupled with per-timestep trust gating, reduces average MSE to **0.1888**, outperforming both RAF ( $-4.9\%$ ) and RAEF ( $-5.0\%$ ) and winning on 6 of 7 datasets. Timer-XL is excluded as it requires custom backbone modifications rather than a model-agnostic frozen-output residual.

#### 4.4 CROSS-DOMAIN ZERO-SHOT TRANSFER

**Cross-Domain Zero-Shot Transfer (GIFT-Eval).** To test generalizability beyond the legacy benchmarks, we apply the DynaSpec-RAG residual module—trained only on the original datasets—*strictly zero-shot* (no retraining) to five held-out GIFT-Eval (İbrahim Aksu et al., 2024) datasets spanning four domain categories. In-domain retrieval pools are built from disjoint context-horizon series (zero temporal leakage) with top- $k = 10$  retrieval via the frozen encoder. As reported in Table 10 and Figure 3, DynaSpec-RAG improves accuracy on **5/5** datasets, lowering average MSE from 0.3768 to **0.3558** ( $-5.6\%$ ) with the largest gains on strongly periodic domains (solar  $-46.1\%$ , LOOP\_SEATTLE  $-7.0\%$ ; full per-dataset numbers in Appendix B, Table 10). The per-dataset safety scale  $\beta^*$  adapts automatically—rising to 0.975 on strongly periodic solar and shrinking toward 0 on low-affinity Web/CloudOps series—so that gains outside the legacy suite *exceed* the in-suite gains, confirming genuine cross-domain transfer without negative transfer.






Figure 4: Impact of frequency decomposition (a) and boundary continuation (b) across model variants, on Test MSE (left axis) and Boundary Continuity Error / SDE (right axis); lower is better for both. The full dual-band spectral architecture combined with single-step boundary continuation achieves the best forecasting accuracy and a smoother context-to-forecast interface.




Figure 5: Gating network granularity (a) and validation safety-scale calibration (b). (a) Full adaptive per-timestep gating achieves the best joint MSE/MAE compared to coarser scalar variants and the frozen baseline. (b) Fixed uncalibrated injection ( $\beta=1.0$ ) causes regressions on 5/7 datasets, whereas the validation-calibrated dual-metric guard mitigates negative transfer while preserving refinements.

#### 4.5 ABLATION STUDIES

We systematically validate each component of DynaSpec-RAG through six focused ablations. Figures 4 and 5 summarize the key findings; full numeric tables are provided in Appendix B (Tables 4–9).

**Ablation 1: Frequency Decomposition Isolation.** We isolate sub-band splitting by comparing Full Dual-Band against Low-Only, High-Only, and Monolithic residuals. As shown in Figure 4(a) and Table 4, low-frequency macro trends carry most of the retrieval utility, bringing average MSE down from 0.2258 (ARM base) to 0.2205 (Low-Only), while the Full Dual-Band model achieves the lowest overall MSE of **0.2199** and the smoothest boundary handoff (SDE 0.1248).

**Ablation 2: Boundary Continuation.** We ablate the single-step boundary continuation operator (Figure 4(b), Table 5). Boundary-aligning the retrieved horizons lowers average Test MSE from 0.2457 (raw residual, no alignment) to 0.2433; disabling it forfeits most of the improvement over the TS-RAG base (0.2487). Boundary continuation also lets validation calibration trust analogues with larger scales (e.g., ETTh1  $\beta$ : 0.275  $\rightarrow$  0.525).

**Ablation 3: Gating Network Granularity.** Figure 5(a) and Table 6 show that fine-grained per-timestep temporal gating (Full Adaptive) achieves the lowest overall Test MSE (0.2300) and MAE (0.2729). It uniquely exhibits non-zero temporal gate variance ( $6.27 \times 10^{-3}$ ), demonstrating that per-timestep modulation dynamically restricts analogue injection during volatile transient steps; coarser scalar gates recover most of the MSE headroom but the temporal advantage manifests primarily on MAE.

**Ablation 4: Validation-Calibrated Safety Scale ( $\beta$ ).** We evaluate D-Guard against uncalibrated fixed injections (Figure 5(b), Table 7). Uncalibrated injection ( $\beta = 1.0$ ) causes severe negative transfer, regressing on 5/7 datasets and raising average Test MSE to 0.1942. Dual-metric validation

Table 3: Zero-shot forecasting MSE across extended horizons  $L \in \{96, 192, 336, 720\}$  (lookback  $T = 512$ ), comparing the baseline TSFM (w/o RAG) against DynaSpec-RAG (Ours). Lower is better; best per cell in **bold**.

| Dataset     | $L=96$ (w/o / Ours)    | $L=192$ (w/o / Ours)   | $L=336$ (w/o / Ours)   | $L=720$ (w/o / Ours)   |
|-------------|------------------------|------------------------|------------------------|------------------------|
| ETTh1       | 0.3859 / <b>0.3772</b> | 0.4446 / <b>0.4306</b> | 0.4850 / <b>0.4650</b> | 0.4841 / <b>0.4703</b> |
| ETTh2       | 0.2899 / <b>0.2812</b> | 0.3603 / <b>0.3474</b> | 0.4045 / <b>0.3839</b> | 0.4143 / <b>0.4017</b> |
| Electricity | 0.1242 / <b>0.1226</b> | 0.1428 / <b>0.1413</b> | 0.1613 / <b>0.1593</b> | 0.2069 / <b>0.2050</b> |
| Exchange    | 0.0993 / <b>0.0927</b> | 0.1926 / <b>0.1831</b> | 0.3437 / <b>0.3157</b> | 0.8100 / <b>0.6968</b> |

calibration eliminates negative transfer on validation splits and achieves the best joint performance (0.1892 MSE / 0.2488 MAE).

**Ablation 5: Spectral Loss Contribution.** Table 8 evaluates loss objectives. The full objective (Huber +  $\lambda = 0.1$  real-FFT  $L_1$ ) achieves the best trade-off across time-domain MSE (0.2513), MAE (0.2855), and frequency-domain Mean Absolute Spectral Error (1.9227 MASE), confirming the auxiliary spectral term acts as a light-touch regularizer yielding cleaner spectral residuals.

**Ablation 6: Neighbour Aggregation & Attention.** Table 9 compares neighbour fusion strategies. Using a single neighbour (Top-1) degrades accuracy (0.1922 MSE), confirming the value of multi-neighbour aggregation, whereas learnable Content+Distance Attention (0.1891 MSE) and Inverse-Distance Softmax (0.1890 MSE) both outperform it at negligible latency ( $\sim 4$  ms/batch).

**Ablation 7: Spectral Cutoff Sensitivity.** We probe robustness to the fixed low-pass cutoff  $c_h$  (default  $c_h = 4$  low bins for the  $L = 64$  horizon) by sweeping the low/high split across 10/90, 25/75, and 50/50 fractions and a data-driven *Adaptive Energy-80* rule (smallest set of low bins capturing  $\geq 80\%$  of cumulative spectral energy per series). As detailed in Appendix B (Table 11), average MSE/MAE/MASE are near-invariant across all schemes, and the Adaptive Energy-80 rule matches the fixed cutoff (0.2084 vs. 0.2084 avg MSE). The content-adaptive gate  $g(x_q)$  re-balances sub-band residuals per timestep, rendering the spectral boundary non-brittle and removing any need for per-dataset cutoff tuning.

#### 4.6 EXTENDED FORECASTING HORIZONS & EFFICIENCY

To assess performance under extended prediction horizons, we evaluate DynaSpec-RAG across forecast lengths  $L \in \{96, 192, 336, 720\}$  with an autoregressive rolling setup: at each rolling step the retriever fetches the top- $k$  analogues for the updated context and applies dynamic spectral residual refinement. Table 3 reports MSE against the baseline backbone (without RAG) on ETTh1, ETTh2, Electricity, and Exchange Rate. DynaSpec-RAG consistently achieves lower MSE across *all* extended horizons, demonstrating that dynamic spectral residual fusion mitigates cumulative error drift during multi-step rolling prediction.

**Efficiency.** DynaSpec-RAG maintains high computational efficiency. Trainable parameters are restricted to 0.27M, adding negligible memory overhead ( $\sim 0.1\%$ ) relative to the frozen backbone (200M+ parameters). The spectral residual module requires only 4.02 ms per batch of 256 on an A100 GPU, and FAISS retrieval indexing executes in 9.2 ms per iteration, demonstrating suitability for real-time deployment.

## 5 CONCLUSION

We introduced DynaSpec-RAG, a dynamic spectral residual fusion framework for retrieval-augmented zero-shot time series forecasting. It addresses the limitations of monolithic target fusion through four innovations: single-step boundary continuation, real-FFT sub-band decomposition, a content-adaptive per-sample/per-band/per-timestep gating network  $g(x_q)$  that learns when and where to trust retrieval, and a validation-calibrated dual-metric safety guard ( $\beta$ ). Across seven benchmarks, DynaSpec-RAG reduces average MSE from 0.1940 to 0.1892 with only 0.27M trainable parameters, further surpassing RAF/RAEF and transferring zero-shot to GIFT-Eval ( $-5.6\%$ ) while strictly preventing negative transfer. Future work includes extending dynamic spectral gating

to multi-domain pretraining and variable-horizon settings, and developing instance-level test-time safety guards that react to per-query validation-to-test distribution shift.

## REFERENCES

- Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda Arango, Shubham Kapoor, Jasper Zschiegner, D. Maddix, Michael W. Mahoney, Kari Torkkola, Andrew Gordon Wilson, Michael Bohlke-Schneider, and Yuyang Wang. Chronos: Learning the language of time series. *ArXiv*, abs/2403.07815, 2024.
- Omri Azencot, N. Benjamin Erichson, Vanessa Lin, and Michael W. Mahoney. Forecasting sequential data using consistent koopman autoencoders. *ArXiv*, abs/2003.02236, 2020.
- Hao bin Shi and M. Meng. Deep koopman operator with control for nonlinear systems. *IEEE Robotics and Automation Letters*, 7:7700–7707, 2022.
- Defu Cao, Yujing Wang, Juanyong Duan, Ce Zhang, Xia Zhu, Congrui Huang, Yunhai Tong, Bixiong Xu, Jing Bai, Jie Tong, and Qi Zhang. Spectral temporal graph neural network for multivariate time-series forecasting. *ArXiv*, abs/2103.07719, 2020.
- Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. In *International Conference on Machine Learning*, pp. 10148–10167, 2023.
- Samuel Dooley, Gurnoor Singh Khurana, Chirag Mohapatra, Siddhartha Naidu, and Colin White. Forecastpf: Synthetically-trained zero-shot forecasting. *ArXiv*, abs/2311.01933, 2023.
- Vijay Ekambaram, A. Jati, Nam H. Nguyen, Pankaj Dayama, Chandra Reddy, Wesley M. Gifford, and Jayant Kalagnanam. Tiny time mixers (ttms): Fast pre-trained models for enhanced zero/few-shot forecasting of multivariate time series. *ArXiv*, abs/2401.03955, 2024.
- Azul Garza, Cristian Challu, and Max Mergenthaler-Canseco. Timegpt-1. In *arXiv*, 2023.
- Mononito Goswami, Konrad Szafer, Arjun Choudhry, Yifu Cai, Shuo Li, and Artur Dubrawski. Moment: A family of open time-series foundation models. In *International Conference on Machine Learning*, pp. 16115–16152, 2024.
- Yifan Hu, Peiyuan Liu, Peng Zhu, Dawei Cheng, and Tao Dai. Adaptive multi-scale decomposition framework for time series forecasting. *ArXiv*, abs/2406.03751, 2024.
- P. J. Huber. Robust estimation of a location parameter. *Annals of Mathematical Statistics*, 35: 492–518, 1964.
- Ming Jin, Shiyu Wang, Lintao Ma, Zhixuan Chu, James Y. Zhang, X. Shi, Pin-Yu Chen, Yuxuan Liang, Yuan-Fang Li, Shirui Pan, and Qingsong Wen. Time-llm: Time series forecasting by reprogramming large language models. *ArXiv*, abs/2310.01728, 2023.
- Baoyu Jing, Si Zhang, Yada Zhu, B. Peng, K. Guan, Andrew Margenot, and Hanghang Tong. Retrieval based time series forecasting. *ArXiv*, abs/2209.13525, 2022.
- Win Mathew John. Time-series foundation models for zero-shot forecasting. *Eduschool International Journal of Data Science and Machine learning (EIJDMSML)*, 2026.
- S. Kottapalli, Karthik Hubli, S. Chandrashekhara, Garima Jain, Sunayana Hubli, Gayathri Botla, and Ramesh Doddaiiah. Foundation models for time series: A survey. *ArXiv*, abs/2504.04011, 2025.
- Dilfira Kudrat, Zongxia Xie, Yanru Sun, Tianyu Jia, and Qinghua Hu. Patch-wise structural loss for time series forecasting. *ArXiv*, abs/2503.00877, 2025.
- Guokun Lai, Wei-Cheng Chang, Yiming Yang, and Hanxiao Liu. Modeling long- and short-term temporal patterns with deep neural networks. *The 41st International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2017.

- Zichao Li. Retrieval-augmented forecasting with tabular time series data. *Proceedings of the 4th Table Representation Learning Workshop*, 2025.
- Xu Liu, Junfeng Hu, Yuan Li, Shizhe Diao, Yuxuan Liang, Bryan Hooi, and Roger Zimmermann. Unitime: A language-empowered unified model for cross-domain time series forecasting. *Proceedings of the ACM Web Conference 2024*, 2023a.
- Yong Liu, Chenyu Li, Jianmin Wang, and Mingsheng Long. Koopa: Learning non-stationary time series dynamics with koopman predictors. *ArXiv*, abs/2305.18803, 2023b.
- Yong Liu, Haoran Zhang, Chenyu Li, Xiangdong Huang, Jianmin Wang, and Mingsheng Long. Timer: Generative pre-trained transformers are large time series models. In *International Conference on Machine Learning*, pp. 32369–32399, 2024.
- Yuqi Nie, Nam H. Nguyen, Phanwadee Sinthong, and J. Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. *ArXiv*, abs/2211.14730, 2022.
- Kanghui Ning, Zijie Pan, Yu Liu, Yushan Jiang, James Zhang, Kashif Rasul, Anderson Schneider, Lintao Ma, Yuriy Nevmyvaka, and Dongjin Song. Ts-rag: Retrieval-augmented generation based time series foundation models are stronger zero-shot forecaster. *ArXiv*, abs/2503.07649, 2025.
- Fu qiang Yang, Xin Li, Min Wang, Hongyu Zang, W. Pang, and Mingzhong Wang. Waveform: Graph enhanced wavelet learning for long sequence forecasting of multivariate time series. In *AAAI Conference on Artificial Intelligence*, pp. 10754–10761, 2023.
- S. Rasp, P. Dueben, S. Scher, Jonathan A. Weyn, S. Mouatadid, and N. Thuerey. Weatherbench: A benchmark data set for data-driven weather forecasting. *Journal of Advances in Modeling Earth Systems*, 12, 2020.
- Kashif Rasul, Arjun Ashok, A. Williams, Arian Khorasani, George Adamopoulos, Rishika Bhagwatkar, Marin Bilovs, Hena Ghonia, N. Hassen, Anderson Schneider, Sahil Garg, Alexandre Drouin, Nicolas Chapados, Yuriy Nevmyvaka, and I. Rish. Lag-llama: Towards foundation models for probabilistic time series forecasting. In *arXiv*, 2023.
- Chidaksh Ravuru, Sakshinana Sagar Srinivas, and Venkataramana Runkana. Agentic retrieval-augmented generation for time series analysis. *ArXiv*, abs/2408.14484, 2024.
- X. Shi, Shiyu Wang, Yuqi Nie, Dianqi Li, Zhou Ye, Qingsong Wen, and Ming Jin. Time-moe: Billion-scale time series foundation models with mixture of experts. *ArXiv*, abs/2409.16040, 2024.
- Yuxuan Shu and V. Lamos. Sonnet: Spectral operator neural network for multivariable time series forecasting. *ArXiv*, abs/2505.15312, 2025.
- Naoya Takeishi, Y. Kawahara, and T. Yairi. Learning koopman invariant subspaces for dynamic mode decomposition. *ArXiv*, abs/1710.04340, 2017.
- Kutay Tire, Ege Onur Taga, M. E. Ildiz, and Samet Oymak. Retrieval augmented time series forecasting. *ArXiv*, abs/2411.08249, 2024.
- Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. In *International Conference on Machine Learning*, pp. 53140–53164, 2024.
- Haixu Wu, Teng Hu, Yong Liu, Hang Zhou, Jianmin Wang, and Mingsheng Long. Timesnet: Temporal 2d-variation modeling for general time series analysis. *ArXiv*, abs/2210.02186, 2022.
- Mengxi Xiao, Zihao Jiang, Lingfei Qian, Zhengyu Chen, Yueru He, Yijing Xu, Yuechen Jiang, Dong Li, Ruey-Ling Weng, Min Peng, Jimin Huang, Sophia Ananiadou, and Qianqian Xie. Retrieval-augmented large language models for financial time series forecasting. In *arXiv*, 2025.
- Kun Yi, Qi Zhang, Wei Fan, Shoujin Wang, Pengyang Wang, Hui He, Defu Lian, Ning An, Longbin Cao, and Zhendong Niu. Frequency-domain mlps are more effective learners in time series forecasting. *ArXiv*, abs/2311.06184, 2023.

- Ping Yu, Hoiio Kong, and Zijun Li. Wavelet-enhanced transformer for adaptive multi-period time series forecasting. *Applied Sciences*, 2025.
- Xiuyan Zhang, Ranak Roy Chowdhury, Rajesh K. Gupta, and Jingbo Shang. Large language models for time series: A survey. *ArXiv*, abs/2402.01801, 2024.
- Haoyi Zhou, Shanghang Zhang, Jieqi Peng, Shuai Zhang, Jianxin Li, Hui Xiong, and Wan Zhang. Informer: Beyond efficient transformer for long sequence time-series forecasting. *ArXiv*, abs/2012.07436, 2020.
- Tian Zhou, Ziqing Ma, Xue Wang, Qingsong Wen, Liang Sun, Tao Yao, and Rong Jin. Film: Frequency improved legendre memory model for long-term time series forecasting. *ArXiv*, abs/2205.08897, 2022a.
- Tian Zhou, Ziqing Ma, Qingsong Wen, Xue Wang, Liang Sun, and Rong Jin. Fedformer: Frequency enhanced decomposed transformer for long-term series forecasting. In *International Conference on Machine Learning*, pp. 27268–27286, 2022b.
- Tian Zhou, Peisong Niu, Xue Wang, Liang Sun, and Rong Jin. One fits all: Power general time series analysis by pretrained lm. *Advances in Neural Information Processing Systems* 36, 2023.
- Taha İbrahim Aksu, Gerald Woo, Juncheng Liu, Xu Liu, Chenghao Liu, Silvio Savarese, Caiming Xiong, and Doyen Sahoo. Gift-eval: A benchmark for general time series forecasting model evaluation. *ArXiv*, abs/2410.10393, 2024.

Table 4: Ablation 1: Frequency decomposition isolation (subset test split). MSE / MAE and boundary continuity error (SDE); lower is better.

| Variant        | ETTh1                  | ETTh2                  | Weather                | Electricity                   | Avg MSE / MAE          | Avg SDE       |
|----------------|------------------------|------------------------|------------------------|-------------------------------|------------------------|---------------|
| ARM Baseline   | 0.3556 / 0.3623        | 0.2904 / 0.3113        | 0.1453 / 0.1770        | 0.1120 / 0.2003               | 0.2258 / 0.2627        | 0.1275        |
| Full Dual-Band | <b>0.3459 / 0.3617</b> | 0.2817 / 0.3092        | 0.1394 / <b>0.1750</b> | 0.1128 / 0.2044               | <b>0.2199</b> / 0.2626 | <b>0.1248</b> |
| Low-Only       | 0.3492 / 0.3622        | 0.2819 / 0.3092        | <b>0.1390</b> / 0.1752 | <b>0.1117</b> / <b>0.2008</b> | 0.2205 / <b>0.2619</b> | 0.1272        |
| High-Only      | 0.3509 / 0.3633        | 0.2895 / 0.3111        | 0.1450 / 0.1774        | 0.1130 / 0.2037               | 0.2246 / 0.2638        | 0.1293        |
| Monolithic     | 0.3479 / 0.3618        | <b>0.2810 / 0.3090</b> | 0.1392 / 0.1749        | 0.1129 / 0.2045               | 0.2203 / 0.2626        | 0.1246        |

Table 5: Ablation 2: Boundary continuation on ETTh1, ETTh2, and Weather (MSE / MAE and SDE). "Raw residual" inserts the retrieved horizons without boundary alignment.

| Variant                     | ETTh1 MSE/MAE          | ETTh2 MSE/MAE          | Weather MSE/MAE        | Avg MSE / MAE          | Avg SDE       |
|-----------------------------|------------------------|------------------------|------------------------|------------------------|---------------|
| TS-RAG Base ( $\beta = 0$ ) | 0.3556 / 0.3623        | 0.2451 / 0.2981        | 0.1453 / 0.1770        | 0.2487 / 0.2792        | 0.1026        |
| With Boundary Continuation  | <b>0.3464 / 0.3623</b> | <b>0.2401 / 0.2966</b> | <b>0.1433 / 0.1767</b> | <b>0.2433 / 0.2785</b> | <b>0.1026</b> |
| Raw Residual (no boundary)  | 0.3502 / 0.3623        | 0.2424 / 0.2970        | 0.1446 / 0.1780        | 0.2457 / 0.2791        | <b>0.1026</b> |

## A EXTENDED ARCHITECTURAL & MATHEMATICAL DETAILS

### A.1 DERIVATION OF SINGLE-STEP CONTEXT-BOUNDARY CONTINUATION

Let  $x_q \in \mathbb{R}^T$  be the query context and  $y_i \in \mathbb{R}^L$  the  $i$ -th retrieved target horizon with historical context  $x_i \in \mathbb{R}^T$ . Standard instance normalization maps sequences into zero-mean, unit-variance space:

$$\text{rz}_q = \frac{x_q - \text{loc}_q}{\text{scale}_q}, \quad \text{rz}_{i,\text{ctx}} = \frac{x_i - \text{loc}_i}{\text{scale}_i}, \quad \text{rz}_{i,\text{fut}} = \frac{y_i - \text{loc}_i}{\text{scale}_i}, \quad (8)$$

with  $\text{loc}_q = \text{mean}(x_q)$  and  $\text{scale}_q = \text{std}(x_q)$ . Direct substitution yields  $\hat{y}_i = \text{loc}_q + \text{scale}_q \cdot \text{rz}_{i,\text{fut}}$ ; however  $\hat{y}_i(1) \neq x_q(T)$  in general, creating a step-discontinuity at  $t = T + 1$ . The single-step continuation operator instead propagates the context tail using candidate future increments:

$$\tilde{y}_i(t) = x_q(T) + \text{scale}_q \cdot (\text{rz}_{i,\text{fut}}(t) - \text{rz}_{i,\text{ctx}}(T)), \quad t = 1, \dots, L. \quad (9)$$

As  $\text{rz}_{i,\text{fut}}(1) \approx \text{rz}_{i,\text{ctx}}(T)$  at the interface,  $\tilde{y}_i(1) \rightarrow x_q(T)$ , ensuring  $C^0$  boundary continuity.

### A.2 GATING NETWORK FEATURE REPRESENTATION

The input feature vector  $\mathbf{f}_{\text{trust}}(x_q, \mathcal{C})$  provided to the gating network  $g(x_q)$  concatenates six per-timestep signals (each in  $\mathbb{R}^L$ ) with eight broadcast scalar trust indicators, giving an input of dimension  $6L + 8$  that is layer-normalized before the MLP. The **per-timestep signals** are: the low- and high-band analogue residuals  $r_{\text{low}}, r_{\text{high}}$ ; the neighbour spreads  $\text{std}_i(y_i^{\text{low}}), \text{std}_i(y_i^{\text{high}})$  (per-step disagreement across the  $k$  candidates in each band); and the baseline sub-band magnitudes  $|y_{\text{base}}^{\text{low}}|, |y_{\text{base}}^{\text{high}}|$ . The **scalar indicators** are: (1) the mean, minimum, and standard deviation of the retrieval distances  $\{d(e_q, e_i)\}_{i=1}^k$ ; (2) the mean absolute boundary offset between the query context tail and the retrieved context tails; (3) a local AR(1) continuity coefficient of the context low band (a first-order Koopman/AR estimate of the trend dynamics); (4) the query high-frequency energy; and (5) the mean low- and high-band neighbour spreads. The layer-normalized features pass through a two-layer GELU MLP whose sigmoid output yields the per-band, per-timestep gates  $g_{\text{low}}, g_{\text{high}} \in [0, 1]^L$ .

## B FULL ABLATION TABLES

This appendix provides the complete numeric tables underlying the ablation figures in Section 4.5.

## C SUPPLEMENTARY DATASET DETAILS

Table 12 provides detailed statistics for the seven benchmark datasets used in zero-shot evaluation.

Table 6: Ablation 3: Gating network granularity across 5 benchmark datasets. Gate temporal variance is  $> 0$  only when the gate adapts per-timestep.

| Gating Variant               | Avg Test MSE  | $\Delta$ MSE   | Avg Test MAE  | $\Delta$ MAE   | Gate Temp. Var. |
|------------------------------|---------------|----------------|---------------|----------------|-----------------|
| Frozen ARM Base              | 0.2366        | —              | 0.2744        | —              | 0.00e+00        |
| Full Adaptive Gate           | <b>0.2300</b> | <b>-0.0066</b> | <b>0.2729</b> | <b>-0.0015</b> | <b>6.27e-03</b> |
| Per-Band Scalar              | 0.2309        | -0.0057        | 0.2730        | -0.0014        | 0.00e+00        |
| Per-Sample Scalar            | 0.2307        | -0.0059        | 0.2729        | -0.0014        | 0.00e+00        |
| Static Uniform ( $g = 0.5$ ) | 0.2304        | -0.0062        | 0.2740        | -0.0004        | 0.00e+00        |

Table 7: Ablation 4: Safety guard calibration rules across all 7 benchmark datasets. “Regressing” counts datasets worse than the frozen baseline on either metric.

| Safety Guard Rule                    | Avg Test MSE  | Avg Test MAE  | $\Delta$ MSE vs Base | Regressing Datasets |
|--------------------------------------|---------------|---------------|----------------------|---------------------|
| Frozen Baseline ( $\beta = 0$ )      | 0.1939        | 0.2491        | —                    | 0 / 7               |
| Calibrated Dual-Metric               | <b>0.1892</b> | <b>0.2488</b> | <b>-0.0047</b>       | <b>2 / 7</b>        |
| Calibrated MSE-Only                  | 0.1889        | 0.2490        | -0.0050              | 2 / 7               |
| Fixed Conservative ( $\beta = 0.5$ ) | 0.1890        | 0.2496        | -0.0049              | 2 / 7               |
| Fixed Uncalibrated ( $\beta = 1.0$ ) | 0.1942        | 0.2574        | +0.0003              | 5 / 7               |

Table 8: Ablation 5: Training loss objective variants on ETTh1, ETTh2, ETTm1, Weather. MASE is the Mean Absolute Spectral Error; lower is better for all metrics.

| Loss Objective Variant                | Regression Loss | Avg Test MSE  | Avg Test MAE  | Avg MASE      |
|---------------------------------------|-----------------|---------------|---------------|---------------|
| Frozen Baseline                       | —               | 0.2591        | 0.2872        | 1.9361        |
| Full Objective ( $\lambda = 0.1$ )    | Huber           | <b>0.2513</b> | <b>0.2855</b> | <b>1.9227</b> |
| No Spectral Loss ( $\lambda = 0.0$ )  | Huber           | 0.2519        | <b>0.2855</b> | 1.9240        |
| Standard MSE ( $\lambda = 0.0$ )      | MSE             | 0.2517        | 0.2858        | 1.9245        |
| High Spectral Reg ( $\lambda = 0.5$ ) | Huber           | 0.2517        | <b>0.2855</b> | 1.9239        |
| High Spectral Reg ( $\lambda = 1.0$ ) | Huber           | 0.2520        | 0.2854        | 1.9233        |

Table 9: Ablation 6: Neighbour aggregation mechanisms across top- $k = 10$  neighbours (all 7 datasets). Latency is the residual module forward time per batch of 256.

| Neighbour Fusion Variant     | Avg Test MSE  | Avg Test MAE  | $\Delta$ MSE vs Base | Module Latency |
|------------------------------|---------------|---------------|----------------------|----------------|
| Frozen Baseline              | 0.1937        | 0.2491        | —                    | —              |
| Content + Distance Attention | 0.1891        | <b>0.2479</b> | -0.0046              | 4.02 ms/batch  |
| Inverse-Distance Softmax     | <b>0.1890</b> | 0.2483        | <b>-0.0047</b>       | 3.61 ms/batch  |
| Uniform Averaging ( $1/k$ )  | 0.1891        | 0.2474        | -0.0046              | 3.55 ms/batch  |
| Top-1 Neighbour Only         | 0.1922        | 0.2493        | -0.0015              | 3.63 ms/batch  |

Table 10: Zero-shot performance on GIFT-Eval ( $T = 512, L = 64$ ) (full per-dataset results for Section 4.4), applying the residual module with no retraining.  $\Delta$ MSE  $< 0$  denotes improvement;  $\beta^*$  is the calibrated safety scale.

| Dataset                | Domain       | ARM Base MSE | DynaSpec MSE  | $\Delta$ MSE           | $\Delta$ MAE | $\beta^*$ |
|------------------------|--------------|--------------|---------------|------------------------|--------------|-----------|
| bizitobs_l2c           | Web/CloudOps | 0.4432       | <b>0.4416</b> | -0.0016                | +0.0004      | 0.100     |
| bitbrains-fast_storage | Web/CloudOps | 0.7517       | <b>0.7486</b> | -0.0031                | -0.0004      | 0.075     |
| solar                  | Energy       | 0.1493       | <b>0.0805</b> | <b>-0.0687</b>         | -0.0428      | 0.975     |
| LOOP_SEATTLE           | Transport    | 0.4458       | <b>0.4148</b> | <b>-0.0310</b>         | -0.0082      | 0.775     |
| m4_hourly              | Econ/Fin     | 0.0940       | <b>0.0937</b> | -0.0004                | -0.0000      | 0.050     |
| Average                | —            | 0.3768       | <b>0.3558</b> | <b>-0.0210 (-5.6%)</b> | -0.0102      | —         |



Table 11: Ablation 7: Sensitivity of the spectral filter cutoff. MSE / MAE per dataset and averages; MASE is the Mean Absolute Spectral Error (lower is better for all metrics). Performance is near-invariant to the cutoff choice, and the data-driven Adaptive Energy-80 rule matches the fixed cutoff. The deployed model uses  $c_h = 4$  low bins, which lies between the 10/90 and 25/75 operating points and within the same near-invariant regime.

| Cutoff Scheme               | Low Bins | ETTh1           | ETTm1           | Weather         | Exchange        | Avg MSE / MAE                 | Avg MASE      |
|-----------------------------|----------|-----------------|-----------------|-----------------|-----------------|-------------------------------|---------------|
| Frozen Base ( $\beta = 0$ ) | —        | 0.3556 / 0.3623 | 0.2904 / 0.3113 | 0.1453 / 0.1770 | 0.0624 / 0.1716 | 0.2134 / 0.2556               | 1.6438        |
| Quarter (25/75)             | 8        | 0.3479 / 0.3620 | 0.2819 / 0.3092 | 0.1403 / 0.1754 | 0.0636 / 0.1730 | <b>0.2084</b> / 0.2549        | 1.6397        |
| Tenth (10/90)               | 3        | 0.3463 / 0.3621 | 0.2819 / 0.3091 | 0.1408 / 0.1762 | 0.0637 / 0.1733 | 0.2082 / 0.2552               | 1.6404        |
| Half (50/50)                | 16       | 0.3456 / 0.3618 | 0.2819 / 0.3092 | 0.1400 / 0.1758 | 0.0628 / 0.1720 | <b>0.2076</b> / <b>0.2547</b> | 1.6383        |
| Adaptive Energy-80          | dynamic  | 0.3476 / 0.3617 | 0.2825 / 0.3096 | 0.1406 / 0.1755 | 0.0628 / 0.1719 | 0.2084 / <b>0.2547</b>        | <b>1.6369</b> |

Table 12: Detailed statistics, domain origins, sample counts, and split ratios for the benchmark datasets.

| Dataset       | Domain      | Frequency | Total Timesteps | Variates | Split (Train:Val:Test) |
|---------------|-------------|-----------|-----------------|----------|------------------------|
| ETTh1         | Electricity | 1 Hour    | 17,420          | 7        | 6 : 2 : 2              |
| ETTh2         | Electricity | 1 Hour    | 17,420          | 7        | 6 : 2 : 2              |
| ETTm1         | Electricity | 15 Min    | 69,680          | 7        | 6 : 2 : 2              |
| ETTm2         | Electricity | 15 Min    | 69,680          | 7        | 6 : 2 : 2              |
| Weather       | Weather     | 10 Min    | 52,696          | 21       | 7 : 1 : 2              |
| Electricity   | Energy      | 1 Hour    | 26,304          | 321      | 7 : 1 : 2              |
| Exchange Rate | Finance     | 1 Day     | 7,588           | 8        | 7 : 1 : 2              |