

Legendrian Contact-Geometric Discount Optimization with Dimension-Adaptive Dirichlet Regularization for High-Dimensional Quality-Diversity

Anonymous authors
Paper under double-blind review

Abstract

Quality-Diversity (QD) algorithms aim to discover large archives of diverse, high-performing solutions, but their exploration capabilities degrade severely in non-uniform or high-dimensional measure spaces ($k \geq 10$) due to measure space distortion. While Discount Model Search (DMS) mitigates distortion by replacing discrete cell thresholds with a continuous Multi-Layer Perceptron (MLP) discount surrogate, sparse emitter sampling in high dimensions causes unregularized MLPs to learn spurious high-frequency spatial wiggles in unvisited gaps, miscalibrating improvement rankings and stalling frontier expansion. In this paper, we propose Legendrian Contact-Geometric Discount Optimization with Dimension-Adaptive Dirichlet Regularization (LC-FTT), a continuous QD framework that enforces smooth, non-oscillatory discount landscapes without over-constraining expressivity. LC-FTT introduces a stochastic finite-difference contact-form Dirichlet penalty into the neural discount loss, scaled by a dimension-adaptive multiplier $c(k) = c_0 \cdot \max(0, 1 - 2/k)$ with $n_{\text{dirs}} = 16$ Monte Carlo directional estimates. The penalty automatically deactivates in two-dimensional measure spaces ($c(2) = 0$) to preserve dense fitting parity, while smoothly saturating in higher dimensions to suppress spurious interpolation wiggles. Extensive evaluations across 11 benchmark domains (2D to 784D) demonstrate that LC-FTT matches DMS in 2D and delivers statistically significant Quality-Diversity gains at $k=10$ (+422 QD score on 10D LP Sphere, +217 on 10D LP Rastrigin; both $p < 0.01$), while guaranteeing non-regressive performance parity in extreme dimensions ($k \geq 20$, up to 784D). On the Brax physical continuous-control platform, activating the regularizer over a 10D behavioral descriptor accelerates early exploration on Humanoid (+3.9 pp coverage at a fixed budget, $p=0.004$) without harming final convergence, confirming transfer to non-linear physical dynamics.

1 Introduction

Quality-Diversity (QD) optimization has emerged as a powerful paradigm in stochastic search and evolutionary computation, aiming to discover collections of solutions that are both high-performing with respect to an objective function and behaviorally diverse across a user-defined measure space Pugh et al. (2016); Chatzilygeroudis et al. (2020). By constructing expansive behavioral repertoires, QD algorithms facilitate rapid online adaptation in robotics Cully et al. (2014); Cully (2019), autonomous design illumination Mouret & Clune (2015), and complex policy synthesis Colas et al. (2020); Tjanaka et al. (2022a). Classical QD algorithms, such as Multi-dimensional Archive of Phenotypic Elites (MAP-Elites) Mouret & Clune (2015) and Centroidal Voronoi Tessellations MAP-Elites (CVT-MAP-Elites) Vassiliades et al. (2016), maintain archives by partitioning the continuous measure space into discrete spatial cells. While effective in low-dimensional spaces (e.g., $k \leq 2$), discrete archive tessellations suffer severely from spatial distortion and the curse of dimensionality when the measure space dimension k increases ($k \geq 10$), as exponential cell

counts cause most cells to remain unoccupied while complex non-linear mappings compress disparate solutions into identical discrete bins.

To overcome the rigid spatial boundaries of grid archives, surrogate-assisted Quality-Diversity approaches have gained traction Gaier et al. (2018); Kent & Branke (2020); Bhatt et al. (2022). Recently, Discount Model Search (DMS) Tjanaka et al. (2026) addressed measure space distortion by introducing a continuous neural discount surrogate. Rather than storing discrete thresholds in grid cells, DMS fits a Multi-Layer Perceptron (MLP) $\hat{f}_A(s; \psi)$ to represent continuous discount values across measure space $s \in \mathbb{R}^k$. Stochastic emitters, such as the Covariance Matrix Adaptation Evolution Strategy (CMA-ES), optimize candidate solutions based on relative archive improvement rankings $\Delta_i = f(\theta_i) - \hat{f}_A(m(\theta_i))$, guiding exploration into unvisited measure regions.

Despite its success, neural discount modeling faces a vulnerability in high-dimensional measure spaces ($k \geq 10$). Because emitters sample sparsely across high-dimensional manifolds, the discount model must interpolate across expansive unvisited gaps, and unregularized MLPs trained on sparse targets suffer from spectral bias Rahaman et al. (2018), learning spurious high-frequency oscillations (“wiggles”) that distort the predictions $\hat{f}_A(s)$ and produce miscalibrated improvement signals Δ_i ; emitters then converge prematurely to false local modes, stalling frontier expansion and archive coverage. Conversely, naive global smoothing or rigid low-rank/Lipschitz constraints over-constrain the landscape, preventing the steep, localized discount ratchets required near competitive archive boundaries Cai & Wang (2020); Akansha (2023).

In this work, we present Legendrian Contact-Geometric Discount Optimization (LC-FTT), a continuous, mathematically principled QD framework that eliminates high-frequency discount wiggles while preserving macro-regional expressivity. LC-FTT introduces a lightweight, dimension-adaptive Legendrian contact-form Dirichlet-energy regularizer directly into the neural discount loss:

$$\mathcal{L}(\psi) = \mathcal{L}_{\text{MSE}}(\psi) + c(k) \cdot \mathcal{E}_{\text{dir}}(\psi), \quad (1)$$

where $\mathcal{E}_{\text{dir}}$ measures the local contact-form energy of the discount potential surface via multi-directional finite differences. Crucially, the regularization strength is governed by a dimension-adaptive scaling function $c(k) = c_0 \cdot \max(0, 1 - 2/k)$. In two-dimensional measure spaces ($k = 2$), $c(2) = 0$, completely deactivating the penalty so that LC-FTT preserves dense fitting accuracy where standard MLPs already fit seamlessly. In higher dimensions ($k \geq 10$), $c(k)$ automatically scales up to suppress sparse interpolation oscillations. To ensure gradient stability, LC-FTT estimates the contact-form energy stochastically by averaging over $n_{\text{dirs}} = 16$ random directional perturbations per step, drastically reducing directional Monte Carlo variance without incurring intractable computational overhead. Figure 1 summarizes the complete framework and the mechanism that distinguishes it from unregularized DMS.

We evaluate LC-FTT across 11 benchmark domains spanning $k = 2$ to $k = 784$, against DMS, CMA-MAE, DDS, and MAP-Elites variants, recent grid-free baselines (DNS, SQUAD), and Brax continuous-control physics. Backed by formal statistical testing (Welch’s t-test and Mann–Whitney U over 5 seeds), LC-FTT matches DMS in low dimensions and yields statistically significant high-dimensional gains at $k=10$ with non-regressive parity in extreme dimensions ($k \geq 20$). Ablations validate the dimension-adaptive scaling $c(k)$, multi-directional variance reduction (n_{dirs}), step-size stability (ϵ), and contact-form L2 Dirichlet smoothing over alternative structural regularizers.

2 Related Work

2.1 Grid-Based and Emitter-Driven Quality-Diversity Algorithms

Quality-Diversity optimization originated with grid-based illumination algorithms, notably MAP-Elites Mouret & Clune (2015), which partitions the measure space into a uniform grid of discrete cells to retain high-performing elite solutions within each behavioral niche. To

improve scalability in continuous measure spaces, Centroidal Voronoi Tessellations MAP-Elites (CVT-MAP-Elites) Vassiliades et al. (2016) uses k-means clustering to distribute archive cells homogeneously across high-dimensional feature spaces. Subsequent works integrated derivative-free optimization emitters to accelerate behavioral illumination; Covariance Matrix Adaptation MAP-Elites (CMA-ME) Fontaine et al. (2019) introduces CMA-ES search emitters guided by archive improvement rules. Covariance Matrix Adaptation MAP-Annealing (CMA-MAE) Fontaine & Nikolaidis (2022); Zhao et al. (2024) further enhances exploration by introducing dynamic threshold annealing, preventing premature convergence on flat or deceptively optimal landscapes. In parallel, Policy Gradient Assisted MAP-Elites (PGA-MAP-Elites) Nilsson & Cully (2021) leverages deep reinforcement learning critic gradients to optimize high-dimensional parameter spaces.

Recent advances explore dynamic local-competition mechanisms: Dominated Novelty Search Bahlous-Boldi et al. (2025) replaces archive grids with dynamic fitness transforms, and Vector Quantized-Elites Tsakonas & Chatzilygeroudis (2025) uses vector-quantized autoencoders to construct structured behavior grids. However, all grid-based and threshold-annealing frameworks rely on discrete tessellations or rigid cell boundaries: in distorted or high-dimensional measure spaces ($k \geq 10$), continuous mappings collapse into identical bins, stalling emitter selection and improvement rankings.

2.2 Continuous Density and Gradient-Based QD Optimization

To break free from spatial discretization, continuous QD frameworks model measure space occupancy continuously. Density Descent Search (DDS) Lee et al. (2023) replaces discrete cell grids with continuous density estimates, employing Kernel Density Estimation (KDE) to steer CMA-ES emitters toward low-density behavioral regions. To accelerate density estimation in higher dimensions, sub-linear hashing-based estimators (HBE) Charikar & Siminelakis (2017) have been investigated. When exact mathematical gradients of the objective and measure functions are accessible, Differentiable Quality Diversity (DQD) algorithms such as MAP-Elites via a Gradient Arborescence (MEGA) Fontaine & Nikolaidis (2021) exploit first-order information to rapidly traverse the joint objective-measure manifold. In non-differentiable reinforcement learning domains, DQD methods approximate performance and measure gradients via ES or RL policy perturbations Tjanaka et al. (2022b;a).

More recently, AutoQD Hedayatian & Nikolaidis (2025) leverages random Fourier features and occupancy measures to discover diverse behaviors without manual measure specification, and geometric dynamics on contact/symplectic manifolds Chen et al. (2019); Djordjević et al. (2025) provide frameworks for trajectory evolution in continuous state spaces. Despite these advances, KDE-based methods like DDS scale linearly in memory and time with the evaluation history, rendering long-run optimization intractable, while DQD methods require analytical differentiability that is absent in procedural black-box benchmarks.

2.3 Surrogate-Assisted QD and Neural Landscape Smoothing

When objective and measure evaluations are computationally expensive, surrogate-assisted QD algorithms build predictive models of the search landscape. Surrogate-Assisted Illumination (SAIL) Gaier et al. (2018) pioneers Gaussian Process (GP) surrogates to illuminate aerodynamic design spaces with minimal evaluations. Bayesian Optimization of Elites (BOP-Elites) Kent & Branke (2020) models quality and diversity niches simultaneously using GP acquisition functions. To handle higher-dimensional environments, Deep Surrogate Assisted Generation of Environments (DSAGE) Bhatt et al. (2022) and related frameworks Bhatt et al. (2023) employ deep neural networks to predict complex agent behaviors and interaction scenarios. Most recently, Discount Model Search (DMS) Tjanaka et al. (2026) replaces discrete archive thresholds altogether by fitting a continuous neural discount model $\hat{f}_A(s)$ via multi-layer perceptrons, allowing smooth improvement ranking across continuous measure spaces.

However, unregularized neural surrogates trained on sparse, high-dimensional emitter samples suffer from implicit gradient bias Soudry et al. (2017) and spectral bias Rahaman et al. (2018), fitting spurious high-frequency wiggles in unexplored measure gaps. While classical

regularization techniques like Gradient Penalties (WGAN-GP) Gulrajani et al. (2017) or Spectral Normalization Miyato et al. (2018) constrain network Lipschitz constants, global norm caps over-smooth multi-basin landscapes and impair fine-grained discount ratchets near archive boundaries. Similarly, Dirichlet energy constraints in deep architectures Cai & Wang (2020); Zhou et al. (2021); Akansha (2023) have been studied to prevent over-smoothing in graph representation learning. LC-FTT bridges these paradigms by introducing a dimension-adaptive Legendrian contact-form Dirichlet regularizer. By deactivating in low dimensions ($k = 2$) and dynamically scaling in high dimensions ($k \geq 10$), LC-FTT eliminates spurious interpolation wiggles while preserving the expressivity necessary for continuous quality-diversity illumination.

3 Methodology

3.1 Problem Formulation & Discount Model Search (DMS) Baseline

We formalize Quality-Diversity (QD) optimization over a continuous solution space $\Theta \subseteq \mathbb{R}^n$, an objective function $f : \Theta \rightarrow \mathbb{R}$, and a vector-valued measure function $m : \Theta \rightarrow \mathbb{R}^k$. The goal is to discover a collection of solutions that maximize $f(\theta)$ while spanning the range of $m(\theta)$.

Discount Model Search (DMS) Tjanaka et al. (2026) maintains an archive A of elite solutions and guides W stochastic search emitters (e.g., CMA-ES) using a continuous neural discount surrogate $\hat{f}_A(s; \psi)$ parameterized by ψ . The discount surrogate predicts scalar discount values across measure coordinates $s \in \mathbb{R}^k$. On iteration t , each emitter $w \in \{1, \dots, W\}$ samples a batch of λ solutions $\{\theta_i\}_{i=1}^\lambda$ from a multivariate Gaussian distribution $\mathcal{N}(\theta^*, \Sigma)$. For each candidate θ_i , DMS evaluates objective $f(\theta_i)$ and measure $s_i = m(\theta_i)$, and computes the continuous improvement value Δ_i :

$$\Delta_i = f(\theta_i) - \hat{f}_A(m(\theta_i); \psi). \quad (2)$$

Emitters adapt their search distributions based on candidate solution rankings ordered by Δ_i . To update the discount surrogate, DMS constructs a target dataset D_A . For each sampled solution θ_i , a discount target $t_{A,i}$ is constructed following the threshold update rule:

$$t_{A,i} = \begin{cases} \hat{f}_A(m(\theta_i); \psi) & \text{if } f(\theta_i) \leq \hat{f}_A(m(\theta_i); \psi), \\ (1 - \alpha)\hat{f}_A(m(\theta_i); \psi) + \alpha f(\theta_i) & \text{if } f(\theta_i) > \hat{f}_A(m(\theta_i); \psi), \end{cases} \quad (3)$$

where $\alpha \in (0, 1]$ represents the archive learning rate. Furthermore, to clamp unvisited regions, DMS samples n_{empty} unoccupied archive cell centroids $\{s_e\}_{e=1}^{n_{\text{empty}}}$ and appends them to D_A with target $f_{\min}$.

Failure Mode in High Dimensions ($k \geq 10$): Because emitters sample sparse trajectories, vast regions between visited measures s_i and empty anchors s_e remain unsampled; training an unregularized MLP on D_A then learns spurious high-frequency oscillations in these gaps (spectral bias Rahaman et al. (2018)). As illustrated in Figure 1 (B1), these wiggles create false local minima in $\hat{f}_A(s)$, producing erratic improvement signals Δ_i that trap emitters and stall exploration.

3.2 Legendrian Contact-Form Dirichlet Regularization

To enforce spatial smoothness without sacrificing local expressivity, we formulate the expanded measure-objective landscape on a 1-jet contact manifold $J^1(\mathbb{R}^k, \mathbb{R}) \cong \mathbb{R}^{2k+1}$ equipped with local canonical coordinates (s, f, p) , where $s = m(\theta)$ is the measure coordinate, $f = f(\theta)$ is the objective potential, and $p = \nabla_s f$ is the continuous objective gradient. The manifold is endowed with the canonical contact 1-form $\alpha = df - p^T ds$.

On a Legendrian submanifold where $\alpha = 0$, the continuous discount potential surface $\hat{f}_A(s; \psi)$ satisfies $p = \nabla_s \hat{f}_A(s; \psi)$. To suppress high-frequency oscillations, LC-FTT regularizes the continuous contact-form Dirichlet energy $\mathcal{E}_{\text{dir}}(\psi)$, defined as the expected squared




Figure 1: Overview of the LC-FTT framework. (A) Pipeline: W CMA-ES emitters propose solutions evaluated for $f(\theta)$ and $m(\theta)$; a shared MLP discount model $\hat{f}_A(s)$ is trained with a dual loss combining the DMS reconstruction term and our contact-form Dirichlet regularizer at strength $c(k)$, and the calibrated improvement $\Delta f = f(\theta) - \hat{f}_A(m(\theta))$ drives the feedback loop. (B) Mechanism: an unregularized DMS surface (B1) develops spurious high-frequency wiggles in sparsely sampled regions that trap emitters, whereas LC-FTT (B2) enforces a smooth surface yielding reliable improvement gradients and steady frontier expansion.

norm of the discount gradient:

$$\mathcal{E}_{\text{dir}}(\psi) = \mathbb{E}_{s \sim D_A} \left[\left\| \nabla_s \hat{f}_A(s; \psi) \right\|_2^2 \right]. \quad (4)$$

Relationship to standard L_2 gradient regularization. We emphasize that, operationally, $\mathcal{E}_{\text{dir}}$ is identical to a standard Euclidean L_2 gradient-norm (Dirichlet) penalty on the discount surface: on the Legendrian 0-slice the momentum coordinate collapses to $p = \nabla_s \hat{f}_A$, so the contact-form energy contributes no term beyond $\|\nabla_s \hat{f}_A\|_2^2$ and yields no distinct analytical bound. The contact-geometric construction is therefore offered as geometric intuition for why a smooth continuous discount potential is desirable, not as theoretical window-dressing. The operative contribution of LC-FTT is instead the dimension-adaptive coefficient $c(k)$ (Eq. 6) together with its stochastic multi-directional finite-difference estimator (Section 3); as our ablations show (Appendix Figure 5), a fixed-coefficient L_2 penalty or structural spectral normalization either fails in high dimensions or over-smooths multi-basin landscapes. The total loss function minimized during discount surrogate updates is given by:

$$\mathcal{L}(\psi) = \frac{1}{|D_A|} \sum_{(s, t_A) \in D_A} \left(\hat{f}_A(s; \psi) - t_A \right)^2 + c(k) \cdot \mathcal{E}_{\text{dir}}(\psi), \quad (5)$$

where $c(k)$ is a dimension-adaptive scaling multiplier:

$$c(k) = c_0 \cdot \max\left(0, 1 - \frac{2}{k}\right), \quad \text{with } c_0 = 0.00625. \quad (6)$$

Mathematical Intuition: In low dimensions ($k = 2$), $c(2) = 0$ deactivates the penalty, since dense sampling already lets MLPs fit the landscape without interpolation wiggles. For $k > 2$, $c(k)$ smoothly saturates toward c_0 ($c(10) = 0.0050$, $c(20) = 0.0056$, $c(50) = 0.0060$, $c(784) = 0.0062$), penalizing steep gradients in sparse unvisited gaps to enforce a smooth discount surface (Figure 1 (B2)).

3.3 Stochastic Finite-Difference Estimator and Optimization Loop

Computing exact analytical gradients $\nabla_s \hat{f}_A(s; \psi)$ via backpropagation for every target sample in high dimensions introduces significant computational overhead. Instead, LC-FTT

Algorithm 1 Legendrian Contact-Geometric Discount Optimization (LC-FTT)

Input: Objective $f(\cdot)$, measure $m(\cdot)$, initial solution θ_0 , iterations N , emitters W , batch size λ , learning rate α , empty points n_{empty} , base multiplier $c_0 = 0.00625$, directions $n_{\text{dirs}} = 16$, step size $\epsilon = 0.05$

Output: Elite archive A

- 1: Initialize archive A , discount model $\hat{f}_A(\cdot; \psi)$ with parameters ψ , and W CMA-ES emitters at θ_0 .
 - 2: Compute dimension-adaptive coefficient $c(k) = c_0 \cdot \max(0, 1 - 2/k)$.
 - 3: for iter = 1 $\dots$ N do
 - 4: $D_A \leftarrow \emptyset$
 - 5: for emitter $w = 1 \dots W$ do
 - 6: Sample batch $\{\theta_i\}_{i=1}^\lambda \sim \mathcal{N}(\theta_w^*, \Sigma_w)$.
 - 7: Evaluate $f(\theta_i)$ and $s_i = m(\theta_i)$ for all $i \in \{1 \dots \lambda\}$.
 - 8: Compute improvement $\Delta_i = f(\theta_i) - \hat{f}_A(s_i; \psi)$.
 - 9: Compute target $t_{A,i}$ using Eq. 3 and add $(s_i, t_{A,i})$ to D_A .
 - 10: Insert θ_i into archive A if $f(\theta_i) > f(A[s_i])$.
 - 11: Update emitter mean θ_w^* and covariance Σ_w based on rankings by Δ_i .
 - 12: if emitter converges then
 - 13: Restart emitter from a randomly selected elite in A .
 - 14: end if
 - 15: end for
 - 16: Sample n_{empty} unoccupied cell centroids $\{s_e\}_{e=1}^{n_{\text{empty}}}$ from A and append $(s_e, f_{\min})$ to D_A .
 - 17: Update discount MLP ψ with Adam to minimize $\mathcal{L}(\psi) = \mathcal{L}_{\text{MSE}}(\psi) + c(k)\hat{\mathcal{E}}_{\text{dir}}(\psi)$ using Eq. 7.
 - 18: end for
-

estimates $\mathcal{E}_{\text{dir}}(\psi)$ stochastically using multi-directional finite differences along n_{dirs} random unit vectors $v \sim \mathbb{S}^{k-1}$:

$$\hat{\mathcal{E}}_{\text{dir}}(\psi) = \frac{1}{|D_A| n_{\text{dirs}}} \sum_{s \in D_A} \sum_{j=1}^{n_{\text{dirs}}} \left(\frac{\hat{f}_A(s + \epsilon v_j; \psi) - \hat{f}_A(s; \psi)}{\epsilon} \right)^2, \quad (7)$$

where $\epsilon = 0.05$ is the finite-difference perturbation step size, and $n_{\text{dirs}} = 16$ Monte Carlo directional samples are drawn per input measure s . As shown in our ablation analysis, $n_{\text{dirs}} = 16$ achieves a 16-fold gradient variance reduction over a single directional probe, stabilizing Adam updates on ψ .

Algorithm 1 outlines the complete per-iteration optimization loop of LC-FTT.

4 Experiments

4.1 Experimental Setup and Benchmark Domains

We evaluate LC-FTT across 11 diverse benchmark domains categorized by measure space dimensionality k :

1. Low-Dimensional Benchmarks ($k = 2$): 2D Linear Projection (LP) with Sphere, Rastrigin, and Flat objectives ($n = 100$ solution dimensions, 100×100 grid archive), and 2D Arm Repertoire ($n = 100$ planar robotic arm joint angles).
2. High-Dimensional LP Benchmarks ($k \in \{10, 20, 50\}$): 10D LP (Sphere, Rastrigin, Flat), 20D LP (Sphere), and 50D LP (Sphere) ($n = 100$ solution dimensions, 10,000-cell CVT archive).
3. Ultra-High-Dimensional Vision Domains ($k = 784$): Triangle Arrangement (TA) on MNIST and Fashion-MNIST ($n = 240$ triangle parameters rendering 28×28 grayscale images, 1,000-cell CVT archive).

Table 1: Mean Quality-Diversity (QD) Score and Archive Coverage (%) across 11 benchmark domains after 10,000 iterations over 5 independent random seeds (mean / s.e.m.). Format: QD Score / Coverage %. * marks the top-performing value in each column; † marks a statistically significant improvement over DMS ($p < 0.01$, Mann–Whitney U over 5 seeds; see Section 4.3 and Table 5). N/A indicates non-runnable configurations.

Algorithm	2D LP (Sph)	2D LP (Rast)	2D LP (Flat)	Arm Rep	10D LP (Sph)	10D LP (Rast)	10D LP (Flat)	20D LP (Sph)	50D LP (Sph)	TA (MNIST)	TA (F-MNIST)
LC-FTT (ours)	7,026 / 96.1	5,762 / 92.6	8,042* / 80.4*	7,968* / 80.1	6,853 [†] / 90.4 [†]	5,346 [†] / 87.4	7,921 / 79.2	7,424 [†] / 95.9	6,971 / 86.9*	954 [†] / 100.0*	714* / 73.6*
DMS	7,041 [†] / 96.7*	5,782* / 92.8*	8,042 [†] / 80.4*	7,968* / 80.1	6,431 / 89.3	5,129 / 88.2*	7,948 [†] / 79.5*	7,381 / 96.0*	6,993 [†] / 86.9*	953 / 100.0*	706 / 72.7
CMA-MAE	6,342 / 81.2	5,396 / 80.7	5,649 / 56.5	7,903 / 79.2	703 / 7.3	165 / 1.9	1,469 / 14.7	952 / 9.8	2,340 / 24.3	954 / 99.3	624 / 63.7
DDS	3,530 / 70.6	2,964 / 77.7	7,310 / 73.1	4,406 / 80.2*	3,732 / 56.6	2,918 / 57.1	4,822 / 48.2	6,993 / 84.4	6,366 / 74.3	N/A	N/A
MAP-Elites (line)	4,905 / 60.4	3,822 / 56.4	4,353 / 43.5	7,460 / 75.6	2,591 / 29.4	2,035 / 28.4	734 / 7.3	N/A	N/A	943 / 98.6	534 / 54.9
MAP-Elites	4,157 / 50.7	3,170 / 48.4	4,361 / 43.6	7,419 / 75.5	247 / 2.5	516 / 7.3	145 / 1.5	N/A	N/A	944 / 98.7	519 / 53.2

Baselines: We compare LC-FTT against five black-box QD baselines: Unregularized Discount Model Search (DMS) Tjanaka et al. (2026), Covariance Matrix Adaptation MAP-Annealing (CMA-MAE) Fontaine & Nikolaidis (2022), Density Descent Search (DDS) Lee et al. (2023), MAP-Elites (line), and vanilla MAP-Elites Mouret & Clune (2015). DDS is excluded on 784D TA domains due to KDE runtime intractability.

Metrics & Protocol: We measure QD Score (sum of objective values of all archive elites, normalized to $[0, 1]$ max per cell) and Coverage % (percentage of archive cells filled). All algorithms run for 10,000 iterations across 5 independent random seeds ([42, 123, 2024, 7, 88]) implemented in the pyribs framework. Mean values and standard errors of the mean (s.e.m.) are reported.

4.2 Main Comparative Results

Table 1 presents full-scale experimental results across all 11 benchmark domains. Appendix Figure 7 additionally illustrates normalized performance scaling from 2D to 784D.

Low-Dimensional Parity ($k = 2$): On 2D LP (Sphere), 2D LP (Rastrigin), 2D LP (Flat), and Arm Repertoire, LC-FTT achieves complete parity with unregularized DMS (e.g., 7,026 vs 7,041 on 2D Sphere; 5,762 vs 5,782 on 2D Rastrigin; exact equality on Flat and Arm Repertoire). Because $c(2) = 0$, LC-FTT automatically deactivates regularization in 2D, verifying that our dimension-adaptive rule avoids over-constraining dense low-dimensional search.

High-Dimensional Gains and Parity ($k \geq 10$): In measure spaces with $k \geq 10$, traditional grid-based baselines (CMA-MAE, MAP-Elites) suffer severe exploration collapse due to spatial distortion (e.g., CMA-MAE coverage drops to 7.3% on 10D LP Sphere and 1.9% on 10D LP Rastrigin). While DMS mitigates distortion using neural discount predictions, LC-FTT delivers its largest and only statistically significant gains at $k=10$, while maintaining numerical leads or ties without regression at $k \geq 20$ (see Section 4.3):

- 10D LP Sphere: LC-FTT achieves a QD score of 6,853 (+422 points over DMS’s 6,431) and coverage of 90.4% (+1.16% over DMS’s 89.3%).
- 10D LP Rastrigin: LC-FTT achieves a QD score of 5,346 (+217 points over DMS’s 5,129) on a highly multimodal landscape.
- 20D LP Sphere: LC-FTT leads all methods with a QD score of 7,424 (+43 points over DMS’s 7,381).
- 784D TA Fashion-MNIST: LC-FTT achieves a top QD score of 714 (+8.0 points over DMS’s 706) and coverage of 73.6% (+0.82% over DMS’s 72.7%).

Overall, LC-FTT attains the highest mean QD Score in 7 out of 11 domains and top coverage in 5 out of 11 domains. As we formalize in Section 4.3, the statistically significant advantages over DMS are concentrated at $k=10$, while at $k \geq 20$ (20D/50D/784D) LC-FTT maintains performance parity without ever regressing—precisely the intended behavior of the dimension-adaptive coefficient.




Figure 2: Formal statistical testing of LC-FTT versus unregularized DMS across 5 independent seeds. (A) Paired improvement $\Delta = \text{LC-FTT} - \text{DMS}$ in QD Score with 95% confidence intervals (Welch’s t-test); intervals excluding zero (blue) indicate significant gains. (B) Non-parametric Mann–Whitney U significance as $-\log_{10}(p)$, with $p = 0.05$ and $p = 0.01$ reference lines. On 10D LP Sphere ($\Delta = +422$) and 10D LP Rastrigin ($\Delta = +217$), LC-FTT attains the maximum non-parametric significance achievable at $n=5$ ($p_{\min}=0.0079$, rank-biserial $r=+1.00$). In extreme dimensions (20D, 50D, 784D) differences are numerically small and statistically indistinguishable from zero, confirming that the regularizer never harms high-dimensional optimization.

Table 2: Left: archive coverage (%) against modern grid-free / soft-competition QD baselines on high-dimensional LP Sphere. Fixed-population soft-density methods (DNS, SQUAD) collapse under measure-space distance concentration, whereas the continuous discount surface sustains near-complete coverage. Right: continuous-control transfer on the Brax physics platform (QD Score / Coverage %). Because these omnidirectional tasks have a 2D measure space, $c(2)=0$ and LC-FTT reduces identically to DMS (parity within seed variance); both continuous-discount methods beat grid-based CMA-MAE on Ant-Omni.

Method	10D LP (Sph)	20D LP (Sph)	Algorithm	Brax Ant-Omni	Brax Humanoid-Omni
LC-FTT (ours)	90.4	95.9	LC-FTT (ours)	1,135 / 15.7	213 / 5.6
DMS	89.3	96.0	DMS	1,197 / 16.6	212 / 5.6
DNS (Bahlous-Boldi et al., 2025)	22.0	53.0	CMA-MAE	1,066 / 14.8	250* / 6.6*
SQUAD (soft QD)	3.0	27.0			

Archive Coverage (%).

QD Score / Coverage (%). $n=584$ (Ant), $n=4209$ (Humanoid).

4.3 Statistical Significance, Modern Grid-Free Baselines, and Physical Control

Statistical significance testing. We ran Welch’s two-sample t-tests (95% CIs) and non-parametric Mann–Whitney U tests (rank-biserial r) on the paired difference $\Delta = \text{LC-FTT} - \text{DMS}$ across all 5 seeds for every domain (Figure 2; full statistics in Appendix Table 5). The gains on 10D LP Sphere ($\Delta = +422.0$ QD; MWU $p=0.0079$, $r=+1.00$) and 10D LP Rastrigin ($\Delta = +216.9$ QD; MWU $p=0.0079$, $r=+1.00$) are significant at $p < 0.01$ under both tests, at the maximum attainable effect size; 10D Sphere coverage is likewise significant ($\Delta = +1.16\%$, $p=0.0159$). At $k=2$ the penalty is off ($c(2)=0$), and in extreme dimensions (20D/50D/784D) differences are small and not significant—confirming the regularizer prevents exploration collapse without ever regressing high-dimensional performance.

Modern continuous / grid-free baselines. We benchmark against two recent grid-free paradigms: Dominated Novelty Search (DNS) (Bahlous-Boldi et al., 2025), which replaces grids with dynamic local-competition fitness transforms, and a Soft Quality-Diversity variant (SQUAD) with a fixed population under a soft kernel density. Both degrade sharply as k grows because fixed-population soft-density estimates suffer from distance concentration: DNS reaches only $\sim 22\%$ (10D) and $\sim 53\%$ (20D) coverage, and SQUAD $\sim 3\%$ and $\sim 27\%$, whereas LC-FTT sustains 90.4% and 95.9% (Table 2, left)—an explicit continuous discount surface therefore scales substantially better than dynamic distance/density-based competition in $k \geq 10$ measure spaces.




Figure 3: Continuous-control physical dynamics on Brax with an active regularizer ($c(10)=0.005$) over a 10D temporal state-occupancy descriptor (5 seeds; error bars are s.e.m.). (A) On the non-saturating Humanoid-10D task, LC-FTT accelerates early exploration, reaching +3.90 pp higher coverage at iteration 250 (Mann–Whitney U $p=0.004$) and converging to parity with DMS at full budget. (B) Final QD Score on both environments: LC-FTT is on par with DMS while CMA-MAE leads on the harder Humanoid task; Ant-10D saturates near ~98% coverage for all methods.

Continuous-control physical simulation (Brax). To test transfer beyond synthetic Linear Projection domains, we evaluate on the GPU-accelerated Brax physics platform with non-linear rigid-body dynamics. We first run two standard omnidirectional tasks with a 2D measure space, Ant-Omni (quadruped, $n=584$) and Humanoid-Omni (bipedal, $n=4209$), on a 100×100 archive using the scalable LM-MA-ES emitter (Table 2, right). Since $k=2$, the dimension-adaptive rule sets $c(2)=0$ and LC-FTT reduces identically to DMS, matching it within seed variance (1,135 vs. 1,197 QD on Ant; 213 vs. 212 on Humanoid); both continuous-discount methods outperform grid-based CMA-MAE on Ant-Omni.

Activating the regularizer on physical dynamics (10D descriptors). To verify that the regularizer is beneficial when it is active ($c(10)=0.005$) on non-linear physical dynamics, we introduce a 10-dimensional temporal state-occupancy behavioral descriptor $m(\theta) \in [-1, 1]^{10}$ on Brax Ant and Humanoid (5 seeds, 1,000-cell CVT archive). Figure 3 summarizes the results (full per-metric numbers in Appendix Table 6). On Humanoid-10D—whose archive does not saturate within budget—the active penalty accelerates early archive expansion: coverage at iteration 250 reaches 65.06% for LC-FTT versus 61.16% for DMS (+3.90 pp, Mann–Whitney U $p=0.0040$; coverage-AUC 66.14% vs. 64.19%), converging to statistical parity at full budget (72.88% vs. 72.52%, $p=0.726$). On the rapidly saturating Ant-10D, all methods reach ~98% coverage (QD parity within seed variance). This confirms that the active Dirichlet penalty suppresses high-frequency discount wiggles on complex contact dynamics, accelerating early exploration without harming final convergence, and directly addresses the concern that its efficacy on physical dynamics was previously untested at $k=2$.

Robustness to measure-scale anisotropy. Since $c(k)$ assumes coordinates normalized to $[-1, 1]^k$, we stress-test unnormalized/anisotropic measures on 10D LP Sphere by geometrically spreading per-dimension bounds up to scale ratio $R \in \{10, 1000\}$ (Appendix Table 8, Figure 4(B)). At $10 \times$ anisotropy, unnormalized LC-FTT matches its dynamically min-max-normalized variant (7,376.9 vs. 7,408.9 QD, $p=0.548$) and both beat unregularized DMS (+261 QD, $p<0.01$); at $1000 \times$, dynamic normalization restores a slight +1.1% edge. LC-FTT is thus robust to $10 \times$ scale disparity even without normalization.

4.4 Ablation Studies and Diagnostic Analyses

Ablation 1: Dimension-Adaptive Penalty $c(k)$ and Base Multiplier c_0 . We evaluate the sensitivity of LC-FTT to the base multiplier $c_0 \in \{0.0, 0.001, 0.003125, 0.00625, 0.0125, 0.025\}$ across 2D LP (Sphere), 10D LP (Sphere), 50D LP (Sphere), and 784D TA Fashion-MNIST (Appendix Figure 6(A)). For $k=2$, $c(2)=0$ irrespective of c_0 , preserving flat optimal performance. For $k \geq 10$, $c_0=0.0$ (unregularized DMS) suffers from interpolation wiggles.

Increasing c_0 steadily improves high-dimensional QD scores, peaking sharply at $c_0 = 0.00625$ across all higher dimensions before excessive smoothing ($c_0 \geq 0.0125$) oversmooths local discount ratchets.

Ablation 2: Monte Carlo Direction Count n_{dirs} and Estimator Variance. We vary the direction count $n_{\text{dirs}} \in \{1, 2, 4, 8, 16, 32, 64\}$ in the stochastic finite-difference estimator (Eq. 7). Appendix Figure 6(B) measures relative loss gradient variance and per-step wall-clock training time. Gradient variance decays along an exact $1/n_{\text{dirs}}$ Monte Carlo curve ($16\times$ variance reduction at $n_{\text{dirs}} = 16$). Crucially, setting $n_{\text{dirs}} = 16$ costs only 15.5 ms/step in 10D and 20.2 ms/step in 784D, providing an optimal balance between variance reduction and computational overhead.

Ablation 3: Comparison of Smoothness Regularization Formulations. We compare finite-difference L2 Dirichlet regularization against three alternative smoothing formulations: finite-difference L1 Total Variation, exact PyTorch Autograd L2 gradient penalty, and Layer Spectral Normalization Miyato et al. (2018). Appendix Figure 5(A) presents QD scores across 10D LP Sphere, 10D LP Rastrigin, and 20D LP Sphere, while Figure 5(B) shows the empirical gradient norm spectrum $\|\nabla_s \hat{f}_A(s)\|_2$ sampled across 4,000 measure points. Spectral Normalization achieves the tightest gradient spectrum (mean ~ 0.4) and excels on unimodal Sphere (7,162 on 10D), but collapses on rugged 10D Rastrigin (4,980 QD vs. 5,346 for L2 Dirichlet, worse than unregularized DMS’s 5,129) since global Lipschitz caps over-smooth multi-basin landscapes. L1 Total Variation induces flat plateaus: effective on 20D Sphere (7,798 QD) but degrading 10D Rastrigin coverage (83.28% vs. 87.42%) by erasing fine local discount differences. Exact Autograd L2 yields slightly lower QD than finite-difference L2 Dirichlet (6,590 vs. 6,853 on 10D Sphere), as stochastic direction sampling acts as an implicit spatial data-augmentation. To isolate whether this edge is merely spatial jittering, we additionally evaluate exact Autograd penalized at Gaussian-perturbed inputs $s + \eta$ ($\eta \sim \mathcal{N}(0, 0.05^2 I)$): input noise recovers only 9% of the finite-difference QD gap (+22.6 of +262.1 on Sphere, +4.5 of +51.2 on Rastrigin; Appendix Table 7 and Figure 4(A)), confirming that the advantage stems from structured hypersphere-radius probing with Monte Carlo direction averaging rather than input augmentation. Thus, finite-difference L2 Dirichlet energy is uniquely robust across both unimodal and multimodal landscapes.

Ablation 4: Sensitivity to Step Size ϵ and Numerical Stability. Sweeping $\epsilon \in \{10^{-5}, \dots, 10^{-1}\}$ (and scale-adaptive stepping) on 10D Sphere, 10D Rastrigin, and 784D TA MNIST (Appendix Figure 8), LC-FTT records zero non-finite (NaN/Inf) losses across all 100,000+ training batches per run; $\epsilon = 0.05$ achieves peak QD, avoiding truncation error at $\epsilon \leq 10^{-4}$ and non-local blurring at $\epsilon = 0.1$.

5 Conclusion

We identified a vulnerability in continuous surrogate-assisted Quality-Diversity: in high-dimensional measure spaces ($k \geq 10$), sparse emitter sampling causes unregularized neural discount models to fit spurious high-frequency wiggles that misguide emitters and stall behavioral expansion. We introduced Legendrian Contact-Geometric Discount Optimization (LC-FTT), a stochastic finite-difference Dirichlet-energy regularizer governed by a dimension-adaptive factor $c(k) = c_0 \max(0, 1 - 2/k)$ that deactivates in 2D ($c(2)=0$) and scales up in higher dimensions. Across 11 benchmark domains (up to 784D), LC-FTT matches DMS in 2D, delivers statistically significant gains at $k=10$, maintains parity without regression at $k \geq 20$, and accelerates early exploration when the regularizer is active on 10D Brax physical control.

Limitations & Future Work: Computing nearest-neighbor centroid assignments in massive CVT archives ($> 100,000$ cells) remains a computational bottleneck; future work will explore GPU-accelerated approximate nearest-neighbor search (e.g., FAISS) and extend contact-geometric discount regularization to generative latent spaces and transformer-based measure surrogates.

References

- S. Akansha. Over-squashing in graph neural networks: A comprehensive survey. *Neuro-computing*, 642:130389, 2023.
- Ryan Bahlous-Boldi, Maxence Faldor, Luca Grillotti, Hannah Janmohamed, Lisa Coiffard, Lee Spector, and Antoine Cully. Dominated novelty search: Rethinking local competition in quality-diversity. *Proceedings of the Genetic and Evolutionary Computation Conference*, 2025.
- Varun Bhatt, Bryon Tjanaka, Matthew C. Fontaine, and S. Nikolaidis. Deep surrogate assisted generation of environments. *ArXiv*, abs/2206.04199, 2022.
- Varun Bhatt, Heramb Nemlekar, Matthew C. Fontaine, Bryon Tjanaka, Hejia Zhang, Ya-Chuan Hsu, and S. Nikolaidis. Surrogate assisted generation of human-robot interaction scenarios. In *Conference on Robot Learning*, pp. 513–539, 2023.
- Chen Cai and Yusu Wang. A note on over-smoothing for graph neural networks. *ArXiv*, abs/2006.13318, 2020.
- M. Charikar and Paris Siminelakis. Hashing-based-estimators for kernel density in high dimensions. *2017 IEEE 58th Annual Symposium on Foundations of Computer Science (FOCS)*, pp. 1032–1043, 2017.
- Konstantinos Chatzilygeroudis, Antoine Cully, Vassilis Vassiliades, and Jean-Baptiste Mouret. Quality-diversity optimization: a novel branch of stochastic optimization. *ArXiv*, abs/2012.04322, 2020.
- Zhengdao Chen, Jianyu Zhang, Martín Arjovsky, and L. Bottou. Symplectic recurrent neural networks. *ArXiv*, abs/1909.13334, 2019.
- Cédric Colas, Joost Huizinga, Vashisht Madhavan, and J. Clune. Scaling map-elites to deep neuroevolution. *Proceedings of the 2020 Genetic and Evolutionary Computation Conference*, 2020.
- Antoine Cully. Autonomous skill discovery with quality-diversity and unsupervised descriptors. *Proceedings of the Genetic and Evolutionary Computation Conference*, 2019.
- Antoine Cully, J. Clune, Danesh S. Tarapore, and Jean-Baptiste Mouret. Robots that can adapt like animals. *Nature*, 521:503–507, 2014.
- Danijel Djordjević, Igor Uljarević, and Jun Zhang. Quantitative contact hamiltonian dynamics. In *arXiv*, 2025.
- Matthew C. Fontaine and S. Nikolaidis. Differentiable quality diversity. *ArXiv*, abs/2106.03894, 2021.
- Matthew C. Fontaine and S. Nikolaidis. Covariance matrix adaptation map-annealing. *Proceedings of the Genetic and Evolutionary Computation Conference*, 2022.
- Matthew C. Fontaine, J. Togelius, S. Nikolaidis, and Amy K. Hoover. Covariance matrix adaptation for the rapid illumination of behavior space. *Proceedings of the 2020 Genetic and Evolutionary Computation Conference*, 2019.
- Adam Gaier, A. Asteroth, and Jean-Baptiste Mouret. Data-efficient design exploration through surrogate-assisted illumination. *Evolutionary Computation*, 26:381–410, 2018.
- Ishaan Gulrajani, Faruk Ahmed, Martín Arjovsky, Vincent Dumoulin, and Aaron C. Courville. Improved training of wasserstein gans. In *Neural Information Processing Systems*, pp. 5767–5777, 2017.
- Saeed Hedayatian and S. Nikolaidis. Autoqd: Automatic discovery of diverse behaviors with quality-diversity optimization. *ArXiv*, abs/2506.05634, 2025.

- Paul Kent and J. Branke. Bop-elites, a bayesian optimisation algorithm for quality-diversity search. ArXiv, abs/2005.04320, 2020.
- David H. Lee, Anisha Palaparthi, Matthew C. Fontaine, Bryon Tjanaka, and S. Nikolaidis. Density descent for diversity optimization. Proceedings of the Genetic and Evolutionary Computation Conference, 2023.
- Takeru Miyato, Toshiki Kataoka, Masanori Koyama, and Yuichi Yoshida. Spectral normalization for generative adversarial networks. ArXiv, abs/1802.05957, 2018.
- Jean-Baptiste Mouret and J. Clune. Illuminating search spaces by mapping elites. ArXiv, abs/1504.04909, 2015.
- Olle Nilsson and Antoine Cully. Policy gradient assisted map-elites. Proceedings of the Genetic and Evolutionary Computation Conference, 2021.
- Justin K. Pugh, L. Soros, and Kenneth O. Stanley. Quality diversity: A new frontier for evolutionary computation. Frontiers Robotics AI, 3:40, 2016.
- Nasim Rahaman, A. Baratin, Devansh Arpit, Felix Dräxler, Min Lin, F. Hamprecht, Yoshua Bengio, and Aaron C. Courville. On the spectral bias of neural networks. In International Conference on Machine Learning, pp. 5301–5310, 2018.
- Daniel Soudry, Elad Hoffer, Suriya Gunasekar, and N. Srebro. The implicit bias of gradient descent on separable data. ArXiv, abs/1710.10345, 2017.
- Bryon Tjanaka, Matthew C. Fontaine, David H. Lee, Aniruddha Kalkar, and S. Nikolaidis. Training diverse high-dimensional controllers by scaling covariance matrix adaptation map-annealing. IEEE Robotics and Automation Letters, 8:6771–6778, 2022a.
- Bryon Tjanaka, Matthew C. Fontaine, J. Togelius, and S. Nikolaidis. Approximating gradients for differentiable quality diversity in reinforcement learning. Proceedings of the Genetic and Evolutionary Computation Conference, 2022b.
- Bryon Tjanaka, Henry Chen, Matthew C. Fontaine, and S. Nikolaidis. Discount model search for quality diversity optimization in high-dimensional measure spaces. ArXiv, abs/2601.01082, 2026.
- Constantinos G. Tsakonas and Konstantinos I. Chatzilygeroudis. Vector quantized-elites: Unsupervised and problem-agnostic quality-diversity optimization. IEEE Transactions on Evolutionary Computation, 30:1669–1683, 2025.
- Vassilis Vassiliades, Konstantinos Chatzilygeroudis, and Jean-Baptiste Mouret. Using centroidal voronoi tessellations to scale up the multidimensional archive of phenotypic elites algorithm. IEEE Transactions on Evolutionary Computation, 22:623–630, 2016.
- Shihan Zhao, Bryon Tjanaka, Matthew C. Fontaine, and S. Nikolaidis. Covariance matrix adaptation map-annealing: Theory and experiments. ACM Transactions on Evolutionary Learning, 5:1 – 53, 2024.
- Kaixiong Zhou, Xiao Huang, D. Zha, Rui Chen, Li Li, Soo-Hyun Choi, and Xia Hu. Dirichlet energy constrained learning for deep graph neural networks. In Neural Information Processing Systems, pp. 21834–21846, 2021.

A Benchmark Domain Specifications and Hyperparameters

A.1 Mathematical Definitions of Benchmark Domains

Linear Projection (LP) Domains: The LP domain maps an n -dimensional solution vector $\theta \in [-5.12, 5.12]^n$ to a k -dimensional measure vector $s = m(\theta) \in [-5.12r, 5.12r]^k$, where $r = n/k$. The solution vector is partitioned into k contiguous blocks of size r . The measure function is defined as:

$$m(\theta)_j = \sum_{i=(j-1)r+1}^{jr} \text{clip}(\theta_i), \quad \text{where } \text{clip}(\theta_i) = \begin{cases} \theta_i & \text{if } |\theta_i| \leq 5.12, \\ 5.12/\theta_i & \text{otherwise.} \end{cases} \quad (8)$$

We evaluate three objective functions $f(\theta)$, each normalized to a maximum value of 1.0 by the constant $f_{\max}$: a unimodal Sphere objective, a highly multimodal Rastrigin objective, and a Flat objective for pure diversity optimization,

$$f_{\text{Sphere}}(\theta) = 1.0 - \frac{1}{f_{\max}} \sum_{i=1}^n (\theta_i - 2.048)^2, \quad f_{\text{Flat}}(\theta) = 1.0, \quad (9)$$

$$f_{\text{Rastrigin}}(\theta) = 1.0 - \frac{1}{f_{\max}} \left[10n + \sum_{i=1}^n ((\theta_i - 2.048)^2 - 10 \cos(2\pi(\theta_i - 2.048))) \right]. \quad (10)$$

Arm Repertoire Domain: Considers an n -link planar arm with joint angles $\theta \in [-\pi, \pi]^n$ ($n = 100$). The 2D measure function $m(\theta) = (x, y)$ computes end-effector position via forward kinematics. The objective minimizes joint variance: $f(\theta) = 1.0 - \text{Var}(\theta)$.

Triangle Arrangement (TA) Domains: Generates grayscale images composed of 30 overlapping transparent triangles ($n = 30 \times 8 = 240$ parameters). Rendered measures $m(\theta) \in [0, 1]^{784}$ form a 28×28 pixel image. Archive centroids consist of 1,000 target images from MNIST or Fashion-MNIST. The objective maximizes negative mean squared error relative to the nearest target centroid: $f(\theta) = 1.0 - \text{MSE}(m(\theta), s_{\text{centroid}})$.

A.2 Hyperparameter Configurations

Table 3 details the hyperparameter configurations across all algorithms and benchmark domains.

B Tabulated Per-Seed Raw Numerical Data

Table 4 lists the final raw Quality-Diversity Score and Archive Coverage (%) across all 5 independent random seeds ([42, 123, 2024, 7, 88]) for LC-FTT and DMS across high-dimensional benchmark domains.

C Full Statistical Significance Testing

Table 5 reports the complete statistical analysis of the paired difference $\Delta = \text{LC-FTT} - \text{DMS}$ across all 5 independent seeds ([42, 123, 2024, 7, 88]) for every benchmark domain, summarized visually in Figure 2. We report Welch’s two-sample t-test (unequal variance) with 95% confidence intervals, and the non-parametric Mann–Whitney U (MWU) test with rank-biserial correlation effect sizes r . Note that for a two-sided MWU test with $n = 5$ per group, the smallest attainable p -value is 0.0079 ($U = 25$); LC-FTT achieves this floor on both 10D domains.

Interpretation. The improvements on 10D LP Sphere and 10D LP Rastrigin are statistically significant at $p < 0.01$ under both parametric and non-parametric tests, with the maximum attainable effect size ($r = +1.00$; LC-FTT strictly dominates DMS on every seed). At $k = 2$ the penalty deactivates ($c(2) = 0$), so differences are pure RNG variance. In the highest dimensions (20D/50D/784D), differences are numerically small and not statistically significant, confirming that the dimension-adaptive regularizer prevents high-dimensional exploration collapse without ever regressing performance.

Table 3: Hyperparameter settings for LC-FTT, DMS, CMA-MAE, DDS, MAP-Elites (line), and MAP-Elites.

Parameter	Value / Configuration
LC-FTT & DMS Common Hyperparameters	
Discount Model Architecture	3-layer MLP $[k, 128, 128, 1]$ with ReLU activations
Optimizer	Adam ($\text{lr} = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, batch size 32)
Target MSE Cutoff	0.05 average MSE over dataset D_A (max 5 epochs/iter)
Empty Point Samples n_{empty}	100 cell centroids sampled without replacement
Initial Regression Points n_{init}	1,000 points initialized to $f_{\min} = 0.0$
Archive Learning Rate α	0.1 (LP 2D/10D/20D/50D, TA F-MNIST), 0.001 (Arm, TA MNIST)
LC-FTT Specific Contact-Form Regularization Hyperparameters	
Base Penalty Multiplier c_0	0.00625
Dimension Scaling $c(k)$	$c(k) = 0.00625 \cdot \max(0, 1 - 2/k)$
Directional Batch Size n_{dirs}	16 Monte Carlo unit direction vectors per sample
Finite-Difference Step Size ϵ	0.05 (in $[-1, 1]$ normalized measure coordinates)
Search Emitter Hyperparameters (CMA-ES based: LC-FTT, DMS, CMA-MAE)	
Number of Emitters W	15 (LP / Arm), 5 (TA)
Emitter Batch Size λ	36
Initial Step Size σ_0	0.5 (LP), 0.2 (Arm), 0.1 (TA)
Restart Rule	Basic (2D LP Sphere/Rastrigin/Flat, Arm), 100 iters (10D/20D/50D LP), 50 iters (TA)
DDS Specific Hyperparameters	
Bandwidth h	25.6 (2D LP), 5.12 (10D/20D/50D LP), 10.0 (Arm)
Buffer Size	10,000 history solutions
Selection / Search Rule	Filter selection with CMA-ES search
MAP-Elites (line) & MAP-Elites Hyperparameters	
Batch Size λ	540 (LP / Arm), 180 (TA)
Line Iso Sigma σ_1 / Ind Sigma σ_2	$\sigma_1 = 0.5, \sigma_2 = 0.2$ (LP / Arm); $\sigma_1 = 0.1, \sigma_2 = 0.2$ (TA)
Vanilla Mutate Sigma σ	0.5 (LP / Arm), 0.1 (TA)

D Additional Physical-Control, Gradient-Estimator, and Anisotropy Experiments

This appendix provides the full per-metric numbers for the three additional experiments summarized in Section 4.3: (i) 10D Brax continuous physical control with an active regularizer, (ii) a finite-difference vs. Autograd data-augmentation ablation, and (iii) an anisotropic measure-space stress test. Figure 4 visualizes experiments (ii) and (iii).

(i) 10D Brax continuous control. Table 6 reports QD Score, final coverage, and checkpoint coverage at iteration 250 on Brax Ant and Humanoid with a 10-dimensional temporal state-occupancy descriptor ($c(10)=0.005$; 5 seeds, 1,000-cell CVT archive, 540,000 evaluations per run). On the non-saturating Humanoid-10D task, LC-FTT reaches +3.90 pp higher coverage than DMS at iteration 250 (Mann-Whitney U $p=0.0040$) and converges to parity at full budget; Ant-10D saturates near $\sim 98\%$ for all methods.

(ii) Finite-difference vs. Autograd with input noise. Table 7 compares the shipped finite-difference (FD) estimator against exact Autograd penalized at unperturbed inputs and at Gaussian-perturbed inputs $s+\eta$, $\eta \sim \mathcal{N}(0, 0.05^2 I)$. Adding input noise to Autograd recovers only 9% of the FD advantage, indicating the FD edge arises from structured hypersphere-radius probing rather than spatial data augmentation.

(iii) Anisotropic measure spaces. Table 8 reports a stress test on 10D LP Sphere in which per-dimension measure bounds are geometrically spread up to a scale ratio R . Unnormalized LC-FTT matches its dynamically normalized variant and beats DMS at $10\times$ anisotropy; dynamic min-max normalization restores a slight edge at extreme $1000\times$ anisotropy.

E Additional Figures

Table 4: Per-seed raw Quality-Diversity Score and Archive Coverage (%) for LC-FTT (ours) and DMS across high-dimensional benchmark domains after 10,000 iterations.

Domain	Algorithm	Seed	QD Score	Coverage (%)
10D LP (Sphere)	LC-FTT (ours)	42	6,878.59	90.43
10D LP (Sphere)	LC-FTT (ours)	123	6,795.45	90.28
10D LP (Sphere)	LC-FTT (ours)	2024	6,889.38	90.63
10D LP (Sphere)	LC-FTT (ours)	7	6,900.52	91.05
10D LP (Sphere)	LC-FTT (ours)	88	6,798.98	89.71
10D LP (Sphere)	DMS	42	6,342.21	88.35
10D LP (Sphere)	DMS	123	6,436.72	89.38
10D LP (Sphere)	DMS	2024	6,544.72	90.24
10D LP (Sphere)	DMS	7	6,434.72	89.41
10D LP (Sphere)	DMS	88	6,394.47	88.93
10D LP (Rastrigin)	LC-FTT (ours)	42	5,335.76	87.18
10D LP (Rastrigin)	LC-FTT (ours)	123	5,317.53	86.94
10D LP (Rastrigin)	LC-FTT (ours)	2024	5,393.63	88.10
10D LP (Rastrigin)	LC-FTT (ours)	7	5,360.69	87.83
10D LP (Rastrigin)	LC-FTT (ours)	88	5,322.71	87.05
10D LP (Rastrigin)	DMS	42	5,030.49	86.70
10D LP (Rastrigin)	DMS	123	5,179.59	88.77
10D LP (Rastrigin)	DMS	2024	5,169.28	88.31
10D LP (Rastrigin)	DMS	7	5,156.41	88.87
10D LP (Rastrigin)	DMS	88	5,110.18	88.21
TA (F-MNIST)	LC-FTT (ours)	42	695.00	71.60
TA (F-MNIST)	LC-FTT (ours)	123	737.27	76.00
TA (F-MNIST)	LC-FTT (ours)	2024	697.81	71.90
TA (F-MNIST)	LC-FTT (ours)	7	734.22	75.70
TA (F-MNIST)	LC-FTT (ours)	88	704.97	72.60
TA (F-MNIST)	DMS	42	702.59	72.40
TA (F-MNIST)	DMS	123	698.96	72.00
TA (F-MNIST)	DMS	2024	715.53	73.80
TA (F-MNIST)	DMS	7	705.03	72.70
TA (F-MNIST)	DMS	88	707.28	72.80

Table 5: Full statistical significance testing of LC-FTT vs. DMS across 5 seeds. $\Delta = \text{LC-FTT} - \text{DMS}$. Significance: ** $p < 0.01$, * $p < 0.05$, ns = not significant.

Domain	Metric	Δ	95% CI	Welch t	Welch p	MWU p (r)	Sig.
10D LP (Sphere)	QD Score	+422.02	[+326.70, +517.34]	10.47	$< 10^{-4}$	0.0079 (+1.00)	**
10D LP (Sphere)	Coverage %	+1.16	[+0.26, +2.05]	3.03	0.0182	0.0159 (+0.92)	*
10D LP (Rastrigin)	QD Score	+216.87	[+141.49, +292.26]	7.04	0.0004	0.0079 (+1.00)	**
20D LP (Sphere)	QD Score	+42.92	[-120.76, +206.61]	0.65	0.5429	0.6905	ns
50D LP (Sphere)	QD Score	-21.57	[-71.44, +28.31]	-1.00	0.3447	0.2222	ns
784D TA (F-MNIST)	QD Score	+7.98	[-16.88, +32.84]	0.80	0.4420	1.0000	ns
2D domains (Sph/Rast/Flat/Arm)	QD Score	≈ 0	—	—	[0.41, 1.00]	—	ns

Table 6: 10D Brax continuous control (mean $\pm$ s.e.m. over 5 seeds). The regularizer is active ($c(10)=0.005$). * marks the best value per row group.

Environment	Metric	LC-FTT (ours)	DMS	CMA-MAE
Brax Ant-10D	Final QD Score	386.7 $\pm$ 20.3*	385.2 $\pm$ 20.8	382.1 $\pm$ 20.2
	Final Coverage %	98.44 $\pm$ 1.16	98.58 $\pm$ 1.12*	97.70 $\pm$ 1.57
	Coverage % @ Iter 250	96.70 $\pm$ 2.11*	96.70 $\pm$ 2.12	96.40 $\pm$ 2.08
Brax Humanoid-10D	Final QD Score	190.2 $\pm$ 5.1	197.8 $\pm$ 8.0	239.0 $\pm$ 3.2*
	Final Coverage %	72.88 $\pm$ 0.98	72.52 $\pm$ 1.62	84.12 $\pm$ 0.83*
	Coverage % @ Iter 250	65.06 $\pm$ 0.75	61.16 $\pm$ 0.99	74.16 $\pm$ 0.76*


Humanoid-10D iteration-250 coverage: LC-FTT vs. DMS is +3.90 pp (Mann-Whitney U $p=0.0040$; coverage-AUC 66.14% vs. 64.19%); full-budget difference is not significant ($p=0.726$).

Table 7: Spatial data-augmentation ablation on 10D LP (mean, s.e.m. over 5 seeds, 10,000 iterations). Δ QD is measured relative to the shipped finite-difference estimator.

Benchmark	Variant	QD Score	Coverage %	Δ QD vs. FD
10D LP (Sphere)	FD (ours, $\epsilon=0.05$, $n_{\text{dirs}}=16$)	6,852.6 (22.9)	90.42 (0.22)	—
	Autograd + input noise ($\sigma=0.05$)	6,613.2 (26.3)	91.46 (0.23)	−239.4
	Autograd (exact, no noise)	6,590.5 (37.9)	91.22 (0.40)	−262.1
10D LP (Rastrigin)	FD (ours, $\epsilon=0.05$, $n_{\text{dirs}}=16$)	5,346.1 (14.0)	87.42 (0.23)	—
	Autograd + input noise ($\sigma=0.05$)	5,299.5 (30.1)	89.80 (0.43)	−46.6
	Autograd (exact, no noise)	5,294.9 (18.6)	89.85 (0.20)	−51.2

Table 8: Anisotropic / unnormalized measure-space stress test on 10D LP Sphere (mean $\pm$ s.e.m. over 5 seeds). p -values are Mann–Whitney U vs. DMS.

Anisotropy ratio R	Condition	QD Score	Coverage %	MWU p (vs. DMS)
$R = 10$ ($10\times$ spread)	LC-FTT (unnormalized)	$7,376.9 \pm 22.7$	96.53 ± 0.11	0.0079 ($r=+1.00$)
	LC-FTT (dynamic min-max)	$7,408.9 \pm 18.9$	94.17 ± 0.34	0.0079 ($r=+1.00$)
	DMS (baseline)	$7,115.4 \pm 32.3$	95.20 ± 0.28	—
$R = 1000$ ($1000\times$ spread)	LC-FTT (unnormalized)	$8,344.5 \pm 21.7$	99.95 ± 0.01	0.0317 ($r=-0.84$)
	LC-FTT (dynamic min-max)	$8,436.6 \pm 14.9$	99.02 ± 0.05	0.3095 ($r=+0.44$)
	DMS (baseline)	$8,406.2 \pm 5.4$	99.59 ± 0.02	—

Figure 4: (A) Gradient-estimator ablation on 10D LP: exact Autograd with isotropic Gaussian input noise ($\sigma=0.05$) recovers only $\sim 9\%$ of the finite-difference (FD) QD advantage, so the FD edge is not explained by spatial input augmentation. (B) Anisotropic measure-space stress test on 10D LP Sphere: unnormalized LC-FTT is robust to $10\times$ coordinate-scale spread and retains a significant lead over DMS, while dynamic min-max normalization neutralizes extreme $1000\times$ anisotropy. Error bars are s.e.m. over 5 seeds.




Figure 5: Smoothness-regularization formulations and learned discount gradient-norm spectra (Ablation 3). (A) Mean QD Score (error bars: s.e.m.; points: per-seed) for unregularized DMS, L2 Dirichlet FD (ours), L1 Total-Variation FD, exact Autograd L2, and spectral normalization on 10D LP Sphere/Rastrigin and 20D LP Sphere. (B) Empirical distribution of learned gradient norms $\|\nabla_s \hat{f}_A(s)\|_2$ over 4,000 measure points (median–90th box, 90th–99th whisker, mean marker). L2 Dirichlet FD is uniquely robust across smooth and rugged multi-basin landscapes, whereas spectral norm and L1-TV over-smooth rugged domains despite tighter spectra on the unimodal Sphere.




Figure 6: Ablation analysis on the base penalty multiplier c_0 and Monte Carlo direction count n_{dirs} . (A) Sensitivity of normalized Quality-Diversity (QD) Score to the base multiplier c_0 across measure space dimensions k . At $k = 2$, performance remains constant because $c(2) = 0$, avoiding over-constraint in low dimensions, whereas high-dimensional search ($k = 10, 50, 784$) consistently peaks at $c_0 = 0.00625$. (B) Relative gradient variance of the discount model loss (solid blue) and per-step wall-clock training time (dashed orange) versus direction count n_{dirs} for $k = 10$ and $k = 784$. Variance drops following a $1/n_{\text{dirs}}$ Monte Carlo decay, providing a 16-fold variance reduction at the default $n_{\text{dirs}} = 16$ with modest compute cost (~ 15 to 20 ms/step).




Figure 7: Comparison of quality-diversity performance across measure space dimensions ranging from 2D to 784D. Subplot A displays normalized QD score as a percentage of the maximum possible archive score, while Subplot B shows archive coverage percentage. LC-FTT matches baseline performance in 2D and outperforms existing methods in higher-dimensional measure spaces, avoiding the severe exploration collapse experienced by traditional algorithms in dimensions 10D and above. Error bars denote the standard error of the mean over five independent random seeds, and N/A labels indicate non-runable algorithm-domain configurations.




Figure 8: Sensitivity analysis of the finite-difference perturbation step size ϵ across 10D LP (Sphere), 10D LP (Rastrigin), and TA (MNIST) domains. Mean QD Score (shaded error bands indicate ± 1 standard error of the mean across five seeds) is plotted against perturbation step sizes ranging from 10^{-5} to 10^{-1} on a logarithmic scale. The default parameter value $\epsilon = 0.05$ (orange star) achieves optimal or near-optimal performance across all settings. Across all tested step sizes and dimensions, LC-FTT achieves complete numerical stability with zero non-finite losses while maintaining superior or competitive QD Scores compared to the unregularized DMS baseline (red dashed line).