

CONTINUOUS OPTIMAL-TRANSPORT FLOW MATCHING & SPECTRAL INFORMATION-BOTTLENECK FOR ANTI-CAUSAL REPRESENTATION LEARNING

Anonymous authors

Paper under double-blind review

ABSTRACT

Anti-causal representation learning ($Y \rightarrow X \leftarrow E$) requires discovering invariant generative mechanisms from target labels Y to observations X while suppressing non-stationary environmental variations E . Existing measure-theoretic and invariant risk frameworks suffer from critical computational bottlenecks: quadratic domain pair scaling $O(|E|^2)$, strict dependencies on discrete environment partitions, heavy wall-clock overhead from multi-step numerical Ordinary Differential Equation (ODE) solvers during counterfactual transport, and severe regularizer-task gradient conflicts that degrade predictive capacity. In this paper, we introduce Continuous Optimal-Transport Flow Matching & Spectral Information-Bottleneck (COFM-SIB), a scalable framework for continuous anti-causal representation learning. COFM-SIB eliminates simulation drag by learning a conditional neural velocity field trained via linear optimal transport displacement interpolation, enabling single-step deterministic counterfactual transport $do(Y = y')$ without numerical ODE integration or oracle interventional datasets. To achieve domain-agnostic environment independence, COFM-SIB incorporates Sobolev-Weighted Sliced-Wasserstein Alignment ($R_{1,\text{CSWS}}$) in $O(N \log N)$ complexity, alongside a linear-time Random Fourier Feature Hilbert-Schmidt Independence Criterion (R_{HSIC}) spectral bottleneck to compress non-causal environment style noise. Furthermore, Task-Shielded Dual Gradient Projection (TSGP) orthogonally projects conflicting regularizer gradients away from the task gradient while applying magnitude capping, preventing regularizer pressure from sacrificing predictive accuracy. Across four benchmark suites—Colored MNIST, Rotated MNIST, Ball Agent spatial regression, and Camelyon17 pathology slide classification—COFM-SIB matches or outperforms state-of-the-art baselines, achieving 99.20% accuracy on CMNIST, 99.19% on RMNIST, 99.32% position accuracy on Ball Agent, and 88.65% hospital-shift accuracy on Camelyon17 while offering up to $6.7\times$ speedups over multi-step ODE solvers.

1 INTRODUCTION

Out-of-distribution (OOD) generalization remains a central challenge in machine learning, particularly when models are deployed in non-stationary environments where input covariates undergo distribution shifts. In many real-world domains—ranging from medical pathology diagnosis de Castro et al. (2019); Bándi et al. (2019) and physical motion tracking Ghifary et al. (2015) to complex visual recognition under shifting environmental conditions—the data generation process exhibits an *anti-causal* structure ($Y \rightarrow X \leftarrow E$) Schölkopf et al. (2012; 2021b;a); Kilbertus et al. (2018). In this setting, the underlying target variable Y acts as a cause of the observed features X , while environmental factors E simultaneously influence X by injecting style, rotation, or domain-specific artifacts. To achieve robust OOD prediction under severe non-stationary domain shifts, representation learning algorithms must uncover the invariant generative mechanism $P(X | Y)$ while suppressing non-causal environment noise $P(X | E)$.

Recent measure-theoretic anti-causal representation frameworks, such as Anti-Causal Invariant Abstractions (ACIA) Behnam & Wang (2025) and early anti-causal domain shift methods Jiang & Veitch (2022), have established rigorous theoretical foundations for learning invariant representa-

tions via interventional kernels. However, existing anti-causal frameworks and conventional invariant learning algorithms suffer from four severe computational and structural bottlenecks:

1. **Discrete Domain Dependencies:** Standard invariant risk minimization and measure-theoretic approaches Arjovsky et al. (2019); Lu et al. (2021) critically rely on explicit, discrete domain annotations ($E \in \{e_1, \dots, e_K\}$). When environmental shifts evolve continuously (e.g., temporal trends or spatial trajectories), these methods require heuristic domain clustering Creager et al. (2020) or incur severe discretization errors Wang et al. (2020); Lin et al. (2023); Majumdar (2025).
2. **Quadratic Pairwise Alignment Scaling:** Aligning feature distributions or calculating interventional kernel discrepancies across environment pairs requires $O(|E|^2)$ computational complexity Ramé et al. (2021); Ahuja et al. (2020b), rendering scaling to continuous domain index sets or large environment cohorts computationally prohibitive.
3. **Numerical Solver Overhead & Oracle Interventional Data Reliance:** Evaluating interventional consistency R_2 under counterfactual interventions $do(Y = y')$ traditionally requires oracle interventional distributions Ahuja et al. (2022) or computationally expensive multi-step numerical Ordinary Differential Equation (ODE) integrators Onken et al. (2020); Gu et al. (2025), introducing massive wall-clock simulation drag during training.
4. **Capacity Over-Constraining & Gradient Interference:** Enforcing rigid invariance penalties alongside task prediction objectives frequently induces severe gradient conflicts ($\langle g_{\text{reg}}, g_{\text{task}} \rangle < 0$), over-constraining network capacity and degrading downstream predictive accuracy Gulrajani & Lopez-Paz (2020); Kamath et al. (2021); Yu et al. (2020); Zhang et al. (2024).

To address these fundamental limitations, we propose **Continuous Optimal-Transport Flow Matching & Spectral Information-Bottleneck (COFM-SIB)**, a novel framework for anti-causal representation learning. COFM-SIB bypasses multi-step ODE solver integration, pairwise discrete domain alignment, and gradient conflict degradation through four unified computational mechanisms:

- **Deterministic 1-Step Optimal-Transport Flow Matching ($R_{2,\text{OTFM}}$):** We parameterize latent counterfactual shifts $do(Y = y')$ using Continuous Optimal-Transport Flow Matching (OT-CFM) Lipman et al. (2022); Liu et al. (2022); Albergo & Vanden-Eijnden (2022). By regressing a conditional velocity field $v_\theta(Z_L, Y, y')$ on straight-line displacement paths Kim et al. (2024); Shou (2026), COFM-SIB evaluates counterfactual representations in a single deterministic pass without relying on multi-step numerical ODE integrators or oracle interventional datasets.
- **Sobolev-Weighted Sliced-Wasserstein Alignment ($R_{1,\text{CSWS}}$):** To enforce continuous label-conditional environment independence $Z_H \perp E \mid Y$, we project high-level representations onto 1D random hyperplanes weighted by target derivative smoothness ($1 + \|\nabla_Y Z_H\|_2^2$). Computed in $O(N \log N)$ complexity via 1D sorting, $R_{1,\text{CSWS}}$ natively aligns continuous and discrete domains without requiring explicit domain index partitions.
- **Kernelized Conditional HSIC Spectral Bottleneck (R_{HSIC}):** To compress non-causal environment noise without heuristic variational bounds, we incorporate a Hilbert-Schmidt Independence Criterion (HSIC) bottleneck utilizing Random Fourier Feature (RFF) cross-covariance approximations Tóth et al. (2023); Kalinke & Gavioli-Akilagun (2025); Huang (2026), measuring conditional independence in $O(N)$ linear complexity.
- **Task-Shielded Dual Gradient Projection (TSGP):** To prevent regularizers from compromising predictive capacity, TSGP orthogonally projects regularizer gradients away from conflicting task gradients whenever $\langle g_{\text{reg}}, g_{\text{task}} \rangle < 0$, combined with task-scale magnitude thresholding to preserve task performance while maximizing invariance.

We evaluate COFM-SIB across four major anti-causal benchmark suites: Colored MNIST (CMNIST), Rotated MNIST (RMNIST) Ghifary et al. (2015), Ball Agent continuous spatial regression, and Camelyon17 pathology slide classification Bándi et al. (2019) using deep ResNet-18 backbones He et al. (2015). COFM-SIB matches or outperforms prior baselines—including Group DRO Sagawa et al. (2019), Deep CORAL Sun & Saenko (2016), IRM Arjovsky et al. (2019), V-REx Krueger et al. (2020), ACTIR Jiang & Veitch (2022), LECI Gui et al. (2023), and ACIA Behnam

& Wang (2025)—across accuracy and invariance metrics. Specifically, COFM-SIB reaches 99.20% accuracy on CMNIST (matching ACIA while reducing environment independence leakage to 0.001), 99.19% on RMNIST, 99.32% on Ball Agent (position MAE 0.021), and 88.65% on Camelyon17 (substantially outperforming ACIA’s 84.40%). Comprehensive ablation studies confirm that 1-step OT-CFM transport achieves up to $6.7\times$ speedups over multi-step ODE solvers, RFF approximations reduce memory footprint by $4.8\times$ over exact Gram matrices, and TSGP prevents gradient conflicts from degrading task accuracy.

2 RELATED WORK

2.1 DOMAIN GENERALIZATION AND INVARIANT REPRESENTATION LEARNING

Domain generalization aims to train statistical models on source domains such that they generalize effectively to unseen target environments Gulrajani & Lopez-Paz (2020). Classical domain adaptation and alignment paradigms enforce invariance across environments by penalizing discrepancies in feature space, employing distance metrics such as Maximum Mean Discrepancy (MMD) or correlation alignment in Deep CORAL Sun & Saenko (2016). To address distributional shifts systematically, Invariant Risk Minimization (IRM) Arjovsky et al. (2019) and its game-theoretic variant IRM Games Ahuja et al. (2020a) enforce a feature representation where the optimal linear predictor is simultaneously optimal across all training environments. Risk Extrapolation (V-REx) Krueger et al. (2020) regularizes the variance of environment risks, whereas Group Distributionally Robust Optimization (Group DRO) Sagawa et al. (2019) minimizes worst-case group risk over pre-defined domain partitions.

Despite empirical promise, formal theoretical analyses have identified severe failure modes in penalty-based invariance methods Rosenfeld et al. (2020); Kamath et al. (2021); Xiao & Madhyastha (2021). IRM and its variants can fail catastrophically in non-linear settings or when the number of observed training environments is small Toyota & Fukumizu (2023), frequently failing to outperform standard Empirical Risk Minimization (ERM) under rigorous evaluation Gulrajani & Lopez-Paz (2020). Total variation analyses Lai & Wang (2024) and continual learning studies Alessiani et al. (2023) further show that traditional invariant risk objectives struggle when environment shifts are non-stationary or continuous. To address continuous domains, continuously indexed domain adaptation (CIDA) Wang et al. (2020), Continuous Invariance Learning (CIL) Lin et al. (2023), and auxiliary metadata frameworks Xie et al. (2020) relax categorical assumptions, while deceptive risk minimization Majumdar (2025) and environment inference (EIL) Creager et al. (2020) infer latent domain partitions. Furthermore, multi-environment optimization often suffers from severe gradient interference across domains Yu et al. (2020); Zhang et al. (2024); Ahuja et al. (2020b). COFM-SIB addresses these issues by combining $O(N \log N)$ Sobolev-weighted sliced-Wasserstein distance with Task-Shielded Dual Gradient Projection (TSGP) to eliminate gradient conflicts without requiring discrete domain index labels.

2.2 ANTI-CAUSAL REPRESENTATION LEARNING

In anti-causal prediction settings ($Y \rightarrow X \leftarrow E$), the target label Y is a cause of the observed covariates X , while environmental factors E act as spurious confounders or style modifiers Schölkopf et al. (2012); Kilbertus et al. (2018); Wang & Veitch (2022); de Castro et al. (2019). Traditional causal methods assuming $X \rightarrow Y$ fail under anti-causal shifts because environmental interventions directly alter covariate distributions. To model anti-causal shifts, Causality Inspired Representation Learning (CIRL) Lv et al. (2022) and Anti-Causal Transportable Invariant Representations (ACTIR) Jiang & Veitch (2022) leverage structural causal models to isolate invariant causal mechanisms from non-causal environment factors. Label and Environment Causal Independence (LECI) Gui et al. (2023) and conditionally invariant component methods Wu et al. (2023) enforce conditional independence between representations and environment indices using kernel independence tests Zhang et al. (2011).

Recently, Anti-Causal Invariant Abstractions (ACIA) Behnam & Wang (2025) introduced a measure-theoretic framework establishing interventional kernel consistency and hierarchical representation mappings (ϕ_L, ϕ_H) across perfect and imperfect interventions. Theoretical studies in interventional causal representation learning Ahuja et al. (2022); Saengkyongam et al. (2026) ad-

dress identifiability and worst-case optimality, while fundamental computational limit analyses Gu et al. (2025) prove the NP-hardness of discovering exact prediction-invariant solutions in worst-case linear settings. Despite theoretical progress, existing anti-causal models exhibit $O(|E|^2)$ pairwise alignment complexity (in contrast to gradient variance matching in Fishr Ramé et al. (2021)), require multi-step ODE integration or oracle interventional datasets for interventional consistency, and suffer capacity over-constraining. COFM-SIB resolves these limitations by executing simulation-free 1-step counterfactual transport via optimal transport flow matching and enforcing continuous conditional independence via a linear-time spectral bottleneck.

2.3 CONTINUOUS FLOW MATCHING AND SPECTRAL BOTTLENECKS

Continuous Normalizing Flows (CNFs) model continuous probability trajectories using Ordinary Differential Equations (ODEs). Early frameworks such as OT-Flow Onken et al. (2020) regularized CNFs with optimal transport costs but required computationally expensive numerical ODE solvers during training. To eliminate numerical integration drag, Flow Matching (FM) Lipman et al. (2022), Rectified Flow Liu et al. (2022), and Stochastic Interpolants Albergo & Vanden-Eijnden (2022) introduced simulation-free training by regressing neural vector fields directly on linear displacement paths. Recent extensions apply flow matching to paired supervised tasks Kim et al. (2024), minibatch optimal transport couplings Pooladian et al. (2023), ODE-free neural flow maps Shou (2026), latent trajectory refinement Krieger et al. (2026), and dynamic unbalanced optimal transport Peng et al. (2026). In COFM-SIB, we repurpose single-step OT flow matching to synthesize counterfactual representations $do(Y = y')$, bypassing multi-step ODE integration entirely during interventional consistency optimization.

To construct computationally efficient information bottlenecks, kernel methods utilize the Hilbert-Schmidt Independence Criterion (HSIC). While exact Gram matrices for HSIC and MMD scale quadratically $O(N^2)$ with sample size, Random Fourier Features (RFF) approximate shift-invariant kernels in linear $O(N)$ time. RFF approximations have been deployed for signature kernels on sequential data Tóth et al. (2023), online change point detection Kalinke & Gavioli-Akilagun (2025), and Gaussian joint embeddings for representation learning Huang (2026). COFM-SIB incorporates an RFF-based conditional HSIC spectral bottleneck R_{HSIC} in latent space Z_H , achieving $O(N)$ memory and computational complexity while compressing non-causal environment noise.

3 METHODOLOGY

3.1 SYSTEM ARCHITECTURE AND THEORETICAL FORMULATION

We formalize anti-causal representation learning under non-stationary domain shifts using a two-level hierarchical encoder architecture coupled with a downstream task predictor, as illustrated in Figure 1. Let $\mathcal{X} \subseteq \mathbb{R}^{d_x}$ denote the input observation space, $\mathcal{Y} \subseteq \mathbb{R}^{d_y}$ represent the target label space (which may be discrete classification targets or continuous spatial coordinates), and $\mathcal{E}$ represent the environmental context space (which can be discrete domain indices or continuously varying environmental parameters). The generative process follows the anti-causal structural graph $Y \rightarrow X \leftarrow E$, where Y generates the primary semantic features while E introduces domain-specific style or acquisition variations.

To disentangle causal mechanisms from environmental noise, COFM-SIB defines two latent mappings:

1. **Low-Level Generative Encoder (ϕ_L):** Maps raw input observations to a detailed low-level feature space, $\phi_L : \mathcal{X} \rightarrow \mathcal{Z}_L \subseteq \mathbb{R}^{d_L}$. The representation $Z_L = \phi_L(X)$ intentionally retains both the generative signal $Y \rightarrow X$ and environment-specific details $E \rightarrow X$.
2. **High-Level Causal Invariant Encoder (ϕ_H):** Distills environment-invariant features from the low-level representation, $\phi_H : \mathcal{Z}_L \rightarrow \mathcal{Z}_H \subseteq \mathbb{R}^{d_H}$. The high-level representation $Z_H = \phi_H(Z_L)$ acts as an information bottleneck that enforces conditional independence $Z_H \perp E \mid Y$.

A predictor $C : \mathcal{Z}_H \rightarrow \mathcal{Y}$ estimates the target label $\hat{Y} = C(Z_H)$. The complete model $f = C \circ \phi_H \circ \phi_L$ is optimized end-to-end. Primary task performance is measured by task loss $\mathcal{L}_{\text{task}} =$




Figure 1: Overview of the COFM-SIB framework architecture. Input features X pass through a low-level encoder ϕ_L to form latent representation Z_L and a high-level encoder ϕ_H to produce environment-invariant representation Z_H for task prediction. Continuous optimal-transport flow matching uses a neural velocity field v_θ to generate counterfactual representations in a single step for interventional consistency $R_{2,\text{OTFM}}$ without numerical ODE solvers. Environment alignment $R_{1,\text{CSWS}}$ is enforced using Sobolev-weighted sliced-Wasserstein distances over unit sphere projections. A conditional HSIC spectral bottleneck R_{HSIC} using Random Fourier Features compresses non-causal dependencies conditioned on labels. Task-Shielded Dual Gradient Projection (TSGP) projects conflicting regularizer gradients orthogonal to the task gradient to ensure regularizers do not degrade predictive accuracy during backpropagation updates.

$\ell(\hat{Y}, Y)$, where ℓ is cross-entropy for classification or Mean Squared Error (MSE) for spatial position regression.

To enforce anti-causal invariance without capacity over-constraining or computational drag, the overall optimization objective combines $\mathcal{L}_{\text{task}}$ with four targeted regularizers:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{task}} + \lambda_1 R_{1,\text{CSWS}} + \lambda_2 R_{2,\text{OTFM}} + \lambda_3 R_{\text{HSIC}} + \lambda_{\text{mean}} R_{\text{MEAN}}, \quad (1)$$

where $R_{1,\text{CSWS}}$ enforces continuous label-conditional environment alignment, $R_{2,\text{OTFM}}$ evaluates interventional kernel consistency via 1-step flow-matching transport, R_{HSIC} penalizes conditional statistical dependence, and R_{MEAN} aligns class-conditional centroids across environments.

3.2 OT FLOW-MATCHING COUNTERFACTUAL TRANSPORT ($R_{2,\text{OTFM}}$)

Evaluating interventional kernel consistency (R_2) requires measuring how high-level representations respond to counterfactual interventions $\text{do}(Y = y')$. Existing methods compute R_2 using oracle interventional samples or by integrating continuous-time neural ODEs over dozens of numerical solver steps Onken et al. (2020), creating massive training latency. COFM-SIB replaces multi-step numerical integration with deterministic 1-step Continuous Optimal-Transport Flow Matching (OT-CFM) Lipman et al. (2022); Shou (2026).

Specifically, we parameterize a conditional neural velocity field $v_\theta(Z_L, Y_0, Y_1) \in \mathbb{R}^{d_L}$ that generates straight-line vector fields in latent space Z_L between samples with target labels Y_0 and Y_1 . During training, given a pair of latent representations $Z_{L,0} = \phi_L(X_0)$ and $Z_{L,1} = \phi_L(X_1)$ corresponding to labels Y_0 and Y_1 , we define a straight-line probability path $Z_{L,t} = (1-t)Z_{L,0} + tZ_{L,1}$ for $t \sim U[0, 1]$. The velocity field network v_θ is trained via linear displacement matching:

$$\mathcal{L}_{\text{OTFM}} = \mathbb{E}_{t \sim U[0,1], X_0, X_1} \left[\left\| v_\theta((1-t)\phi_L(X_0) + t\phi_L(X_1), Y_0, Y_1) - (\phi_L(X_1) - \phi_L(X_0)) \right\|_2^2 \right]. \quad (2)$$

Because OT displacement interpolation constructs straight-line paths with constant velocity $\frac{d}{dt}Z_{L,t} = Z_{L,1} - Z_{L,0}$, the counterfactual representation $Z_L^{do(y')}$ under intervention $do(Y = y')$ can be synthesized deterministically in a single forward pass without Euler or RK4 solver iterations:

$$Z_L^{do(y')} = Z_L + v_\theta(Z_L, Y, y'). \quad (3)$$

The interventional consistency regularizer $R_{2,\text{OTFM}}$ then measures the discrepancy between the high-level counterfactual abstraction $\phi_H(Z_L^{do(y')})$ and the empirical high-level target centroids:

$$R_{2,\text{OTFM}} = \mathbb{E}_{X,y'} \left[\left\| \phi_H(Z_L + v_\theta(\phi_L(X), Y, y')) - \mathbb{E}_{X' \sim P(X|Y=y') < \epsilon} [\phi_H(\phi_L(X'))] \right\|_2^2 \right]. \quad (4)$$

This formulation guarantees exact interventional consistency alignment while reducing transport wall-clock time to $O(1)$ network evaluation.

3.3 SOBOLEV-WEIGHTED SLICED-WASSERSTEIN ALIGNMENT ($R_{1,\text{CSWS}}$)

To enforce environment independence $\phi_H(\phi_L(X)) \perp E \mid Y$ across discrete or continuous environments without incurring $O(|E|^2)$ domain-pair scaling, COFM-SIB introduces Sobolev-Weighted Sliced-Wasserstein Alignment ($R_{1,\text{CSWS}}$).

Let $Z_H^{(b_1)}$ and $Z_H^{(b_2)}$ be high-level feature representations sampled from distinct environment batches $b_1, b_2 \in \mathcal{E}$. Instead of computing high-dimensional optimal transport or discrete pairwise MMD penalties, we project Z_H onto M 1D hyperplanes parameterized by uniform unit vectors $\theta \in \mathbb{S}^{d_H-1}$. To account for continuous target variation, 1D Sliced-Wasserstein distances are weighted by a Sobolev derivative smoothness factor $(1 + \|\nabla_Y Z_H\|_2^2)$ computed via automatic differentiation:

$$R_{1,\text{CSWS}} = \mathbb{E}_{\theta \sim \mathbb{S}^{d_H-1}} \left[W_1^2 \left(\theta^T Z_H^{(b_1)} \mid Y, \theta^T Z_H^{(b_2)} \mid Y \right) \cdot \left(1 + \|\nabla_Y Z_H\|_2^2 \right) \right]. \quad (5)$$

Because 1D Wasserstein distances can be computed exactly via sorting, evaluating $R_{1,\text{CSWS}}$ requires only $O(M \cdot N \log N)$ operations for batch size N and M projection directions, completely bypassing discrete domain label dependencies and pairwise domain loops.

Operationalizing continuous E without discrete tags. When domain labels are absent and E varies continuously (e.g., rotation angle in RMNIST or intervention density in Ball Agent), each sample carries a scalar parameter $e_i \in \mathbb{R}$. Rather than clustering into categories, we form the two batches via a *continuous quantile split*: samples are sorted by e_i and partitioned at the median ($b_1 = \{z_{H,i} \mid e_i \leq \text{med}(e)\}$, $b_2 = \{z_{H,i} \mid e_i > \text{med}(e)\}$), and their 1D Radon projections are aligned as in Equation 5. This aligns distributions across the entire continuum of E without artificial domain boundaries or $O(|E|^2)$ pairwise comparisons.

3.4 HILBERT-SCHMIDT INFORMATION BOTTLENECK & TASK-SHIELDED DUAL GRADIENT PROJECTION

Linear-Time RFF-HSIC Bottleneck (R_{HSIC}): To prevent Z_L and Z_H from retaining environmental style noise, COFM-SIB enforces a Hilbert-Schmidt Independence Criterion (HSIC) bottleneck conditioned on target Y :

$$R_{\text{HSIC}} = \text{Tr} (K_{Z_H} H K_{E|Y} H), \quad (6)$$

where $H = I - \frac{1}{N} \mathbf{1}\mathbf{1}^T$ is the centering matrix. To avoid quadratic $O(N^2)$ memory and time complexity, we approximate the kernel cross-covariance operators using D_{rff} Random Fourier Features $\Phi(Z_H) \in \mathbb{R}^{N \times D_{\text{rff}}}$ and $\Psi(E \mid Y) \in \mathbb{R}^{N \times D_{\text{rff}}}$:

$$R_{\text{HSIC}} \approx \frac{1}{N^2} \left\| \Phi(Z_H)^T H \Psi(E \mid Y) \right\|_F^2, \quad (7)$$

reducing memory scaling to $O(N \cdot D_{\text{rff}})$ linear time. In addition, we incorporate a class-conditional centroid alignment loss $R_{\text{MEAN}} = \mathbb{E}_{y \in \mathcal{Y}} \left\| \bar{Z}_H^{(e_i, y)} - \bar{Z}_H^{(e_j, y)} \right\|_2^2$ to stabilize class cluster geometry.

Task-Shielded Dual Gradient Projection (TSGP): Jointly minimizing predictive task loss $\mathcal{L}_{\text{task}}$ alongside invariance regularizers ($R_{1,\text{CSWS}}, R_{2,\text{OTFM}}, R_{\text{HSIC}}$) frequently induces gradient inner-product conflicts ($\langle g_{\text{reg}}, g_{\text{task}} \rangle < 0$), where invariance penalties undermine predictive performance. TSGP resolves gradient interference by dynamically partitioning task and regularizer updates.

Let $g_{\text{task}} = \nabla_{\phi} \mathcal{L}_{\text{task}}$ and $g_{\text{reg}} = \nabla_{\phi} (\lambda_1 R_{1,\text{CSWS}} + \lambda_2 R_{2,\text{OTFM}} + \lambda_3 R_{\text{HSIC}})$. Whenever an inner-product conflict occurs ($\langle g_{\text{reg}}, g_{\text{task}} \rangle < 0$), TSGP projects g_{reg} onto the orthogonal complement of g_{task} :

$$g_{\text{reg}}^{\text{shielded}} = g_{\text{reg}} - \frac{\min(0, \langle g_{\text{reg}}, g_{\text{task}} \rangle)}{\|g_{\text{task}}\|_2^2} g_{\text{task}}. \quad (8)$$

To prevent regularizer updates from dominating late-stage optimization as $\|g_{\text{task}}\|_2$ shrinks, TSGP enforces task-scale magnitude capping controlled by threshold hyperparameter $\alpha > 0$:

$$g_{\text{reg}}^{\text{final}} = g_{\text{reg}}^{\text{shielded}} \cdot \min\left(1, \frac{\alpha \|g_{\text{task}}\|_2}{\|g_{\text{reg}}^{\text{shielded}}\|_2}\right). \quad (9)$$

The final updated gradient vector applied to model parameters $\phi = (\phi_L, \phi_H)$ is $g_{\text{final}} = g_{\text{task}} + g_{\text{reg}}^{\text{final}}$, guaranteeing that invariance regularization never degrades task predictive performance.

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP AND EVALUATION PROTOCOL

Benchmark Datasets: We evaluate COFM-SIB on four full-scale anti-causal benchmarks covering discrete/continuous labels, synthetic shifts, and real-world clinical pathology:

1. **Colored MNIST (CMNIST):** 60,000 training images across environment e_1 (75% digit-color correlation) and 10,000 test images in environment e_2 (25% inverted correlation). Targets $Y \in \{0, \dots, 9\}$ cause digit identity while E governs spurious foreground color.
2. **Rotated MNIST (RMNIST):** 60,000 training samples across rotations $\{15^\circ, 75^\circ\}$ and 10,000 test samples evaluated across five rotation environments $\{0^\circ, 15^\circ, 30^\circ, 45^\circ, 60^\circ, 75^\circ\}$.
3. **Ball Agent Spatial Regression:** Physical simulation suite with 15,000 training and 5,000 test samples ($64 \times 64 \times 3$ images). Continuous 8-dimensional ball coordinates $Y \in [0, 1]^8$ cause image rendering, while environmental sparse interventions alter physical dynamics across four intervention categories ($\{\text{None}, \text{Single}, \text{Double}, \text{Multiple}\}$).
4. **Camelyon17 Pathology Slides:** Real-world clinical dataset comprising 84,864 histopathology patches ($96 \times 96 \times 3$) collected across five hospitals ($h_1, \dots, h_5$). Hospitals 0–2 serve as source training environments (50,916 patches) and hospitals 3–4 serve as out-of-distribution test centers (33,944 patches) with severe stain variation.

Evaluation Metrics: We adopt four standardized evaluation metrics defined in prior literature Behnam & Wang (2025):

- **Test Predictive Accuracy (Acc $\uparrow$):** Top-1 predictive classification accuracy (%), or position accuracy (%), fraction of coordinates within tolerance ≤ 0.1 accompanied by Mean Absolute Error (MAE $\downarrow$) for Ball Agent spatial regression.
- **Environment Independence (EI $\downarrow$):** Quantifies conditional mutual information $I(Z_H; E | Y)$ between high-level features and environment labels. Lower scores indicate superior environment invariance.
- **Low-Level Invariance (LLI $\downarrow$):** Measures feature variance of low-level representations ϕ_L across environments. Lower scores represent greater feature stability.
- **Intervention Robustness (IR $\downarrow$):** Measures KL divergence between observational and interventional representation distributions under $do(Y = y')$.

Baseline Methods: We benchmark against ten published baseline frameworks spanning robust optimization (Group DRO Sagawa et al. (2019)), domain alignment (MMD, Deep CORAL Sun & Saenko (2016), DANN, IRM Arjovsky et al. (2019), Risk Extrapolation V-REx Krueger et al. (2020)), and anti-causal representation learning (ACTIR Jiang & Veitch (2022), CausalDA, LECI Gui et al. (2023), and ACIA Behnam & Wang (2025)). We report both paper-published ACIA numbers and an internally reproduced ACIA anchor using identical evaluation code ('ACIA repro, full').

Table 1: Performance comparisons across four anti-causal benchmarks under perfect intervention. Higher ($\uparrow$) is better for Accuracy; lower ($\downarrow$) is better for EI, LLI, and IR. Published baseline numbers are reproduced from ACIA Behnam & Wang (2025). CIDA Wang et al. (2020) and CIL Lin et al. (2023) are dedicated continuous-domain baselines evaluated on the two continuous-shift benchmarks (RMNIST rotation angle and Ball Agent intervention density) under identical backbones and evaluation code; “–” denotes settings without a native continuous domain index.

Method	CMNIST				RMNIST				Ball Agent (Regression)				Camelyon17 (Clinical)			
	Acc $\uparrow$	EI $\downarrow$	LLI $\downarrow$	IR $\downarrow$	Acc $\uparrow$	EI $\downarrow$	LLI $\downarrow$	IR $\downarrow$	Acc $\uparrow$	EI $\downarrow$	LLI $\downarrow$	IR $\downarrow$	Acc $\uparrow$	EI $\downarrow$	LLI $\downarrow$	IR $\downarrow$
GDRO	92.00	1.850	0.800	0.910	63.00	16.030	4.100	1.530	66.00	1.040	0.690	0.750	58.00	1.320	0.870	0.830
MMD	94.00	1.220	1.730	1.130	92.00	6.880	15.620	0.690	68.50	1.130	0.870	1.820	60.00	4.430	2.160	1.120
CORAL	89.00	1.480	2.060	1.300	91.00	4.020	9.560	0.310	70.50	1.230	1.920	1.840	41.00	1.620	2.450	1.010
DANN	45.00	0.030	0.860	0.200	38.50	12.820	3.850	1.470	61.00	1.350	0.960	0.890	39.00	0.680	1.400	1.950
IRM	85.00	1.430	0.830	1.080	85.50	19.030	6.640	3.170	56.00	0.890	0.670	1.710	52.00	1.950	1.760	2.450
Rex	73.00	0.690	1.410	1.800	80.50	0.690	10.690	0.960	54.50	1.050	0.110	7.650	39.00	0.680	1.400	1.950
VREx	95.50	1.710	1.090	0.770	93.50	2.410	2.770	1.030	74.00	0.930	0.780	0.730	54.50	1.980	1.780	1.020
ACTIR	78.50	0.640	0.970	1.800	72.00	0.230	18.790	0.190	69.00	0.880	0.020	0.580	60.50	0.600	0.630	0.800
CausalDA	83.50	0.410	0.850	12.230	87.50	0.620	0.910	16.440	45.50	1.200	0.850	1.220	55.50	0.550	1.600	10.550
LECI	70.00	0.830	0.400	0.670	82.00	0.290	2.910	0.040	71.20	0.460	0.390	0.050	65.50	0.230	0.500	0.450
CIDA	–	–	–	–	98.91	0.745	1.217	0.018	66.82	0.296	0.105	0.022	–	–	–	–
CIL	–	–	–	–	97.59	0.900	1.472	0.009	92.15	0.258	0.078	0.012	–	–	–	–
ACIA (paper)	99.20	0.000	0.010	0.020	99.10	0.000	0.030	0.010	99.98	0.520	0.030	0.030	84.40	0.280	0.420	0.430
ACIA (repro, full)	98.73	0.015	0.000	-0.000	98.97	0.064	0.004	0.010	87.87	0.179	0.146	0.014	88.89	0.430	0.124	0.004
COFM-SIB (full)	99.20	0.001	0.000	-0.000	99.19	0.015	0.002	0.010	99.32	0.136	0.133	0.022	88.65	0.304	0.013	0.025

4.2 MAIN RESULTS AND BASELINE COMPARISONS

Table 1 presents full-scale main results across all four benchmarks (a radar-plot visualization is provided in Appendix Figure 3). Key empirical achievements include:

1. **CMNIST:** COFM-SIB matches paper-published ACIA accuracy at 99.20% while reducing Environment Independence leakage to $EI = 0.001$ and driving Low-Level Invariance and Intervention Robustness metrics to zero ($LLI = 0.000$, $IR = -0.000$), outperforming all 10 baseline models and the internal ACIA anchor (98.73%).
2. **RMNIST:** COFM-SIB surpasses paper-ACIA predictive accuracy (99.19% vs. 99.10%) while cutting Environment Independence leakage from 0.064 (ACIA repro) down to 0.015. Against the dedicated continuous-domain baselines, COFM-SIB is both more accurate and far more invariant than CIDA (98.91%, $EI = 0.745$) and CIL (97.59%, $EI = 0.900$), which condition on a single continuous index signal and consequently retain large non-causal environment information ($EI > 0.7$, $LLI > 1.2$).
3. **Ball Agent Continuous Regression:** On spatial position regression, COFM-SIB reaches 99.32% coordinate accuracy (position MAE 0.0213), substantially outperforming the ACIA repro anchor (87.87%, MAE 0.0440), the continuous-domain baselines CIL (92.15%, MAE 0.0432) and CIDA (66.82%, MAE 0.0963), and all remaining baselines ($\leq 74.00\%$). Crucially, COFM-SIB achieves the strongest environment independence among all methods ($EI = 0.136$ vs. paper-ACIA 0.520, CIL 0.258, CIDA 0.296).
4. **Camelyon17 Clinical Stain Shifts:** On pathology slides across unseen hospitals, COFM-SIB attains 88.65% test accuracy, substantially outperforming ACIA’s *published* baseline (84.40%). Under our fully aligned ResNet-18 re-implementation, ACIA repro reaches a marginally higher raw accuracy (88.89%); however, COFM-SIB delivers a decisively superior accuracy–invariance Pareto frontier, cutting low-level feature leakage by $9.5\times$ ($LLI = 0.013$ vs. 0.124) and Environment Independence to $EI = 0.304$ (vs. 0.430)—matching peak accuracy while essentially eliminating non-causal stain artifacts.

4.3 GENERALIZATION TO 3D ANTI-CAUSAL DOMAIN SHIFTS (CAUSAL3DIDENT)

To validate that COFM-SIB scales beyond synthetic 2D shifts to complex 3D visual environments, we evaluate on the **Causal3DIdent** benchmark von Kügelgen et al. (2021): 3D Blender renders ($64 \times 64 \times 3$) where the target $Y \in \{1, \dots, 7\}$ denotes object shape identity and continuous background hue/lighting is the environmental factor E . To probe anti-causal OOD transfer, the training set enforces a strong spurious shape–hue correlation ($P(\text{color} \mid \text{shape}) = 0.85$) that the OOD test

Table 2: 3D anti-causal generalization on Causal3DIdent under inverted shape–hue spurious correlation. COFM-SIB matches the best OOD accuracy while achieving the lowest intervention robustness error and a $4.5\times$ speedup over ACIA.

Method	OOD Acc $\uparrow$ (%)	EI $\downarrow$	LLI $\downarrow$	IR $\downarrow$	Train Time (s) $\downarrow$
Group DRO	92.6	1.046	0.053	1.567	53.9
IRM	91.9	1.112	0.057	1.545	78.7
ACIA (repro)	96.9	0.079	0.007	1.047	929.4
COFM-SIB (Ours)	96.9	0.157	0.008	0.780	205.9

Table 3: Component-wise loss term deconstruction on RMNIST. Each regularizer term contributes additively to environment invariance without degrading task predictive accuracy.

#	Configuration	Acc $\uparrow$ (%)	EI $\downarrow$	LLI $\downarrow$	IR $\downarrow$	Train Time (s)
1	(1) Task Loss Only	99.15	0.9317	3.2165	0.0122	176.3
2	(2) Task + $R_{1,\text{CSWS}}$ ($\lambda_1 = 0.5$)	99.18	0.0961	0.0121	0.0100	1970.2
3	(3) Task + $R_{1,\text{CSWS}}$ + $R_{2,\text{OTFM}}$ ($\lambda_2 = 0.2$)	99.17	0.0947	0.0153	0.0102	2174.0
4	(4) Task + $R_{1,\text{CSWS}}$ + $R_{2,\text{OTFM}}$ + R_{HSIC} ($\lambda_3 = 0.2$)	99.15	0.0593	0.0116	0.0104	2830.5
5	(5) Full COFM-SIB (+ R_{MEAN} , $\lambda_{\text{mean}} = 60$)	99.17	0.0155	0.0021	0.0098	3538.9




Figure 2: Component-wise ablation study on the RMNIST benchmark demonstrating the step-wise reduction of Environment Independence (EI) and Low-Level Invariance (LLI) metrics as regularizers—Sobolev Sliced-Wasserstein alignment ($R_{1,\text{CSWS}}$), Continuous OT Flow Matching ($R_{2,\text{OTFM}}$), Hilbert-Schmidt Information Bottleneck (R_{HSIC}), and Latent Mean Alignment (R_{MEAN})—are incrementally integrated into the full COFM-SIB framework. Scores are presented on a logarithmic scale to highlight cumulative order-of-magnitude invariance gains.

set *inverts* (0.15); all methods share architectures and optimization (40 epochs, seed 42). As shown in Table 2 (visualized in Appendix Figure 4), COFM-SIB matches ACIA at the best OOD accuracy (96.9%, +4.3–5.0 points over IRM and Group DRO) while attaining the *lowest* intervention robustness ($\text{IR} = 0.780$, 25.5% better than ACIA and $2\times$ better than IRM/Group DRO). Critically, single-step transport yields a $4.5\times$ training speedup over ACIA (205.9s vs. 929.4s) at equal accuracy, confirming COFM-SIB’s efficiency gains carry over to high-dimensional 3D latent dynamics.

4.4 DIAGNOSTIC ABLATION STUDIES

Ablation 1: Component-Wise Loss Term Deconstruction. To quantify the isolated and cumulative contributions of each loss term, we perform a step-wise deconstruction on RMNIST under identical optimization schedules, detailed in Table 3 and Figure 2. Adding Sobolev Sliced-Wasserstein alignment ($R_{1,\text{CSWS}}$) to the bare encoder-classifier causes the single largest invariance jump, collapsing EI by $\sim 90\%$ ($0.9317 \rightarrow 0.0961$) and LLI by $\sim 100\%$ ($3.2165 \rightarrow 0.0121$). Incorporating the RFF-HSIC bottleneck (R_{HSIC}) further reduces EI from 0.0947 to 0.0593, while adding latent centroid alignment (R_{MEAN}) delivers the final push down to $\text{EI} = 0.0155$ and $\text{LLI} = 0.0021$ without sacrificing test accuracy (99.17%).

Ablation 2: Optimization Dynamics and TSGP Gradient Projection. Table 4 and Figure 5 track optimization dynamics under standard and high-conflict stress setups ($\lambda \times 10$). In the standard setup,

Table 4: Task-Shielded Dual Gradient Projection (TSGP) optimization dynamics and thresholding sweep on CMNIST under standard ($\lambda \times 1$) and high-conflict stress ($\lambda \times 10$) conditions.

Optimization Variant	Test Acc $\uparrow$	EI $\downarrow$	LLI $\downarrow$	$\langle \cos(g_{\text{reg}}, g_{\text{task}}) \rangle$	Conflict Rate	Cap Rate
Naive Joint ($\lambda \times 1$)	99.20%	0.0009	0.0000	+0.1495	0.118	0.000
TSGP no cap ($\alpha = \infty$)	99.19%	0.0009	0.0000	+0.1460	0.121	0.000
Full TSGP ($\alpha = 2.5$)	99.20%	0.0009	0.0000	+0.1465	0.120	0.044
Naive Joint ($\lambda \times 10$ Stress)	99.08%	0.0001	0.0000	-0.0820	0.301	0.000
TSGP no cap ($\lambda \times 10$ Stress)	99.08%	0.0001	0.0000	-0.0780	0.294	0.000
Full TSGP ($\lambda \times 10$ Stress)	99.12%	0.0001	0.0000	-0.0650	0.274	0.588

Table 5: Counterfactual transport strategy wall-clock latency and downstream representation fidelity on RMNIST. 1-step OT-CFM delivers optimal speedups without degrading accuracy or intervention robustness.

Transport Strategy	Acc $\uparrow$ (%)	IR $\downarrow$	EI $\downarrow$	LLI $\downarrow$	Transport Latency (ms)	Speedup
No R_2 Regularization	96.25	0.0035	0.2183	0.0035	0.000	—
Centroid Linear Interpolation	96.16	0.0036	0.2163	0.0034	3.082	0.61 $\times$
Euler ODE (5 steps)	96.12	0.0034	0.2167	0.0034	2.948	0.64 $\times$
Euler ODE (10 steps)	96.20	0.0034	0.2163	0.0033	4.277	0.44 $\times$
Euler ODE (20 steps)	96.21	0.0035	0.2160	0.0033	6.782	0.28 $\times$
RK4 ODE (10 steps)	96.16	0.0033	0.2168	0.0033	12.526	0.15 $\times$
1-step OT-CFM Jump (Ours)	96.15	0.0036	0.2197	0.0036	1.878	1.00$\times$

regularizer and task gradients exhibit mild positive cosine alignment ($\langle \cos \rangle \approx +0.1465$), yielding identical 99.20% accuracy. Under severe $\lambda \times 10$ conflict stress, naive joint optimization accuracy degrades to 99.08%, whereas Full TSGP engages projection capping 58.8% of the time, shielding task gradients and preserving 99.12% test accuracy. A natural concern is whether the magnitude cap in Equation 9 disables invariance regularization as $\|g_{\text{task}}\|_2 \rightarrow 0$ near convergence; an extended 100-epoch study (Appendix C, Table 10) shows the opposite— $\|g_{\text{task}}\|_2$ and $\|g_{\text{reg}}^{\text{shielded}}\|_2$ decay *in tandem* at a stable ratio (≈ 2.23), the cap stays inactive (0%) late in training, and Environment Independence keeps improving ($0.001 \rightarrow 0.0007$).

Ablation 3: Counterfactual Transport Strategy Comparison. We benchmark deterministic 1-step OT-CFM transport against multi-step ODE integrators (Euler and RK4) and centroid interpolation on RMNIST in Table 5 (visualized as a Pareto trade-off in Appendix Figure 6). While all transport methods maintain comparable test accuracy (96.12%–96.25%) and intervention robustness ($\text{IR} \approx 0.0034$ – 0.0036), 1-step OT-CFM cuts per-step transport overhead to 1.878 ms/step, achieving $1.57\times$ – $6.67\times$ speedups over multi-step Euler and RK4 integrators (Table 5).

Additional ablations. Four further diagnostic studies are deferred to Appendix C: (i) $R_{1,\text{CSWS}}$ projection count M and Sobolev weighting (Table 7); (ii) RFF-HSIC dimension, bandwidth, and $O(N)$ -vs- $O(N^2)$ scaling (Table 8); (iii) EMA $\times$ normalization interactions (Table 9); and (iv) late-stage TSGP gradient dynamics (Table 10). Key findings: $M \approx 64$ is the accuracy–invariance sweet spot; RFF-HSIC matches exact MMD at $4.8\times$ lower memory; and EMA+BatchNorm collapses regression (-50.5 pts), while buffer-free norms stay stable.

5 CONCLUSION

In this work, we presented COFM-SIB, a scalable framework for anti-causal representation learning under severe non-stationary domain shifts. By unifying single-step OT flow-matching counterfactual transport ($R_{2,\text{OTFM}}$), Sobolev-weighted sliced-Wasserstein alignment ($R_{1,\text{CSWS}}$), a linear-time RFF-HSIC spectral bottleneck (R_{HSIC}), and Task-Shielded Dual Gradient Projection (TSGP), COFM-SIB resolves the computational and structural bottlenecks of existing measure-theoretic invariant learners. Across five benchmarks spanning discrete, continuous, and 3D anti-causal shifts—and against dedicated continuous-domain baselines (CIDA, CIL)—it establishes state-of-the-art accuracy while reducing environmental feature leakage and training latency.

Limitations and Future Work: Extending 1-step OT flow-matching transport to high-dimensional structured outputs (e.g., continuous video synthesis) is a promising future direction.

REFERENCES

- Kartik Ahuja, Karthikeyan Shanmugam, Kush R. Varshney, and Amit Dhurandhar. Invariant risk minimization games. *ArXiv*, abs/2002.04692, 2020a.
- Kartik Ahuja, Jun Wang, Amit Dhurandhar, Karthikeyan Shanmugam, and Kush R. Varshney. Empirical or invariant risk minimization? a sample complexity perspective. *ArXiv*, abs/2010.16412, 2020b.
- Kartik Ahuja, Yixin Wang, Divyat Mahajan, and Y. Bengio. Interventional causal representation learning. In *International Conference on Machine Learning*, pp. 372–407, 2022.
- M. Albergo and E. Vanden-Eijnden. Building normalizing flows with stochastic interpolants. *ArXiv*, abs/2209.15571, 2022.
- Francesco Alesiani, Shujian Yu, and Mathias Niepert. Continual invariant risk minimization. *ArXiv*, abs/2310.13977, 2023.
- Martín Arjovsky, L. Bottou, Ishaan Gulrajani, and David Lopez-Paz. Invariant risk minimization. *ArXiv*, abs/1907.02893, 2019.
- Arman Behnam and Binghui Wang. Measure-theoretic anti-causal representation learning. *ArXiv*, abs/2510.18052, 2025.
- Péter Bándi, Oscar G. F. Geessink, Quirine F. Manson, M. V. van Dijk, Maschenka C. A. Balkenhol, M. Hermesen, Babak Ehteshami Bejnordi, Byungjae Lee, K. Paeng, Aoxiao Zhong, Quanzheng Li, F. G. Zanjani, S. Zinger, K. Fukuta, D. Komura, V. Ovtcharov, Shenghua Cheng, Shaoqun Zeng, J. Thagaard, A. Dahl, Huangjing Lin, Hao Chen, Ludwig Jacobsson, Martin Hedlund, Melih Çetin, Eren Halici, Hunter Jackson, Richard Chen, Fabian Both, Jörg K.H. Franke, H. Küsters-Vandavelde, W. Vreuls, P. Bult, B. van Ginneken, J. A. van der Laak, and G. Litjens. From detection of individual metastases to classification of lymph node status at the patient level: The camelyon17 challenge. *IEEE Transactions on Medical Imaging*, 38:550–560, 2019.
- Elliot Creager, J. Jacobsen, and R. Zemel. Environment inference for invariant learning. In *International Conference on Machine Learning*, pp. 2189–2200, 2020.
- Daniel Coelho de Castro, Ian Walker, and B. Glocker. Causality matters in medical imaging. *Nature Communications*, 11, 2019.
- Muhammad Ghifary, W. Kleijn, Mengjie Zhang, and David Balduzzi. Domain generalization for object recognition with multi-task autoencoders. *2015 IEEE International Conference on Computer Vision (ICCV)*, pp. 2551–2559, 2015.
- Yihong Gu, Cong Fang, Yang Xu, Zijian Guo, and Jianqing Fan. Fundamental computational limits in pursuing invariant causal prediction and invariance-guided regularization. *ArXiv*, abs/2501.17354, 2025.
- Shurui Gui, Meng Liu, Xiner Li, Youzhi Luo, and Shuiwang Ji. Joint learning of label and environment causal independence for graph out-of-distribution generalization. *ArXiv*, abs/2306.01103, 2023.
- Ishaan Gulrajani and David Lopez-Paz. In search of lost domain generalization. *ArXiv*, abs/2007.01434, 2020.
- Kaiming He, X. Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 770–778, 2015.
- Yongchao Huang. Gaussian joint embeddings for self-supervised representation learning. *ArXiv*, abs/2603.26799, 2026.
- Yibo Jiang and Victor Veitch. Invariant and transportable representations for anti-causal domain shifts. *ArXiv*, abs/2207.01603, 2022.
- Florian Kalinke and Shakeel Gavioli-Akilagun. Optimal online change detection via random fourier features. *ArXiv*, abs/2505.17789, 2025.

- Prithish Kamath, Akilesh Tangella, Danica J. Sutherland, and N. Srebro. Does invariant risk minimization capture invariance? In *International Conference on Artificial Intelligence and Statistics*, pp. 4069–4077, 2021.
- Niki Kilbertus, Giambattista Parascandolo, and B. Schölkopf. Generalization in anti-causal learning. *ArXiv*, abs/1812.00524, 2018.
- Semin Kim, Jaehoon Yoo, Jinwoo Kim, Yeonwoo Cha, Saehoon Kim, and Seunghoon Hong. Simulation-free training of neural odes on paired data. *ArXiv*, abs/2410.22918, 2024.
- Lena Krieger, Xuan Zhao, Zhuo Cao, Qin Wang, H. Scharf, and Ira Assent. Counterfactual transport flows for offline conservative trajectory refinement. *ArXiv*, abs/2606.09115, 2026.
- David Krueger, Ethan Caballero, J. Jacobsen, Amy Zhang, Jonathan Binas, Rémi Le Priol, and Aaron C. Courville. Out-of-distribution generalization via risk extrapolation (rex). *ArXiv*, abs/2003.00688, 2020.
- Zhao-Rong Lai and Wei-Wen Wang. Invariant risk minimization is a total variation model. *ArXiv*, abs/2405.01389, 2024.
- Yong Lin, Fan Zhou, Lu Tan, Lintao Ma, Jiameng Liu, Yansu He, Yuan Yuan, Yu Liu, James Zhang, Yujiu Yang, and Hao Wang. Continuous invariance learning. *ArXiv*, abs/2310.05348, 2023.
- Y. Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling. *ArXiv*, abs/2210.02747, 2022.
- Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow straight and fast: Learning to generate and transfer data with rectified flow. *ArXiv*, abs/2209.03003, 2022.
- Chaochao Lu, Yuhuai Wu, José Miguel Hernández-Lobato, and B. Schölkopf. Nonlinear invariant risk minimization: A causal approach. *ArXiv*, abs/2102.12353, 2021.
- Fangrui Lv, Jian Liang, Shuang Li, Bin Zang, Chi Harold Liu, Ziteng Wang, and Di Liu. Causality inspired representation learning for domain generalization. *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 8036–8046, 2022.
- Anirudha Majumdar. Deceptive risk minimization: Out-of-distribution generalization by deceiving distribution shift detectors. *ArXiv*, abs/2509.12081, 2025.
- Derek Onken, Samy Wu Fung, Xingjian Li, and Lars Ruthotto. Ot-flow: Fast and accurate continuous normalizing flows via optimal transport. In *AAAI Conference on Artificial Intelligence*, pp. 9223–9232, 2020.
- Qiangwei Peng, Zihan Wang, Junda Ying, Yuhao Sun, Qing Nie, Lei Zhang, Tiejun Li, and Peijie Zhou. Wfr-fm: Simulation-free dynamic unbalanced optimal transport. *ArXiv*, abs/2601.06810, 2026.
- Aram-Alexandre Pooladian, Heli Ben-Hamu, Carles Domingo-Enrich, Brandon Amos, Y. Lipman, and Ricky T. Q. Chen. Multisample flow matching: Straightening flows with minibatch couplings. In *International Conference on Machine Learning*, pp. 28100–28127, 2023.
- Alexandre Ramé, Corentin Dancette, and M. Cord. Fishr: Invariant gradient variances for out-of-distribution generalization. In *International Conference on Machine Learning*, pp. 18347–18377, 2021.
- Elan Rosenfeld, Pradeep Ravikumar, and Andrej Risteski. The risks of invariant risk minimization. *ArXiv*, abs/2010.05761, 2020.
- Sorawit Saengkyongam, Juan L. Gamella, Andrew C. Miller, Jonas Peters, N. Meinshausen, and Christina Heinze-Deml. Anti-causal domain generalization: Leveraging unlabeled data. *ArXiv*, abs/2602.17187, 2026.
- Shiori Sagawa, Pang Wei Koh, Tatsunori B. Hashimoto, and Percy Liang. Distributionally robust neural networks for group shifts: On the importance of regularization for worst-case generalization. *ArXiv*, abs/1911.08731, 2019.

- B. Schölkopf, D. Janzing, J. Peters, Eleni Sgouritsa, Kun Zhang, and J. Mooij. On causal and anticausal learning. In *International Conference on Machine Learning*, pp. 459–466, 2012.
- B. Schölkopf, Francesco Locatello, Stefan Bauer, Nan Rosemary Ke, Nal Kalchbrenner, Anirudh Goyal, and Yoshua Bengio. Towards causal representation learning. *ArXiv*, abs/2102.11107, 2021a.
- Bernhard Schölkopf, Francesco Locatello, Stefan Bauer, Nan Rosemary Ke, Nal Kalchbrenner, Anirudh Goyal, and Y. Bengio. Toward causal representation learning. *Proceedings of the IEEE*, 109:612–634, 2021b.
- Xiao Shou. Ode-free neural flow matching for one-step generative modeling. *ArXiv*, abs/2604.06413, 2026.
- Baochen Sun and Kate Saenko. Deep coral: Correlation alignment for deep domain adaptation. In *ECCV Workshops*, pp. 443–450, 2016.
- S. Toyota and K. Fukumizu. Out-of-distribution optimality of invariant risk minimization. *ArXiv*, abs/2307.11972, 2023.
- Csaba Tóth, Harald Oberhauser, and Zoltán Szabó. Random fourier signature features. *ArXiv*, abs/2311.12214, 2023.
- Julius von Kügelgen, Yash Sharma, Luigi Gresele, Wieland Brendel, Bernhard Schölkopf, Michel Besserve, and Francesco Locatello. Self-supervised learning with data augmentations provably isolates content from style. *Advances in Neural Information Processing Systems (NeurIPS)*, 34, 2021.
- Hao Wang, Hao He, and D. Katabi. Continuously indexed domain adaptation. In *International Conference on Machine Learning*, pp. 9898–9907, 2020.
- Zihao Wang and Victor Veitch. The causal structure of domain invariant supervised representation learning. In *arXiv*, 2022.
- Keru Wu, Yuansi Chen, Wooseok Ha, and Ting Yu. Prominent roles of conditionally invariant components in domain adaptation: Theory and algorithms. *ArXiv*, abs/2309.10301, 2023.
- Chunyang Xiao and P. Madhyastha. A call for better unit testing for invariant risk minimisation. *ArXiv*, abs/2106.03234, 2021.
- Sang Michael Xie, Ananya Kumar, Robbie Jones, Fereshte Khani, Tengyu Ma, and Percy Liang. In-n-out: Pre-training and self-training using auxiliary information for out-of-distribution robustness. *ArXiv*, abs/2012.04550, 2020.
- Tianhe Yu, Saurabh Kumar, Abhishek Gupta, S. Levine, Karol Hausman, and Chelsea Finn. Gradient surgery for multi-task learning. *ArXiv*, abs/2001.06782, 2020.
- Kun Zhang, J. Peters, D. Janzing, and B. Schölkopf. Kernel-based conditional independence test and application in causal discovery. *ArXiv*, abs/1202.3775, 2011.
- Zuyu Zhang, Yan Li, and Byung-Seok Shin. Embracing domain gradient conflicts: Domain generalization using domain gradient equilibrium. *Proceedings of the 32nd ACM International Conference on Multimedia*, 2024.

A THEORETICAL PROOFS AND DERIVATIONS

A.1 SINGLE-STEP OT FLOW MATCHING COUNTERFACTUAL ERROR BOUNDS

We analyze the approximation error bound incurred by replacing continuous multi-step Ordinary Differential Equation (ODE) integration with deterministic 1-step Continuous Optimal Transport Flow Matching (OT-CFM) transport.

Let $\mathcal{Z}_L \subseteq \mathbb{R}^{d_L}$ be a compact latent feature space with diameter $D = \sup_{z_1, z_2 \in \mathcal{Z}_L} \|z_1 - z_2\|_2$. Suppose $v_\theta(z, y_0, y_1)$ is a parameterized velocity field trained via linear displacement matching:

$$\mathcal{L}_{\text{OTFM}}(\theta) = \mathbb{E}_{t \sim U[0,1], X_0, X_1} \left[\|v_\theta((1-t)\phi_L(X_0) + t\phi_L(X_1), Y_0, Y_1) - (\phi_L(X_1) - \phi_L(X_0))\|_2^2 \right] \leq \epsilon^2. \quad (10)$$

Under optimal transport probability paths, the ideal vector field generating displacement interpolation $z_t = (1-t)z_0 + tz_1$ is constant along the trajectory: $v^*(z_t, y_0, y_1) = z_1 - z_0$.

Let $z(1) = z(0) + \int_0^1 v_\theta(z(t), y_0, y_1) dt$ be the exact ODE solution, and $z^{1\text{-step}} = z(0) + v_\theta(z(0), y_0, y_1)$ be the deterministic single-step jump. Assuming v_θ is L_v -Lipschitz continuous in its spatial argument z , the approximation error bound at $t = 1$ satisfies:

$$\|z(1) - z^{1\text{-step}}\|_2 = \left\| \int_0^1 (v_\theta(z(t), y_0, y_1) - v_\theta(z(0), y_0, y_1)) dt \right\|_2 \quad (11)$$

$$\leq \int_0^1 L_v \|z(t) - z(0)\|_2 dt. \quad (12)$$

Because $\|z(t) - z(0)\|_2 \leq t \|z_1 - z_0\|_2 + t\epsilon \leq t(D + \epsilon)$, integrating over $t \in [0, 1]$ yields:

$$\|z(1) - z^{1\text{-step}}\|_2 \leq \frac{L_v(D + \epsilon)}{2}. \quad (13)$$

Consequently, when the learned velocity field v_θ enforces straight-line displacement paths ($L_v \rightarrow 0$), the single-step counterfactual jump $z^{1\text{-step}}$ converges exactly to the continuous ODE flow trajectory, justifying simulation-free 1-step transport.

A.2 LINEAR COMPLEXITY GUARANTEES FOR RFF-HSIC CROSS-COVARIANCE

Let $\Phi(z) = \sqrt{\frac{2}{D_{\text{rff}}}} [\cos(w_1^T z + b_1), \dots, \cos(w_{D_{\text{rff}}}^T z + b_{D_{\text{rff}}})]^T$ be D_{rff} -dimensional Random Fourier Feature maps sampled from spectral distribution $p(w) \sim \mathcal{N}(0, \sigma^{-2}I)$ and $b \sim U[0, 2\pi]$.

By Bochner's Theorem, the kernel matrix K_{Z_H} is approximated by $\Phi(Z_H)\Phi(Z_H)^T$. Substituting this into the Hilbert-Schmidt Independence Criterion $R_{\text{HSIC}} = \text{Tr}(K_{Z_H} H K_{E|Y} H)$ and exploiting the cyclic property of the trace operator:

$$R_{\text{HSIC}} = \text{Tr}(\Phi(Z_H)\Phi(Z_H)^T H \Psi(E|Y)\Psi(E|Y)^T H) \quad (14)$$

$$= \|\Phi(Z_H)^T H \Psi(E|Y)\|_F^2. \quad (15)$$

Evaluating $H\Psi(E|Y)$ requires $O(N \cdot D_{\text{rff}})$ operations, and matrix multiplication with $\Phi(Z_H)^T$ requires $O(N \cdot D_{\text{rff}}^2)$ time and $O(N \cdot D_{\text{rff}})$ memory. For fixed feature dimension $D_{\text{rff}} \ll N$, this guarantees strict linear $O(N)$ computational complexity with respect to batch size N , in contrast to exact Gram matrix evaluation $O(N^2)$.

B IMPLEMENTATION DETAILS AND HYPERPARAMETER SPECIFICATIONS

B.1 NETWORK ARCHITECTURE SPECIFICATIONS

Across all experiments, we employ standard neural network architectures tuned for high-capacity representation extraction:

- **Low-Level Encoder (ϕ_L):** For CMNIST and RMNIST, ϕ_L consists of a 3-layer Convolutional Neural Network (ConvNet with 32, 64, and 128 filters, 3×3 kernels, ReLU activations) followed by dense layers outputting a 256-dimensional latent representation Z_L . For Ball Agent spatial regression, ϕ_L uses a 4-layer ConvNet mapping $64 \times 64 \times 3$ inputs to a 256-dimensional Z_L . For Camelyon17 pathology patches ($96 \times 96 \times 3$), ϕ_L uses a ResNet-18 backbone pre-trained on ImageNet, outputting a 512-dimensional Z_L .
- **High-Level Encoder (ϕ_H):** Implemented as a 2-layer Multi-Layer Perceptron (MLP: $256 \rightarrow 128 \rightarrow 64$ for synthetic benchmarks, and $512 \rightarrow 256 \rightarrow 64$ for Camelyon17) with LayerNorm and LeakyReLU activations (negative_slope = 0.1).
- **Velocity Field Network (v_θ):** Implemented as a 3-layer residual MLP accepting concatenated inputs (Z_L, Y_0, Y_1) with hidden width 256 and SiLU activations.
- **Classifier / Regressor (C):** Linear head mapping 64-dimensional high-level representations Z_H to target label logits or 8-dimensional ball position outputs.

B.2 HYPERPARAMETER REFERENCE TABLE

Table 6 lists complete hyperparameter choices across all four benchmark datasets to ensure exact reproducibility.

Table 6: Detailed hyperparameter configurations for COFM-SIB across benchmark datasets.

Hyperparameter	CMNIST	RMNIST	Ball Agent	Camelyon17
Optimizer	Adam	Adam	Adam	Adam
Learning Rate	1×10^{-4}	1×10^{-4}	1×10^{-4}	5×10^{-5}
Batch Size	128	128	128	128
Training Epochs	30	15	40	5
λ_1 ($R_{1,\text{CSWS}}$ weight)	0.60	0.50	0.30	0.50
λ_2 ($R_{2,\text{OTFM}}$ weight)	0.60	0.20	0.05	0.10
λ_3 (R_{HSIC} weight)	0.10	0.20	0.10	0.20
λ_{mean} (R_{MEAN} weight)	8.00	60.00	0.30	10.00
TSGP Cap (α)	2.50	1.50	0.30	2.00
EMA Decay Rate	0.9995	0.9990	Disabled (0.0)	0.9990
RFF Dimensions (D_{rff})	256	256	256	256
Random Projections (M)	64	64	64	64
Normalization Layer	GroupNorm	GroupNorm	LayerNorm	BatchNorm

C ADDITIONAL ABLATION STUDIES

We report further diagnostic studies that complement the main-text ablations (Section 4.4). All runs use seed 42 with deterministic cuDNN and share the tuned per-dataset COFM-SIB configuration of Table 6; only the studied factor is varied.

Ablation 4: Sobolev Weighting and Projection Density (M) in CSWS. On Ball Agent continuous spatial regression, we evaluate projection count $M \in \{8, 32, 64, 128, 512\}$, Sobolev weighting, and 1D sorting implementations in Table 7. Increasing projection density from $M = 8$ to $M = 64$ improves environment independence ($\text{EI} = 0.1423 \rightarrow 0.1359$), with diminishing returns at $M = 512$ ($\text{EI} = 0.1355$). Discrete label binning achieves lower position MAE than continuous label-distance sorting (0.0213 vs. 0.0223).

Ablation 5: RFF-HSIC Bottleneck Dimensionality and Kernel Scaling. Evaluating R_{HSIC} on Camelyon17 across RFF dimensions $D_{\text{rff}} \in \{16, 64, 256, 1024\}$ and bandwidth scale factor σ (Table 8) shows that peak GPU memory remains virtually constant (437.8–439.1 MB), confirming that RFF feature buffers add minimal overhead. Micro-benchmarking at batch size $N = 8192$ reveals that exact $O(N^2)$ MMD consumes 577.9 MB GPU memory ($4.8\times$ higher than RFF’s 121.0 MB) and requires 29.9 ms vs. RFF’s 7.4 ms, proving $O(N)$ linear scalability.

Ablation 6: EMA $\times$ Normalization Layer Interactions. Evaluating Exponential Moving Average (EMA decay = 0.999) across BatchNorm, LayerNorm, and GroupNorm on CMNIST classification




Figure 3: Multi-metric performance comparison of COFM-SIB against ACIA and competitive base-lines across four anti-causal benchmarks: Continuous CMNIST, Rotated MNIST, Ball Agent, and Camelyon17. Each radar plot displays normalized scores scaled from 0 to 100 across four evaluation criteria, where larger radial values represent superior performance: Test Accuracy, Environment Independence, Low-Level Invariance, and Intervention Robustness. COFM-SIB demonstrates optimal trade-offs between predictive accuracy and causal representation invariance across diverse classification, continuous spatial regression, and real-world domain-shift settings.

Table 7: Hyperparameter sensitivity of projection directions M , Sobolev derivative weighting, and sorting formulations in $R_{1,CSWS}$ on Ball Agent spatial position regression.

Setting	M	Sobolev	Sort Mode	Acc $\uparrow$ (%)	MAE $\downarrow$	EI $\downarrow$	LLI $\downarrow$	s/epoch
$M = 8$, Sobolev, discrete	8	On	Discrete	99.40	0.0202	0.1423	0.1355	6.90
$M = 32$, Sobolev, discrete	32	On	Discrete	99.38	0.0209	0.1368	0.1336	6.80
$M = 64$, Sobolev, discrete	64	On	Discrete	99.32	0.0213	0.1359	0.1327	7.14
$M = 128$, Sobolev, discrete	128	On	Discrete	99.31	0.0213	0.1354	0.1334	7.16
$M = 512$, Sobolev, discrete	512	On	Discrete	99.27	0.0213	0.1355	0.1332	7.10
$M = 64$, No Sobolev, discrete	64	Off	Discrete	99.32	0.0213	0.1362	0.1329	6.64
$M = 64$, Sobolev, continuous	64	On	Continuous	99.19	0.0223	0.1372	0.1347	5.59

and Ball Agent regression reveals a critical interaction (Table 9 and Figure 7). On discrete CMNIST classification, EMA provides consistent accuracy boosts across all norm layers (+0.08% to +0.19%). However, on Ball Agent continuous spatial regression, enabling EMA under BatchNorm triggers a severe -50.49% accuracy collapse ($99.32\% \rightarrow 48.83\%$, MAE $0.0213 \rightarrow 0.1321$). Isolating the mechanism proves that averaging BatchNorm’s running feature statistics introduces severe




Figure 4: 3D anti-causal generalization on Causal3DIdent (visualization of Table 2). COFM-SIB matches ACIA’s best out-of-distribution accuracy (left), achieves the lowest intervention robustness error among all methods (center), and does so with a 4.5× reduction in training time relative to ACIA’s multi-step interventional consistency solver (right, log scale).




Figure 5: Optimization dynamics and evaluation of Task-Shielded Dual Gradient Projection (TSGP) on CMNIST across training epochs and conflict regimes. Panel (a) shows the mean cosine similarity between regularizer and task gradients across training epochs for Naive Joint Optimization and Full TSGP. Panel (b) tracks the L_2 gradient norms of task and regularizer components, showing similar convergence behavior for both optimization approaches. Panel (c) illustrates the conflict rate and cap-activation rate over epochs, highlighting how projection capping engages as the task gradient magnitude decreases in later stages of training. Panel (d) compares gradient conflict rates under standard and high-conflict stress setups with ten-fold regularizer weighting, demonstrating that TSGP effectively bounds regularizer gradient interference to protect predictive performance.

phase lag during spatial regression. Buffer-free normalization layers (LayerNorm and GroupNorm) completely prevent this collapse, maintaining $\geq 99.4\%$ position accuracy regardless of EMA state.

Ablation 7: Late-Stage TSGP Gradient Dynamics Near Convergence. To directly address whether the task-scale magnitude cap (Equation 9) systematically drives $g_{\text{reg}}^{\text{final}} \rightarrow 0$ as the task loss approaches a local minimum ($\|g_{\text{task}}\|_2 \rightarrow 0$), we extend CMNIST training to 100 epochs (3.3× the




Figure 6: Counterfactual transport strategy wall-clock overhead and fidelity comparison on RMNIST (visualization of Table 5). Left: Per-step transport overhead across different counterfactual transport implementations, showing that 1-step OT-CFM achieves the lowest latency at 1.878 ms per step. Right: Trade-off between transport overhead and test accuracy, demonstrating that 1-step OT-CFM lies within the Pareto-optimal efficiency zone and matches the downstream predictive accuracy of computationally intensive multi-step ODE solvers at a fraction of the wall-clock cost.

Table 8: Random Fourier Feature dimension D_{rff} sweep, bandwidth scale σ , and linear $O(N)$ vs. quadratic $O(N^2)$ MMD computational scaling comparison on Camelyon17.

Configuration / Feature Dim	Acc $\uparrow$ (%)	LLI $\downarrow$	EI $\downarrow$	IR $\downarrow$	Peak GPU Mem (MB)	Time/epoch (s)
RFF-HSIC ($D_{\text{rff}} = 16$)	88.44	0.018	0.326	0.016	437.8	8.5
RFF-HSIC ($D_{\text{rff}} = 64$)	87.57	0.016	0.292	0.019	438.8	8.4
RFF-HSIC ($D_{\text{rff}} = 256, \sigma = 1.0\sigma_0$)	88.20	0.016	0.311	0.012	438.9	8.4
RFF-HSIC ($D_{\text{rff}} = 1024$)	87.21	0.018	0.300	0.016	439.1	8.4
Exact MMD Bottleneck ($O(N^2)$)	88.45	0.017	0.315	0.016	438.8	11.9

Table 9: Interaction of Exponential Moving Average (EMA) parameter tracking and normalization layer selection across discrete classification (CMNIST) and continuous spatial position regression (Ball Agent).

Norm Layer	EMA State	CMNIST Classification		Ball Agent Spatial Regression		
		Acc $\uparrow$ (%)	EI $\downarrow$	Acc $\uparrow$ (%)	MAE $\downarrow$	EI $\downarrow$
BatchNorm	Disabled	99.13	0.006	99.32	0.0213	0.136
BatchNorm	Enabled	99.20	0.002	48.83	0.1321	0.154
LayerNorm	Disabled	99.13	0.003	99.89	0.0133	0.138
LayerNorm	Enabled	99.30	0.001	99.48	0.0215	0.140
GroupNorm	Disabled	99.21	0.002	99.80	0.0146	0.130
GroupNorm	Enabled	99.40	0.001	99.41	0.0212	0.129

nominal schedule) and track gradient norms, the regularizer scaling factor, conflict/cap activation rates, and invariance metrics (Table 10, Figure 8). Two mechanisms prevent invariance collapse. First, as representations consolidate into environment-invariant clusters, $\|g_{\text{reg}}^{\text{shielded}}\|_2$ decays *in tandem* with $\|g_{\text{task}}\|_2$, so their ratio remains stable at ≈ 2.23 throughout late training. Second, since the cap in Equation 9 activates only when $\|g_{\text{reg}}^{\text{shielded}}\|_2 > \alpha \|g_{\text{task}}\|_2$ with $\alpha = 2.5$, the balanced ratio keeps the cap inactive (0.0%) after epoch 10. Consequently, Environment Independence continues to improve (EI : 0.084 $\rightarrow$ 0.0007) rather than regressing, and the conflict rate steadily shrinks (12.0% $\rightarrow$ 1.8%), confirming that invariance regularization stays fully active and stable near convergence.




Figure 7: Interaction of normalization layer selection and Exponential Moving Average (EMA) parameter updates on Ball Agent position regression accuracy. Combining BatchNorm with EMA leads to a 50.49% performance collapse (falling to 48.83%) due to running statistic lag during continuous spatial regression. In contrast, buffer-free normalization layers such as LayerNorm and GroupNorm remain highly stable, maintaining over 99.4% position accuracy regardless of EMA state.

Table 10: Late-stage TSGP gradient dynamics over 100 epochs on CMNIST. Task and regularizer gradient norms decay in tandem at a stable ratio; the magnitude cap becomes inactive while invariance (EI) keeps improving, showing regularization is not disabled near convergence.

Epoch Range	$\ g_{\text{task}}\ _2$	$\ g_{\text{reg}}^{\text{shielded}}\ _2$	Ratio	Active Cap Rate	Conflict Rate	Test Acc $\uparrow$	EI $\downarrow$
Epochs 1–10	1.241	0.582	2.13	4.4%	12.0%	98.15%	0.0840
Epochs 20–30	0.984	0.441	2.23	0.0%	5.2%	99.20%	0.0010
Epochs 50–60	0.812	0.365	2.22	0.0%	3.1%	99.21%	0.0008
Epochs 90–100	0.743	0.331	2.24	0.0%	1.8%	99.20%	0.0007




Figure 8: Late-stage TSGP gradient dynamics on CMNIST over 100 epochs. (a) Task and shielded-regularizer gradient norms decay in tandem, keeping their ratio stable (green, right axis) below the $\alpha = 2.5$ cap threshold. (b) Environment Independence (EI, log scale) continues to improve while both the active cap rate and the conflict rate fall toward zero, demonstrating that invariance regularization remains active and stable rather than being disabled as $\|g_{\text{task}}\|_2$ shrinks near convergence.